

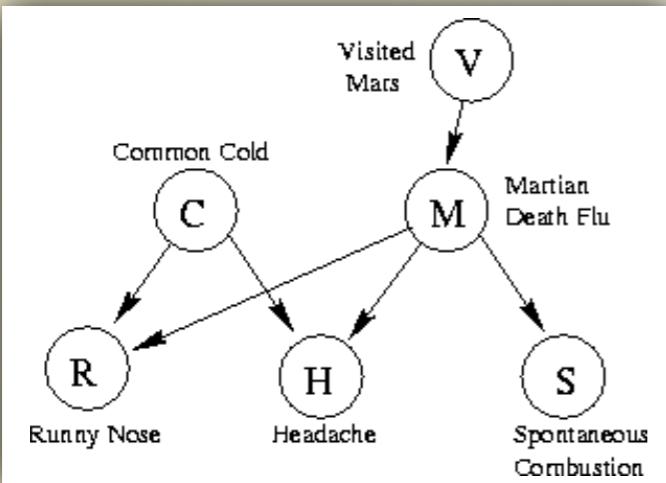
The surprising effectiveness of lenses in machine learning

Fabio Zanasi
University College London

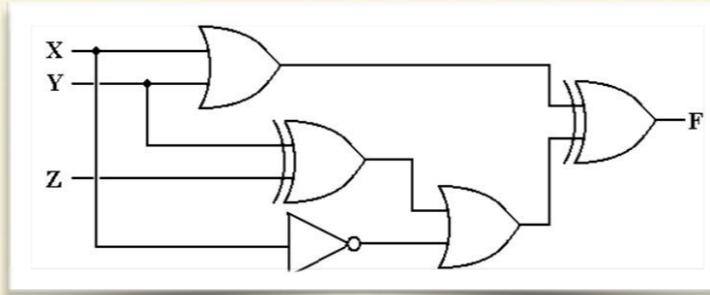
Based on joint work with
Geoff Cruttwell, Bruno Gavranovic, Neil Ghani, Paul Wilson

Background: Categorical Foundations of Component-based systems

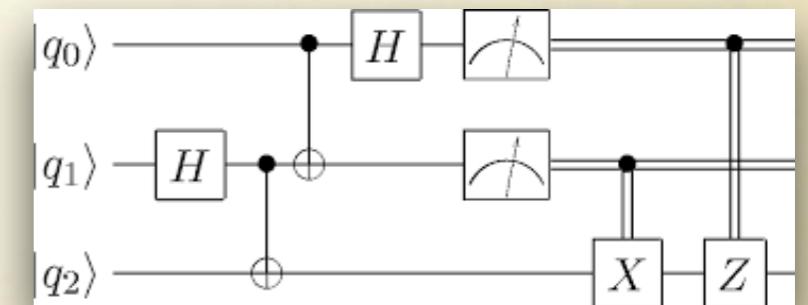
Bayesian networks



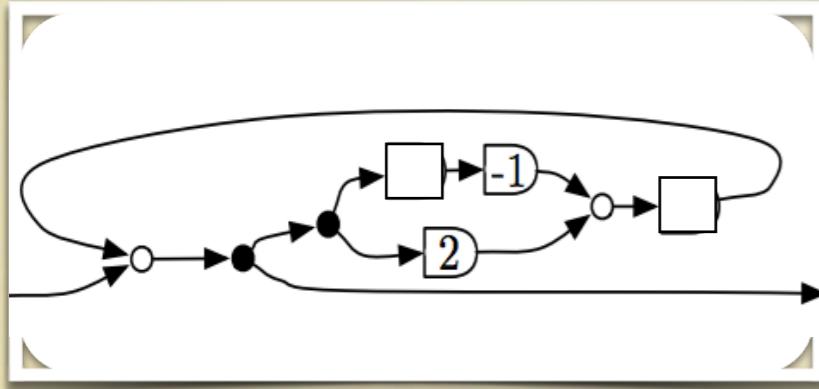
Digital Circuits



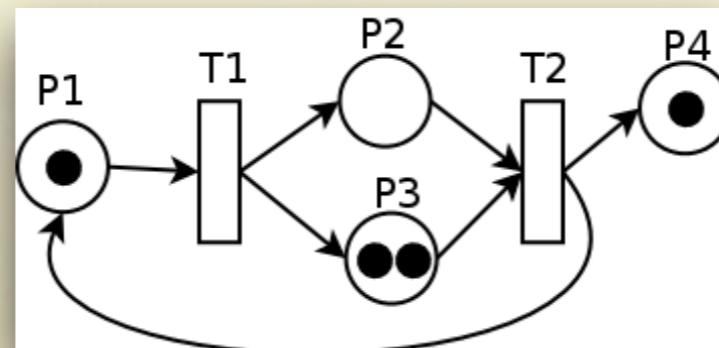
Quantum circuits



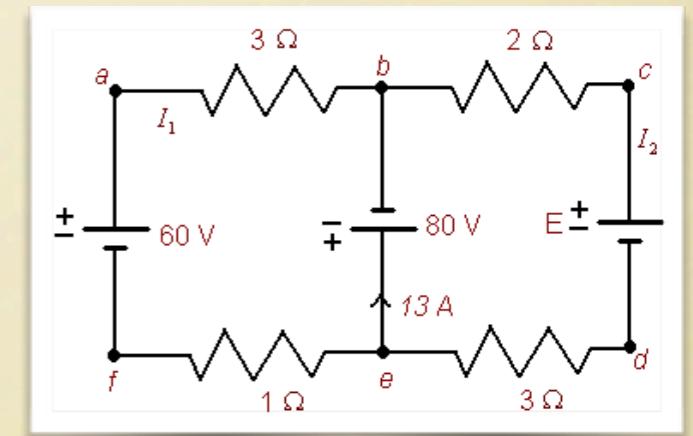
Signal flow graphs



Petri nets



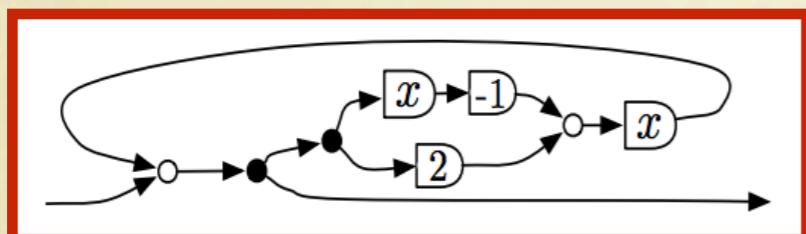
Electrical Circuits



Key Ideas

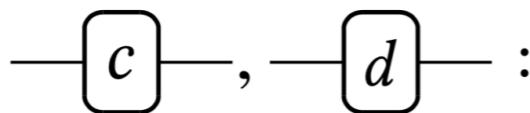
Graphical notation becomes **formal syntax**
and it is given **compositional** semantics

Example: Signal Flow Graphs



Key Ideas

Graphical notation becomes **formal syntax**
and it is given **compositional** semantics

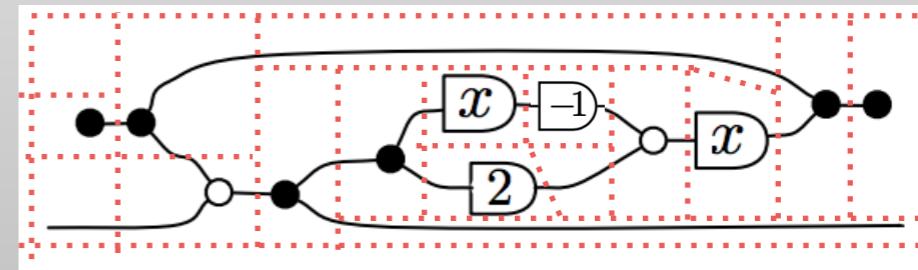
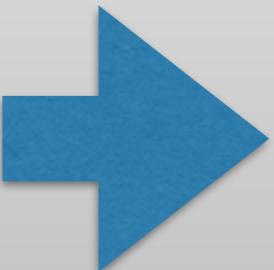
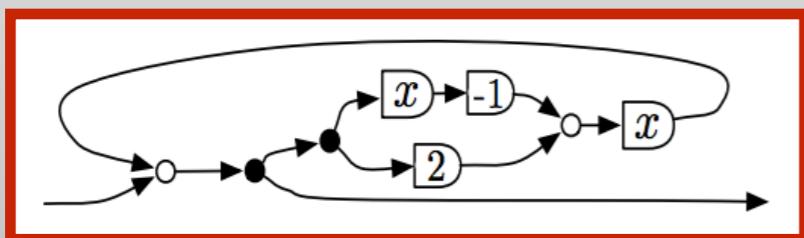


$\cdots ::= \bullet | \circ | \text{---} | \text{---} | \text{---}$

$| \bullet | \circ | \text{---} | \text{---} | \text{---}$

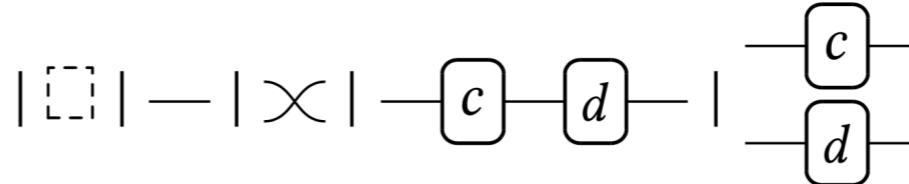
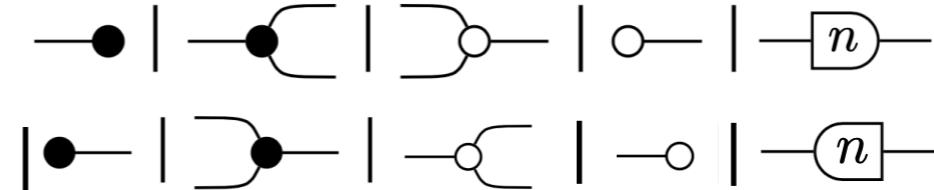
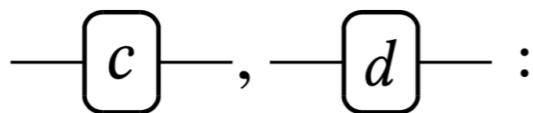
$| \square | \times | \square | \square | \square$

Example:
Signal Flow Graphs

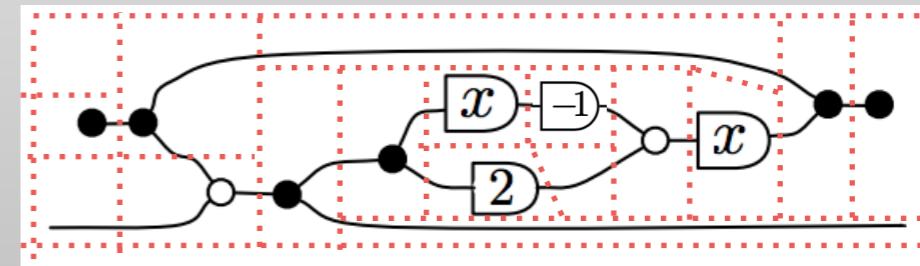
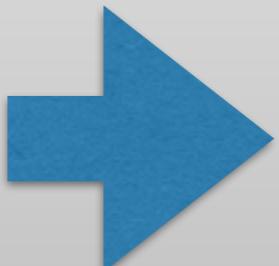
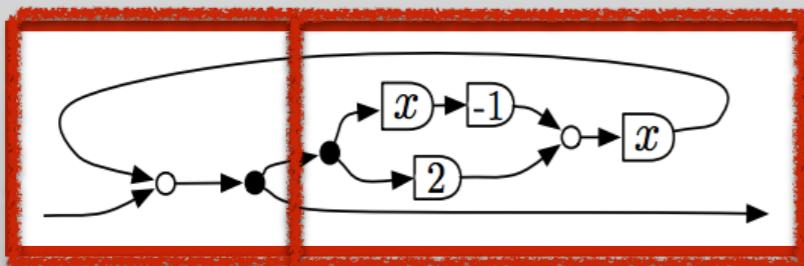


Key Ideas

Graphical notation becomes **formal syntax**
and it is given **compositional** semantics

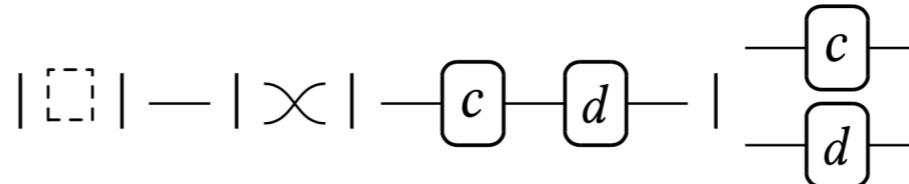
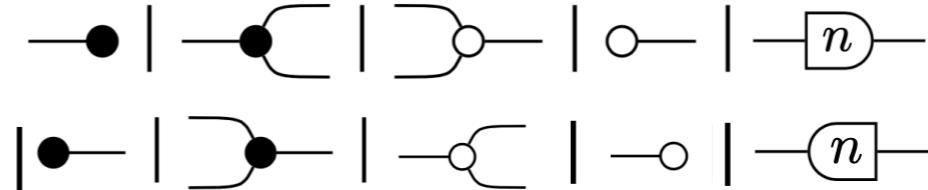
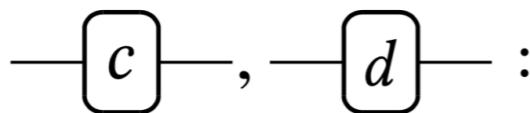


Example:
Signal Flow Graphs

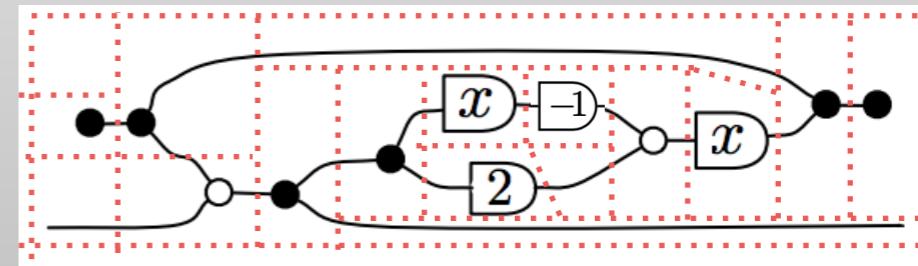
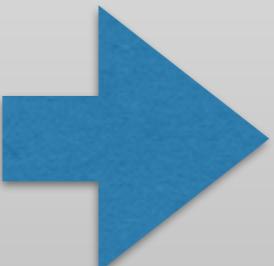
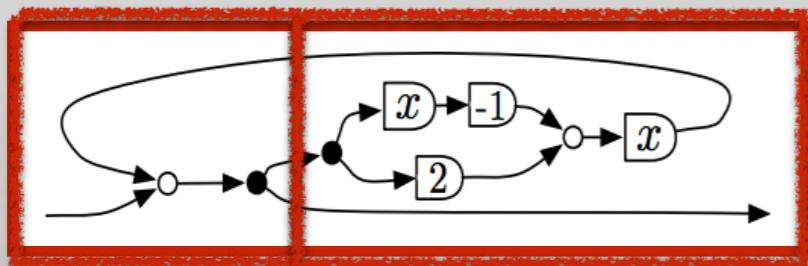


Key Ideas

Graphical notation becomes **formal syntax**
and it is given **compositional** semantics



Example:
Signal Flow Graphs

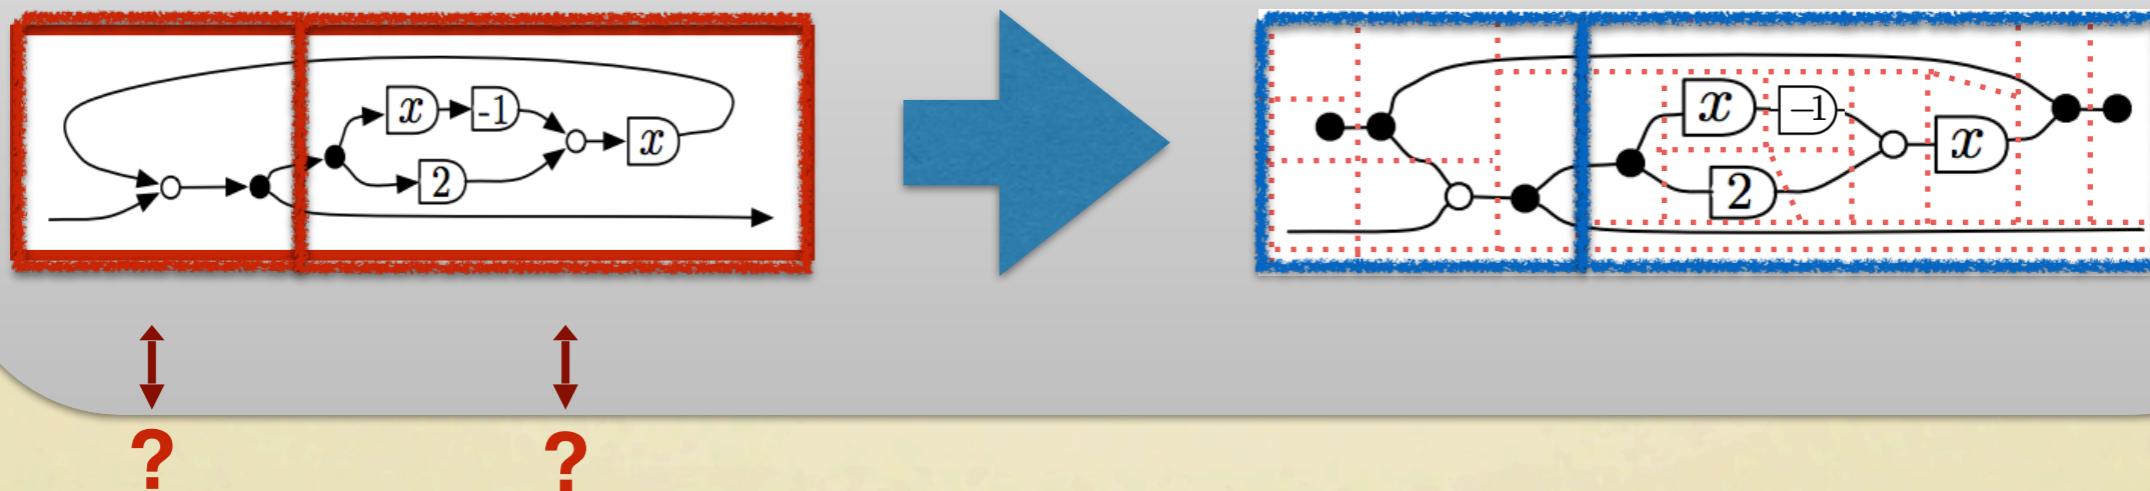
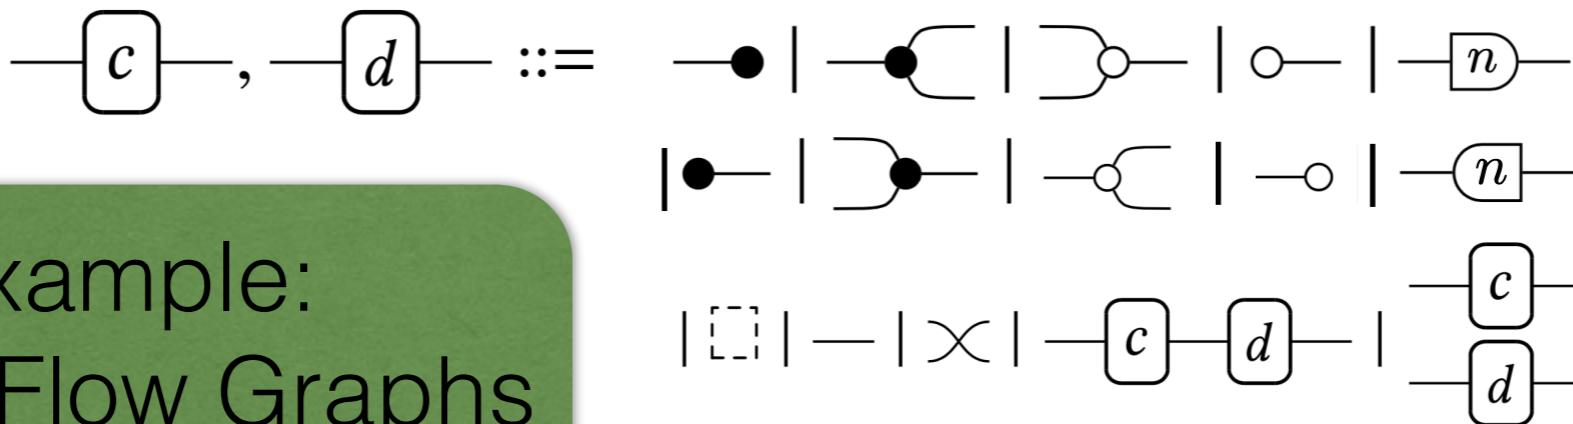


?

?

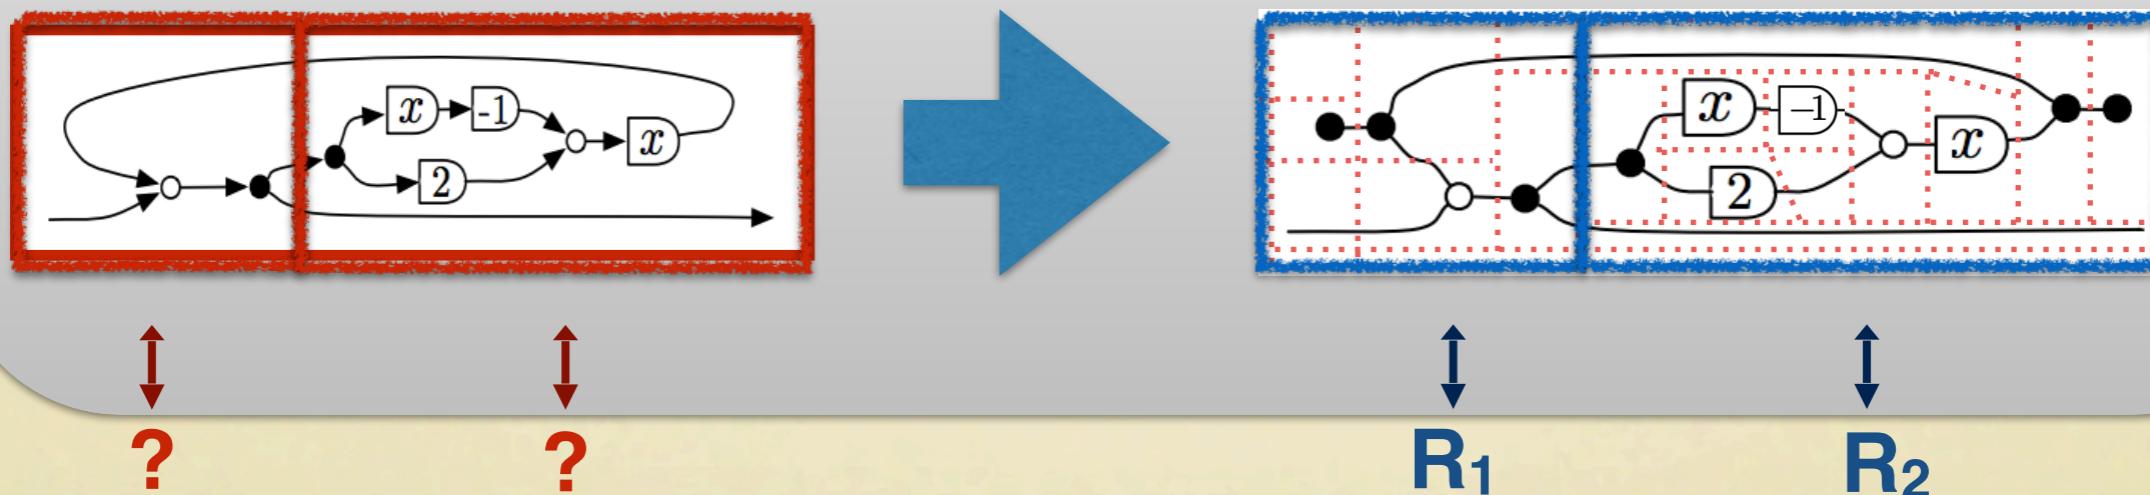
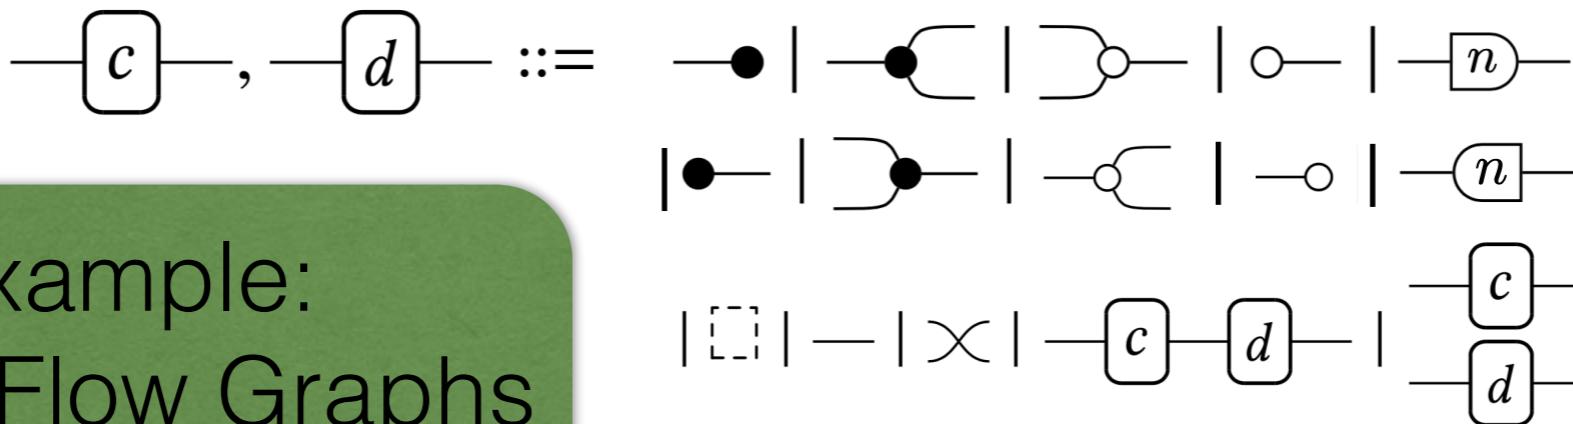
Key Ideas

Graphical notation becomes **formal syntax**
and it is given **compositional** semantics



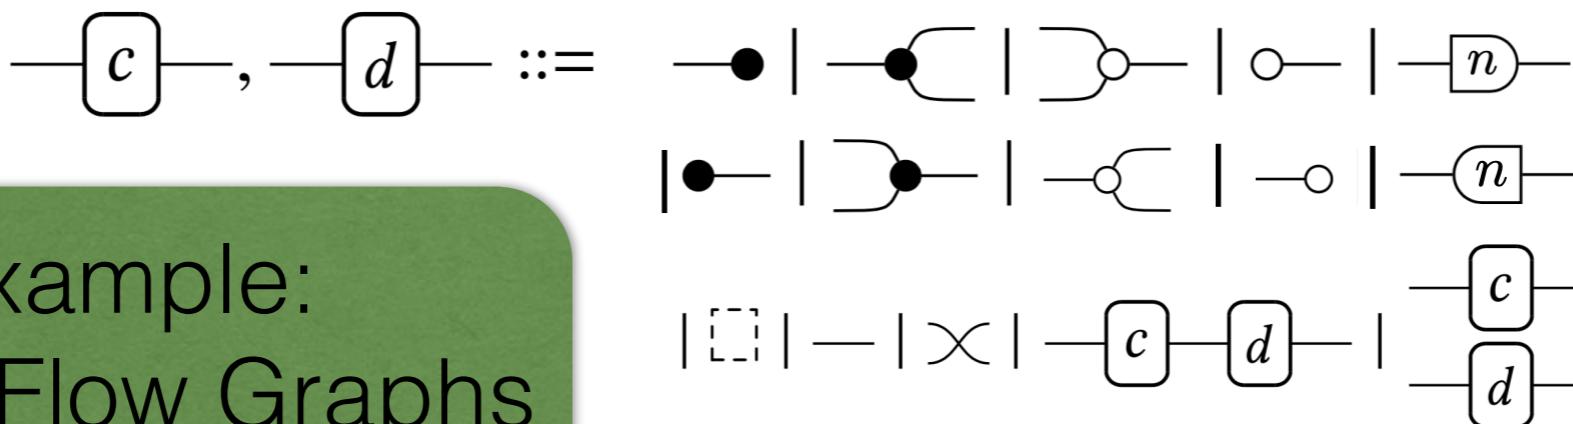
Key Ideas

Graphical notation becomes **formal syntax**
and it is given **compositional** semantics

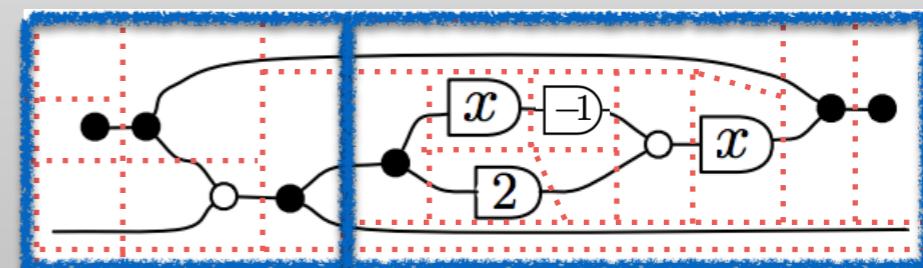
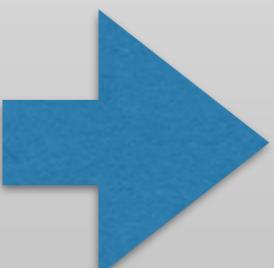
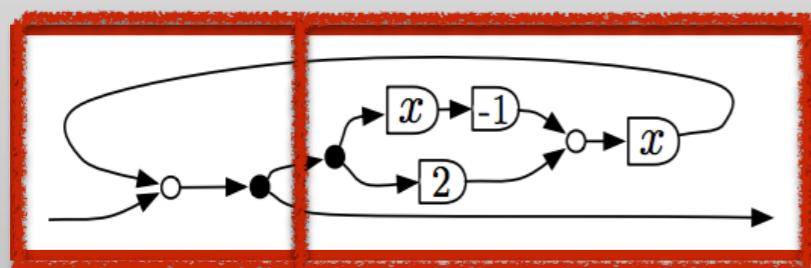


Key Ideas

Graphical notation becomes **formal syntax**
and it is given **compositional** semantics



Example:
Signal Flow Graphs



?

?

R_1

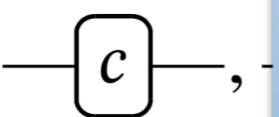
;

R_2

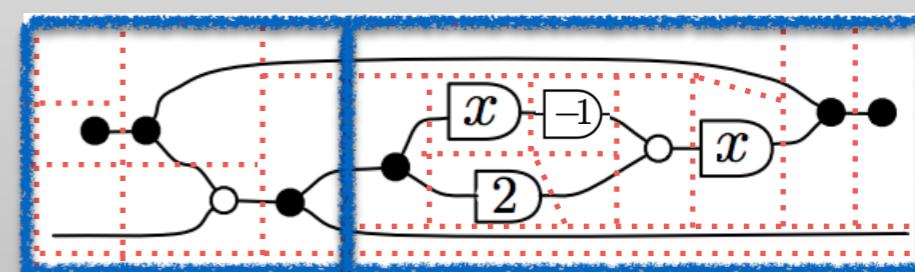
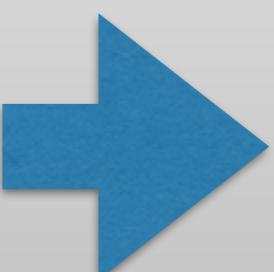
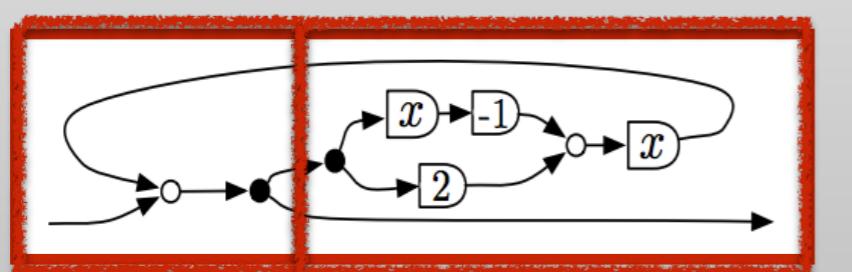
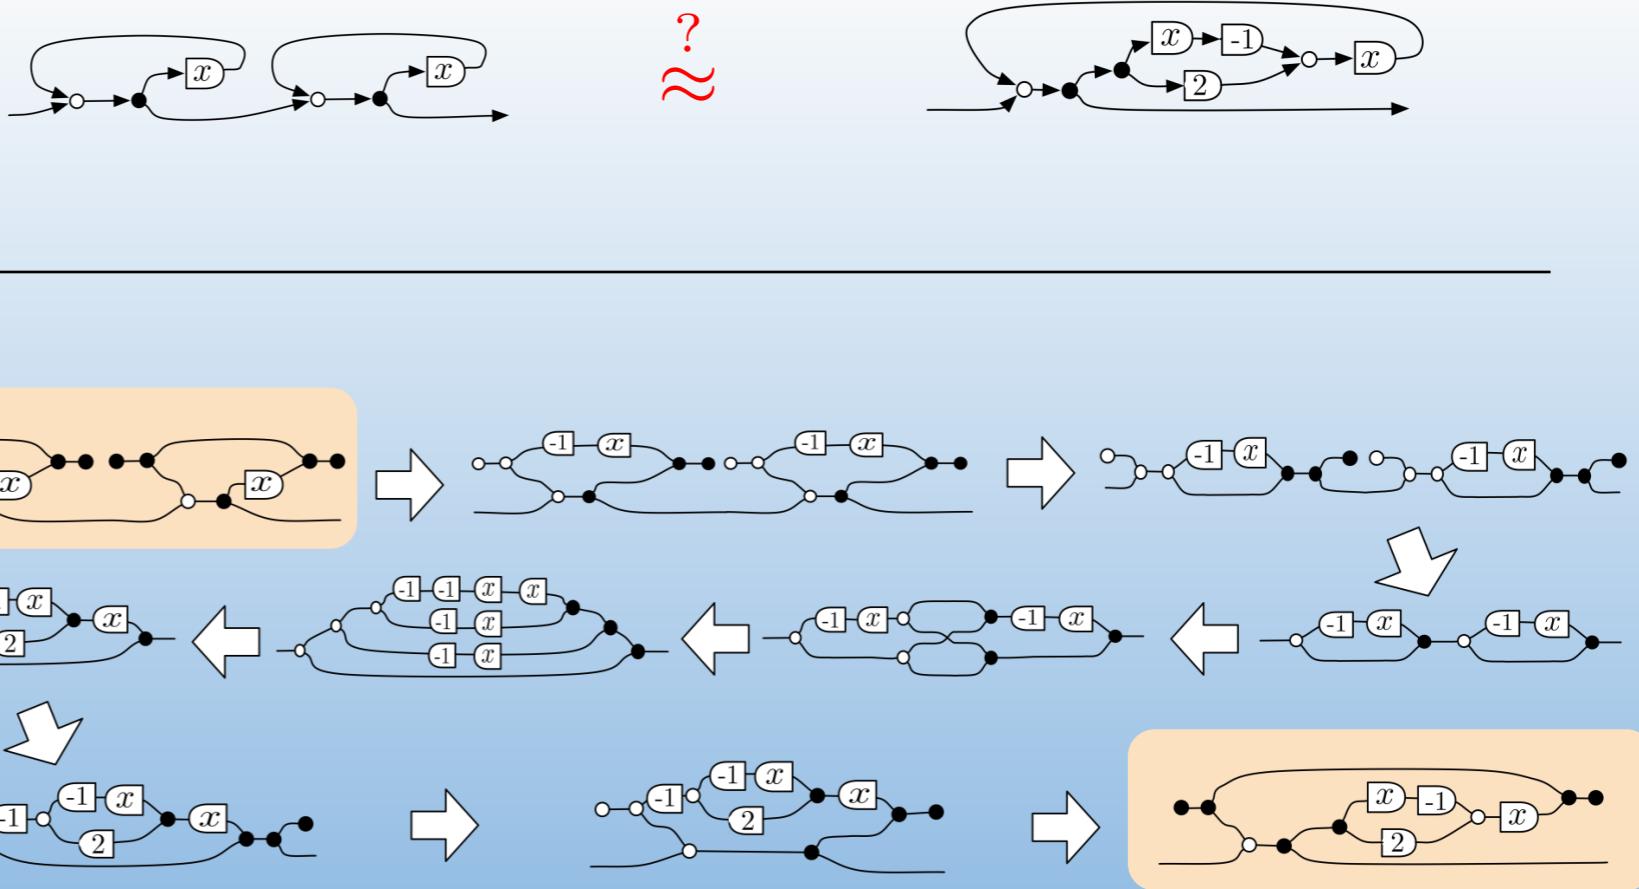
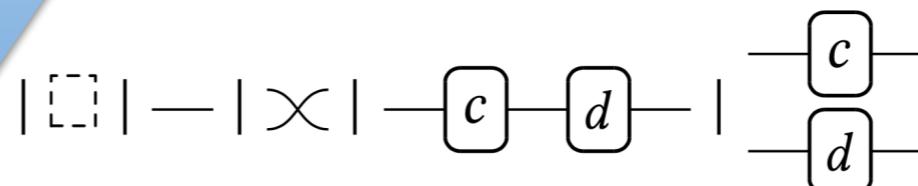
=

R

Graphica and it is



Example:
Signal Flow Graphs



?

?

R_1

;

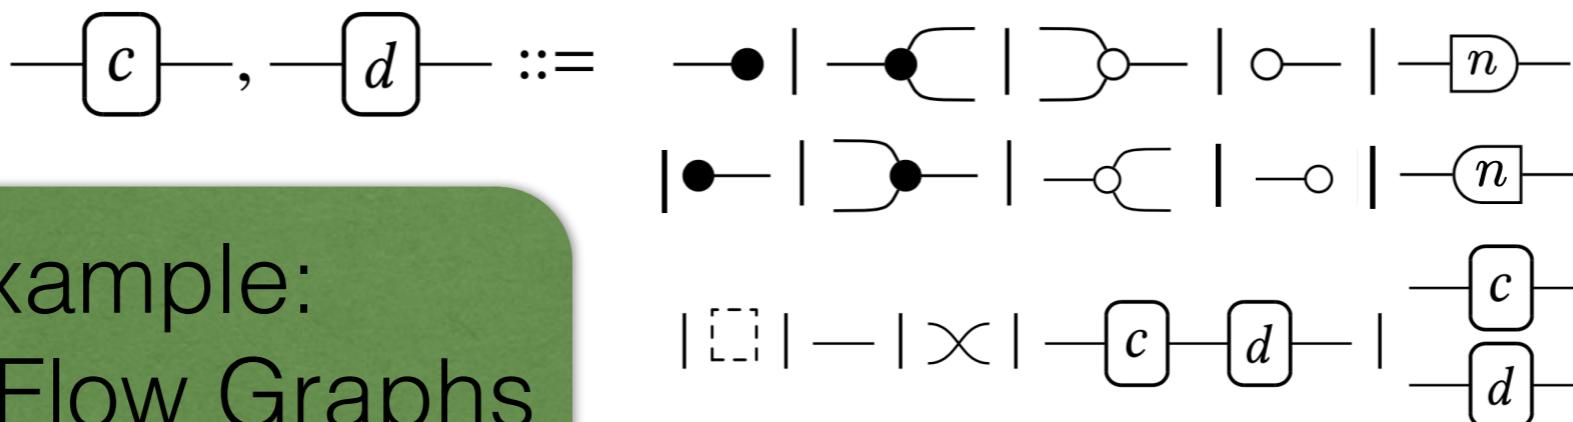
R_2

=

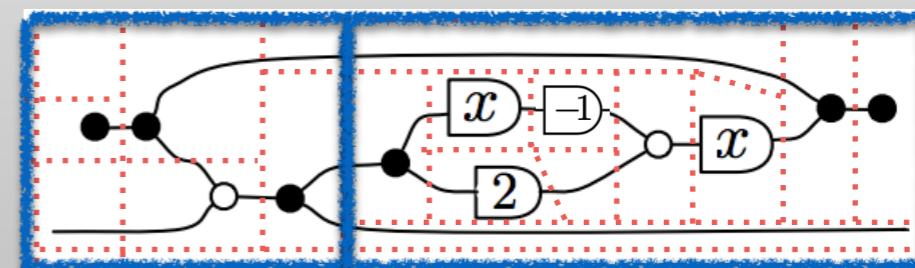
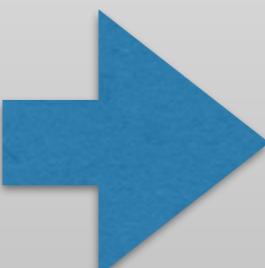
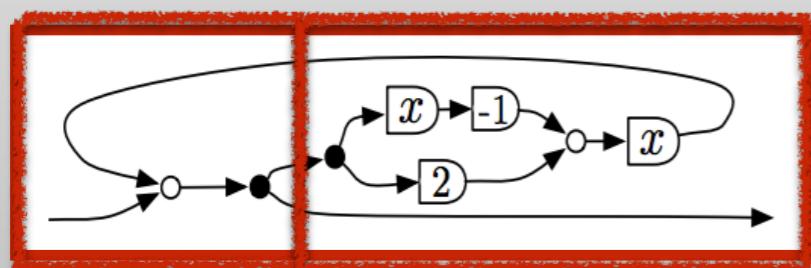
R

Key Ideas

Graphical notation becomes **formal syntax**
and it is given **compositional** semantics



Example:
Signal Flow Graphs



?

?

R_1

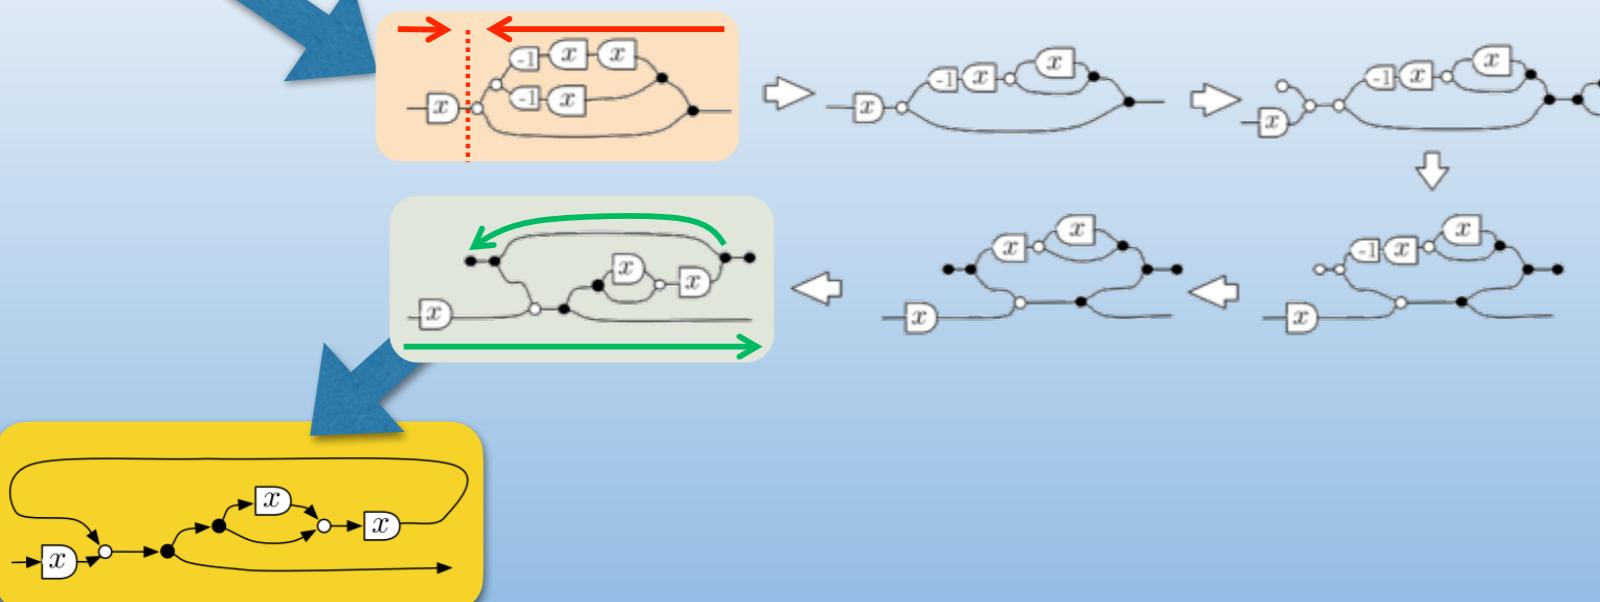
;

R_2

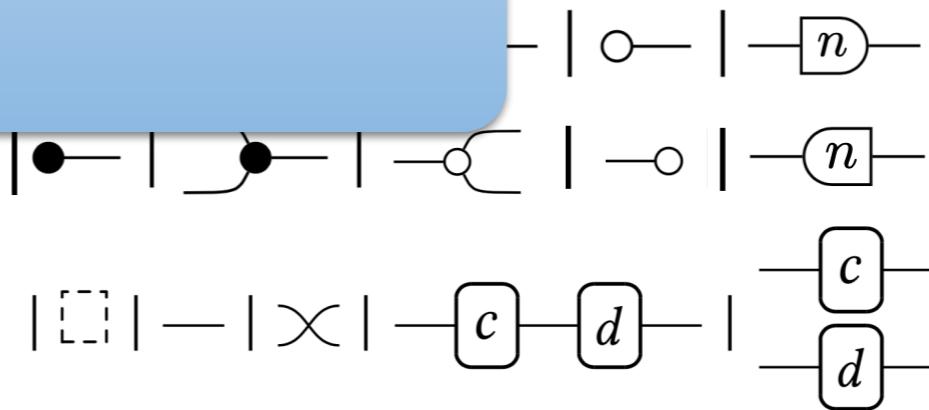
=

R

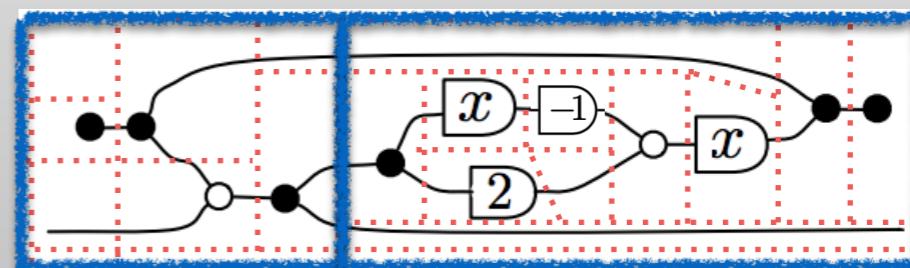
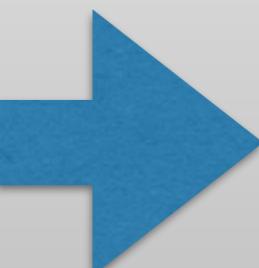
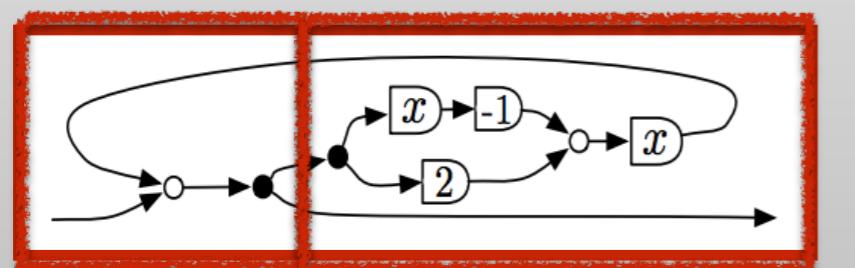
$$\frac{x}{1 - x - x^2}$$



S
ormal syntax
al semantics

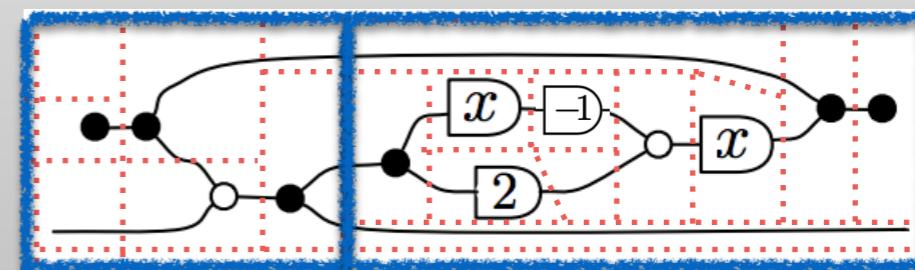
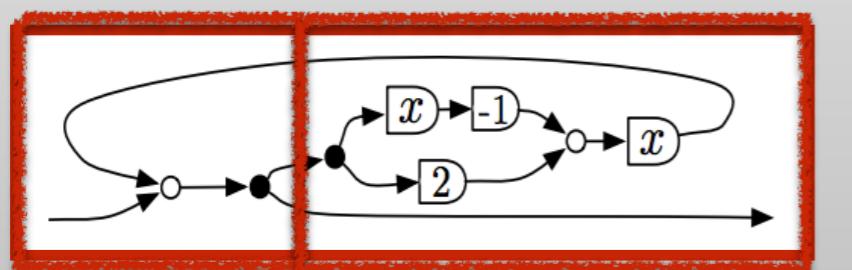
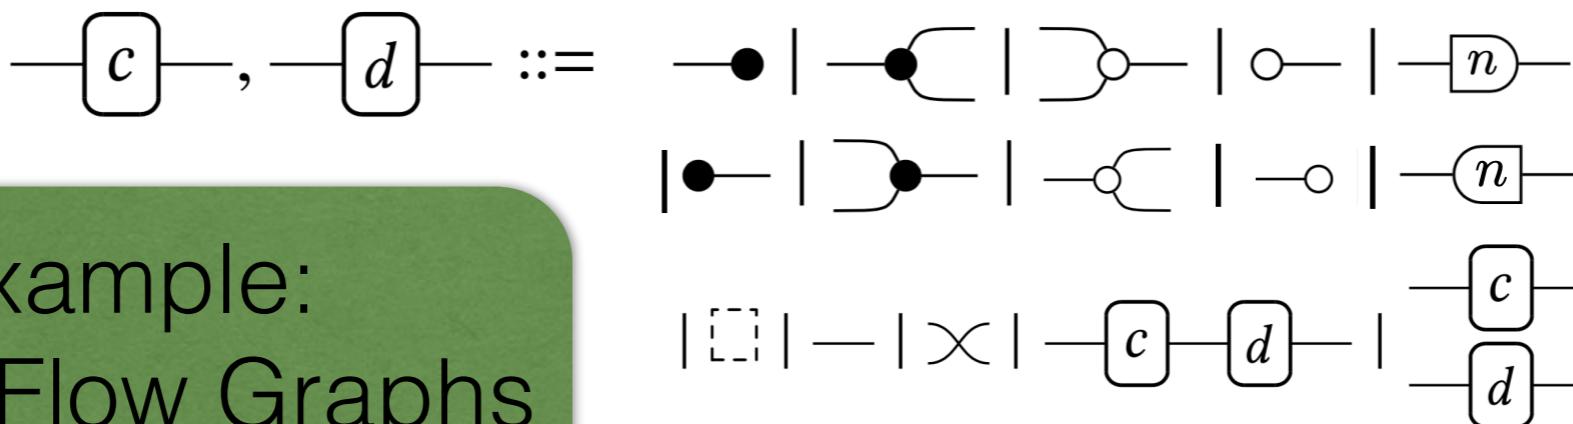


Example:
Signal Flow Graphs



Key Ideas

Graphical notation becomes **formal syntax**
and it is given **compositional** semantics



?

?

R_1

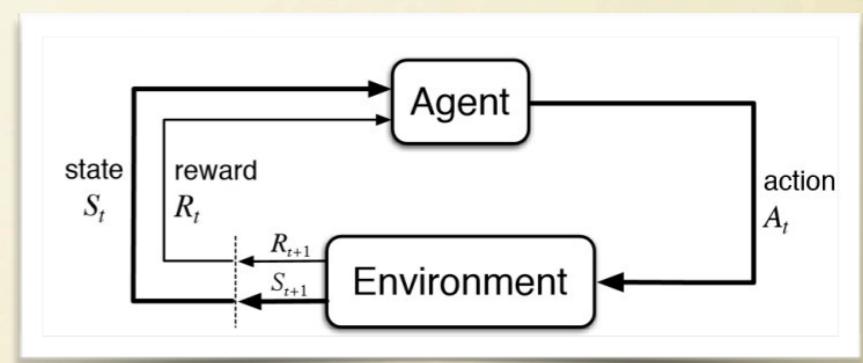
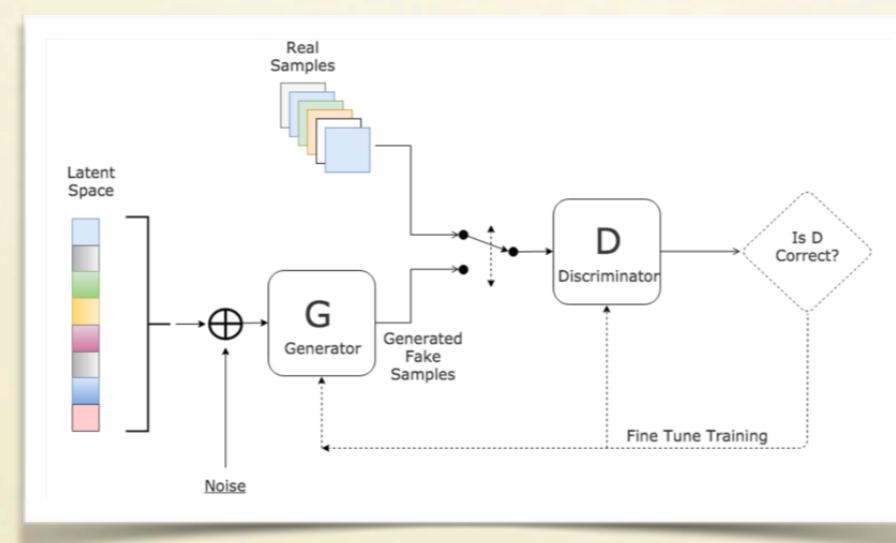
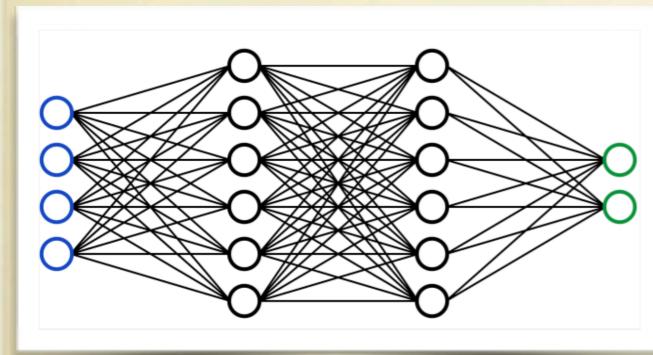
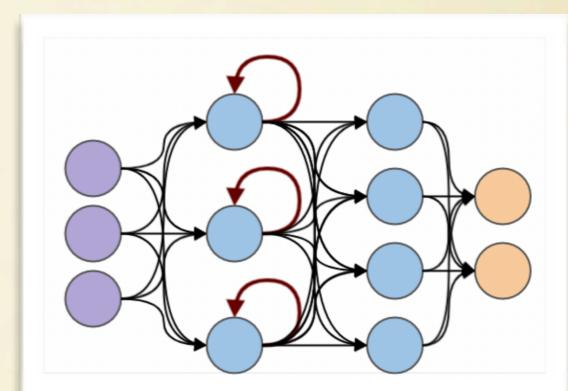
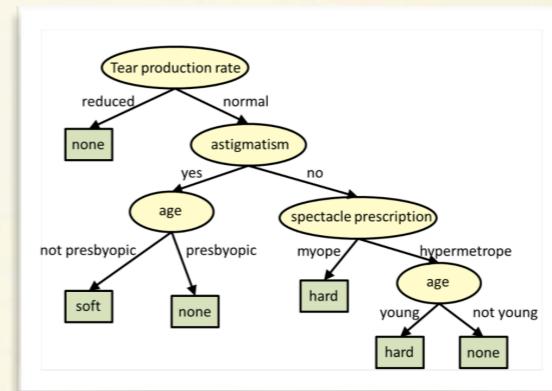
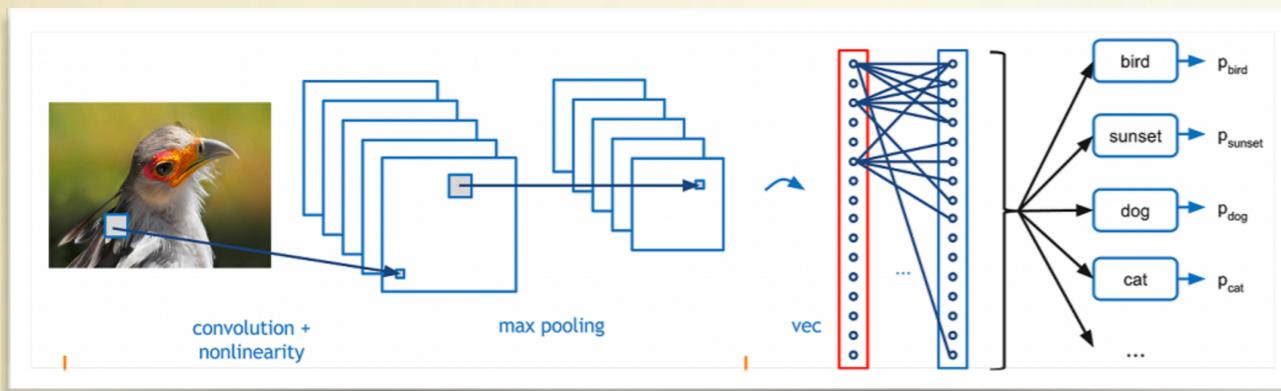
;

R_2

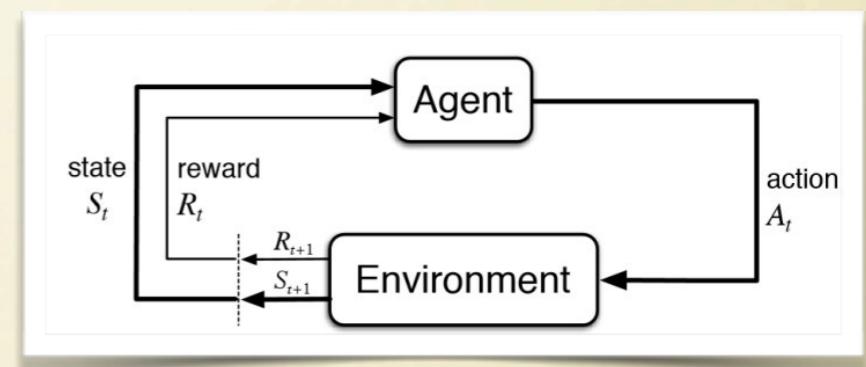
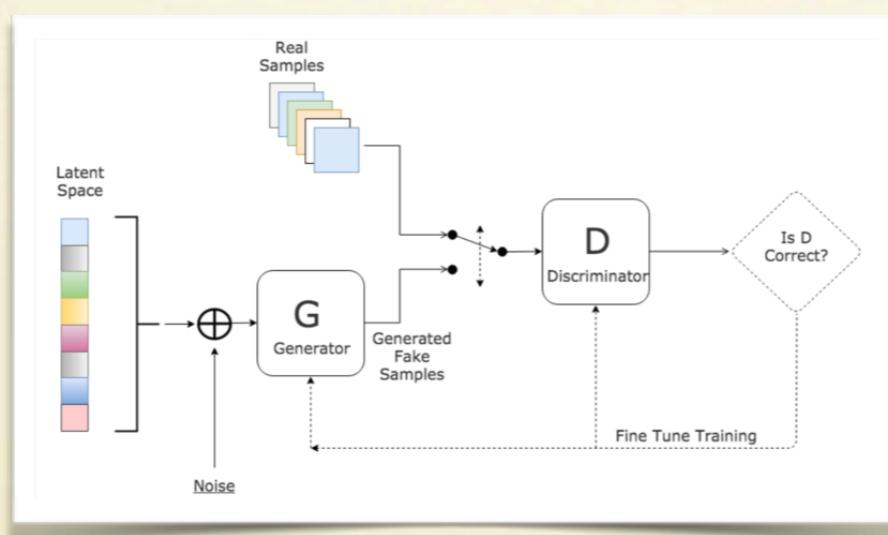
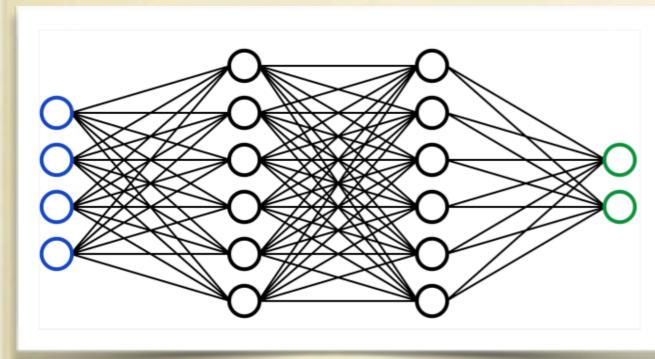
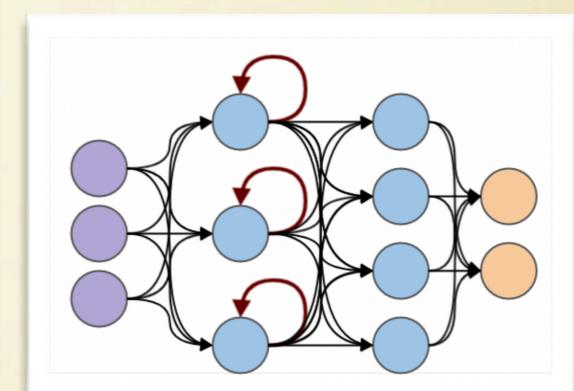
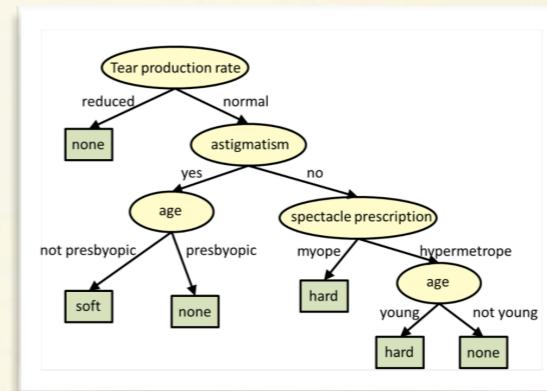
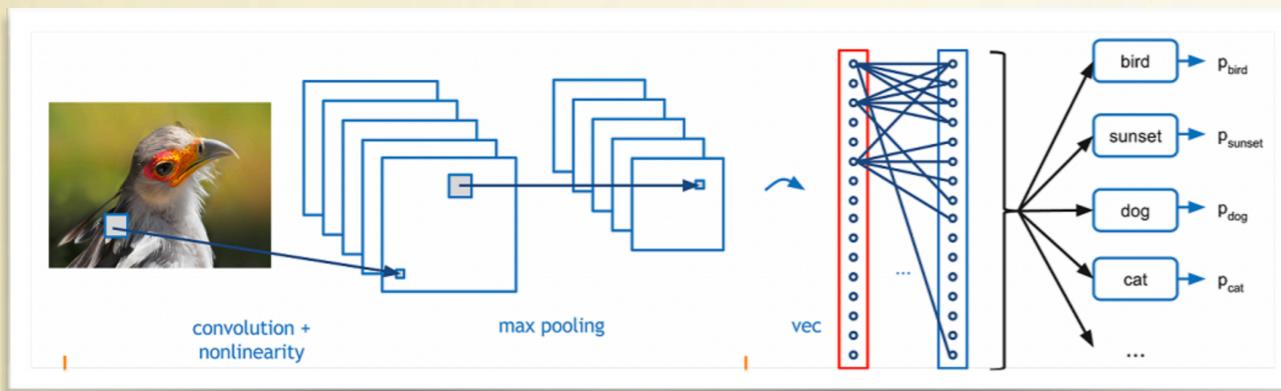
=

R

Categorical Foundations of Machine Learning?



Categorical Foundations of Machine Learning?



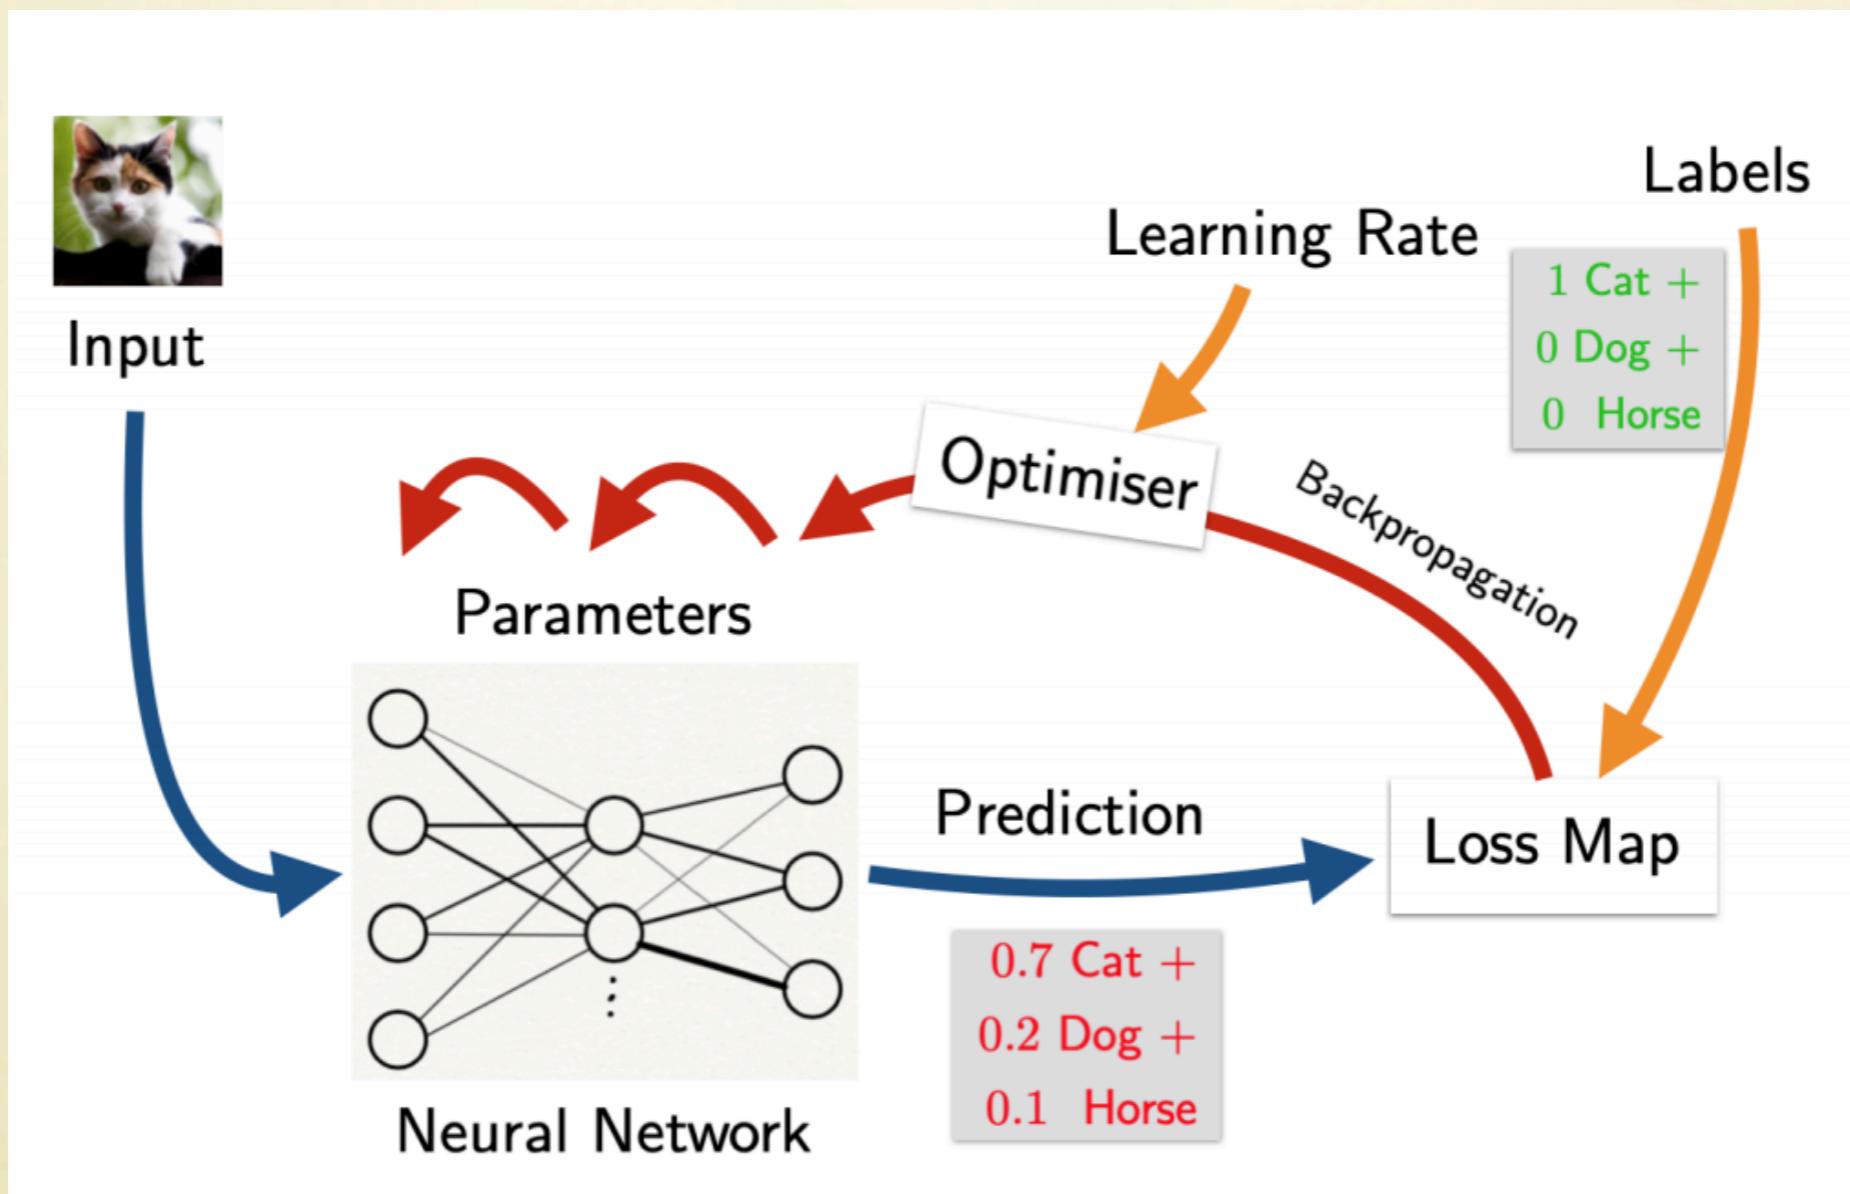
Uniformity

Explainability

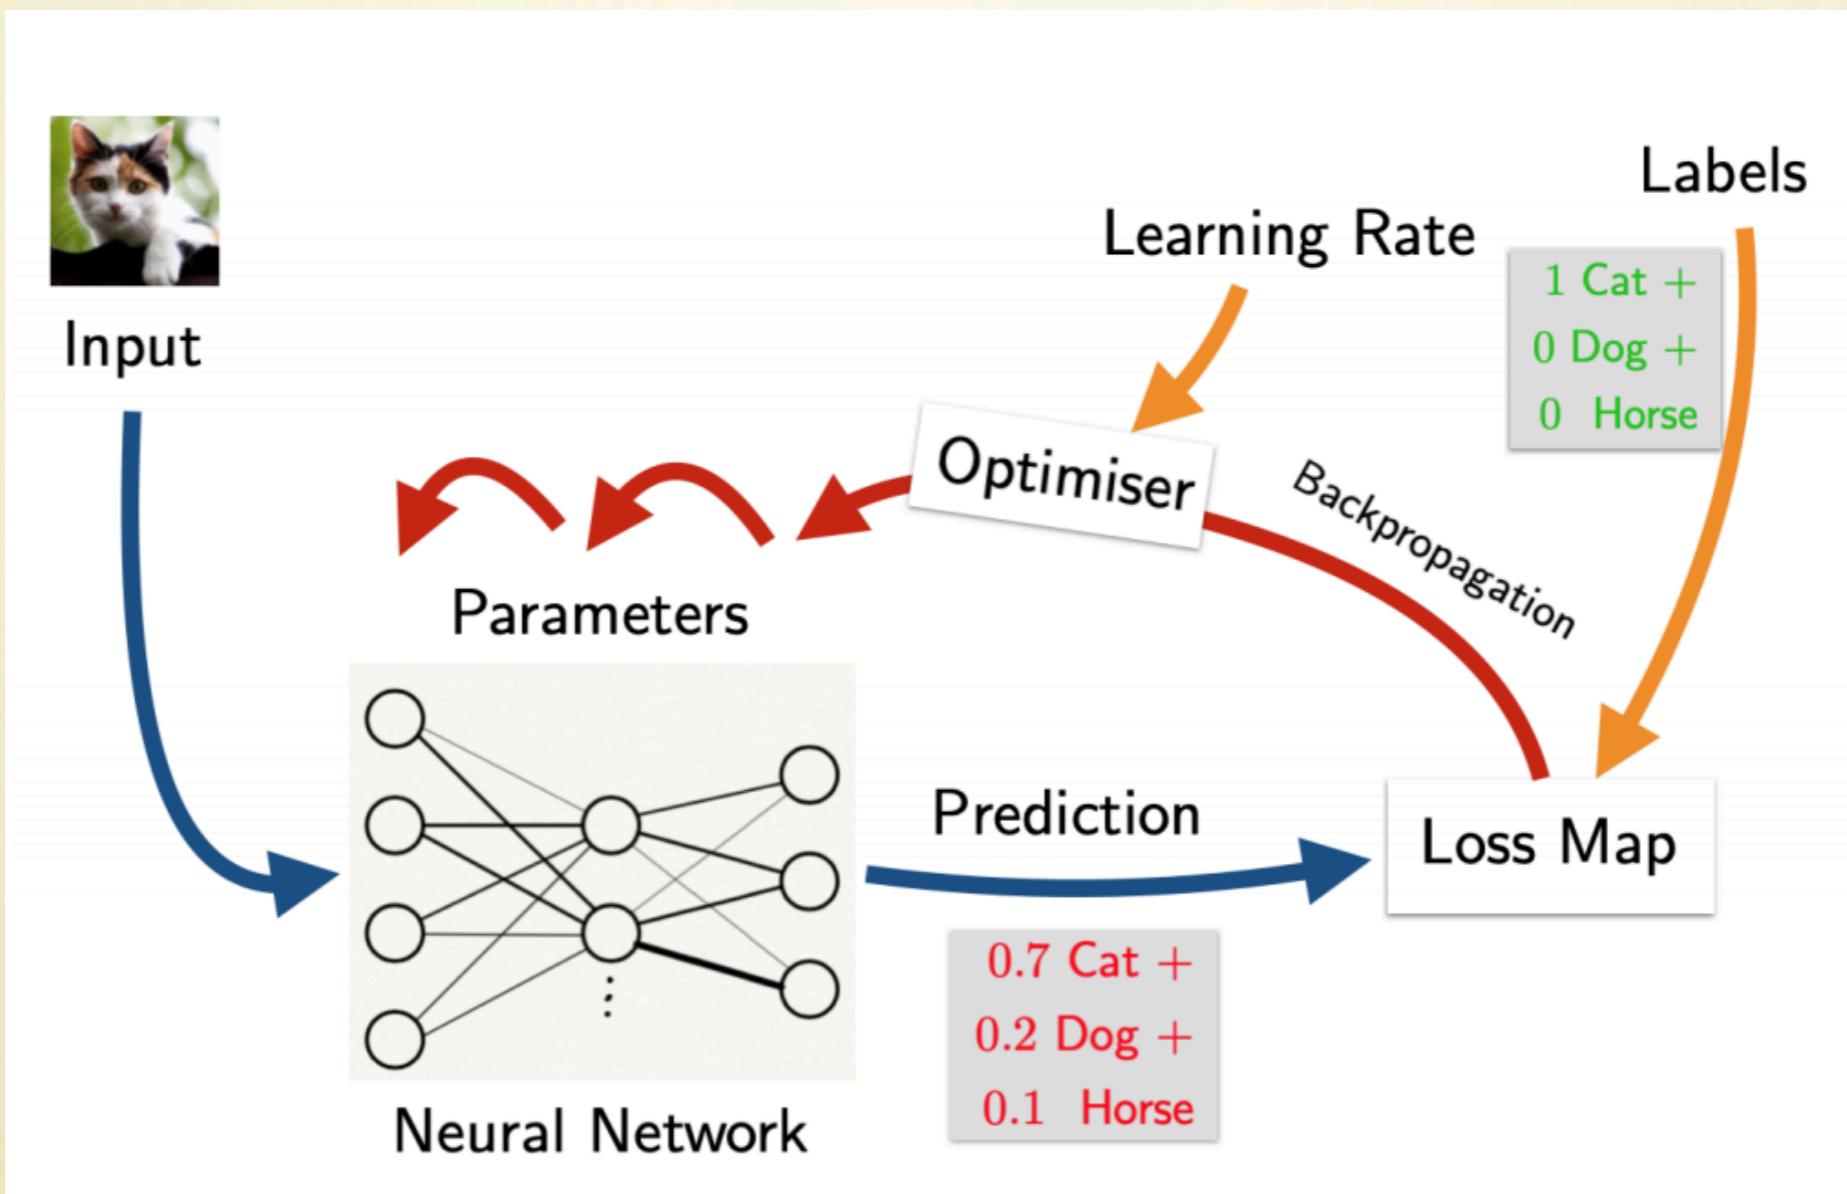
Prog. Principles

Trustworthiness

Case study: Gradient-Based Learning

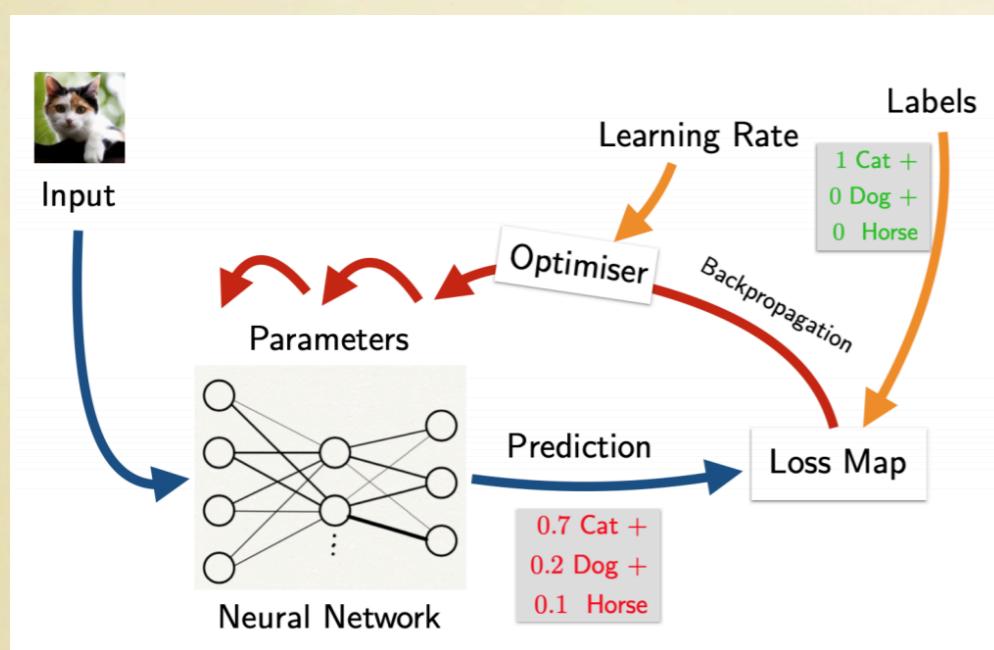


Case study: Gradient-Based Learning

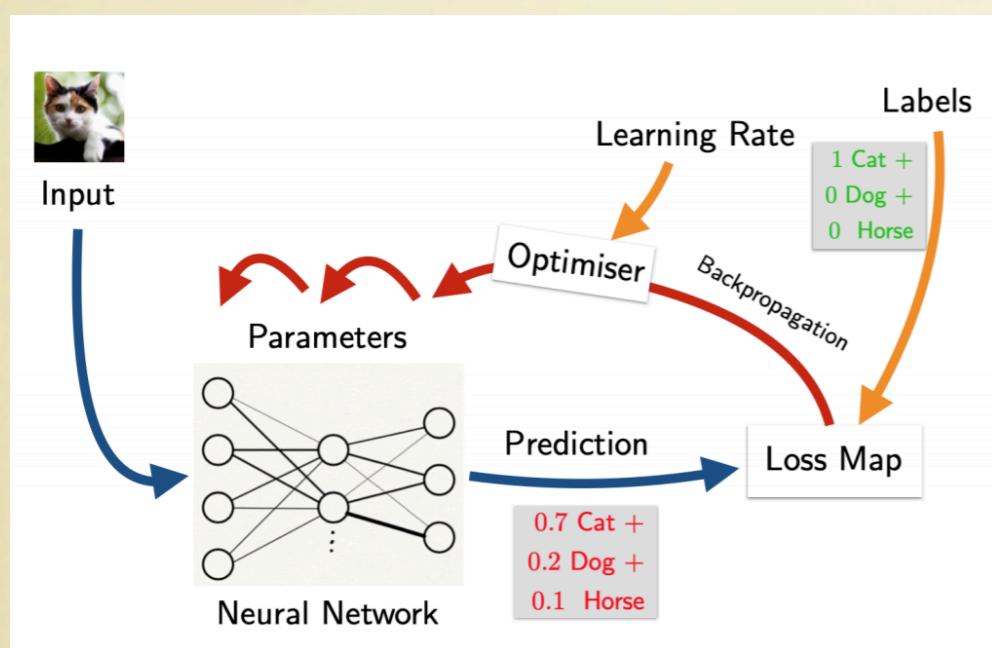


What are the fundamental mathematical structures underpinning gradient-based learning?

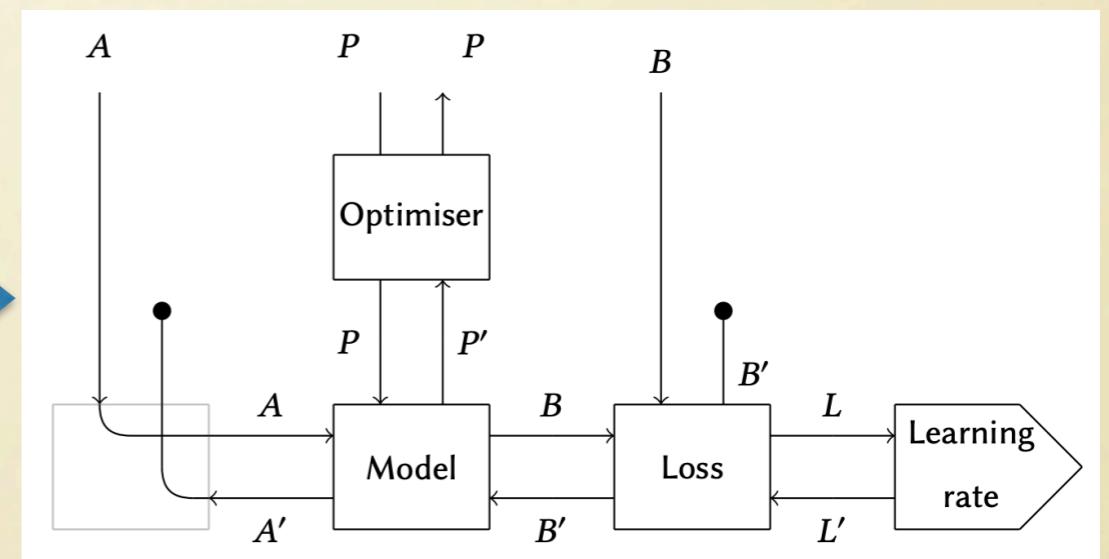
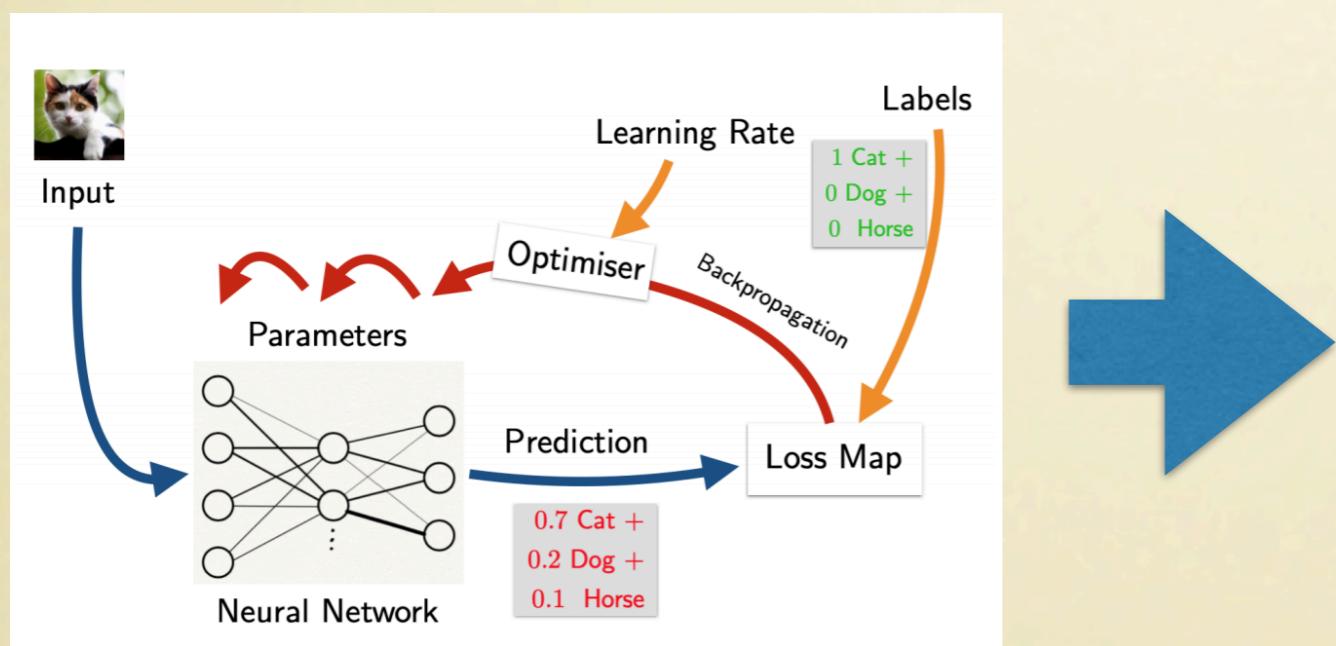
Case study: Gradient-Based Learning



Roadmap



Roadmap



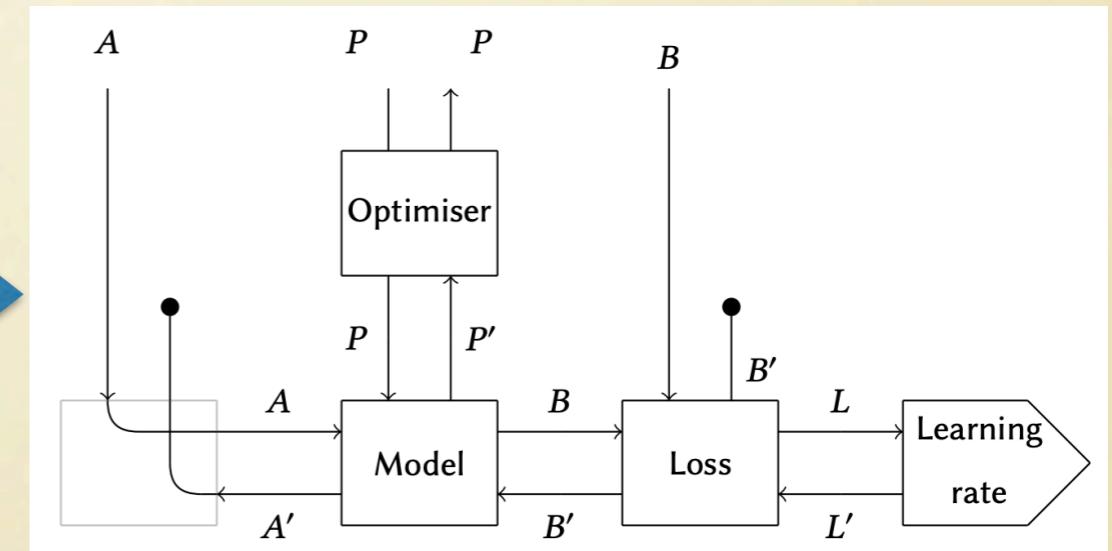
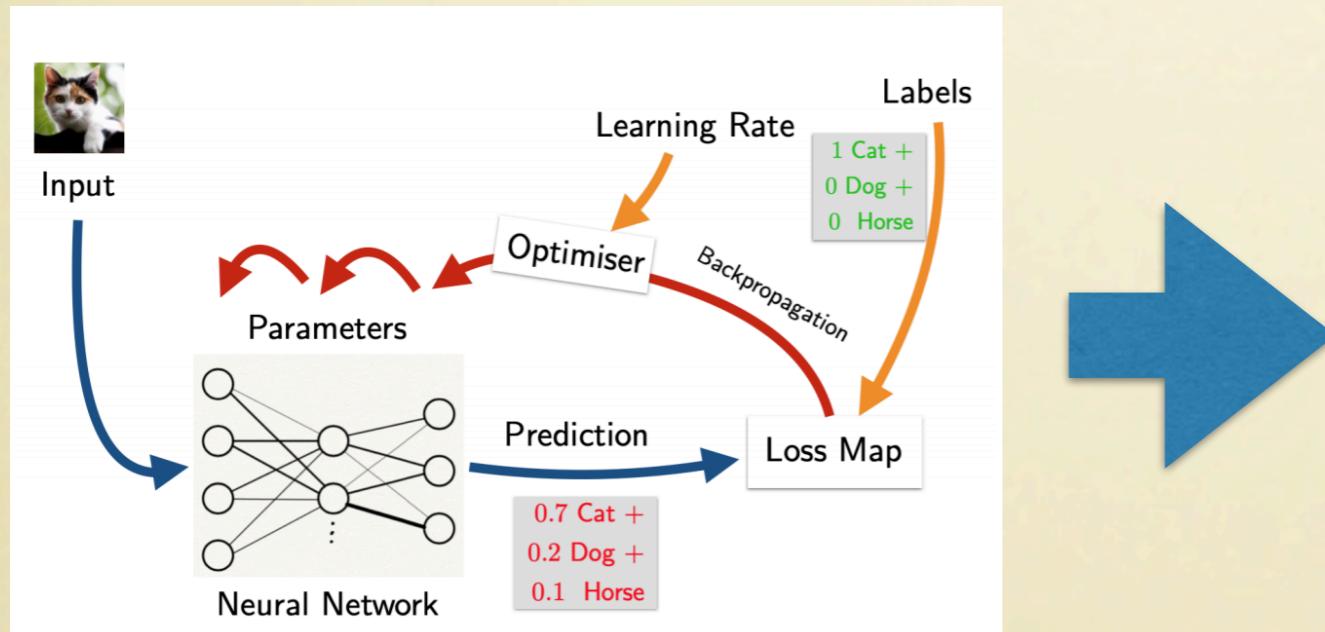
Roadmap

Plan:

Deconstruct learning into simple, uniform components, called **parametric lenses**

(We express parametric lenses as *string diagrams* in a suitable category)

Learning will be defined by ***lens composition***.



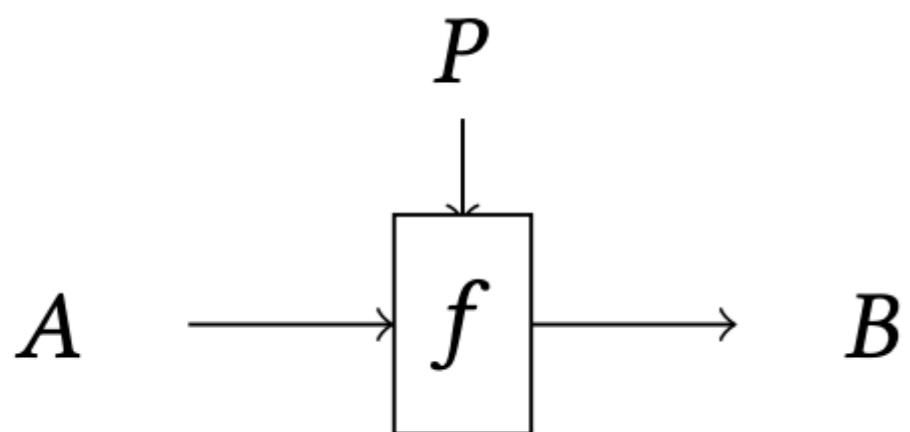
Parametric Lenses

Parametric lens = parametric map with lens structure

Parametric Lenses

Parametric lens = parametric map with lens structure

$$f: P \times A \rightarrow B$$

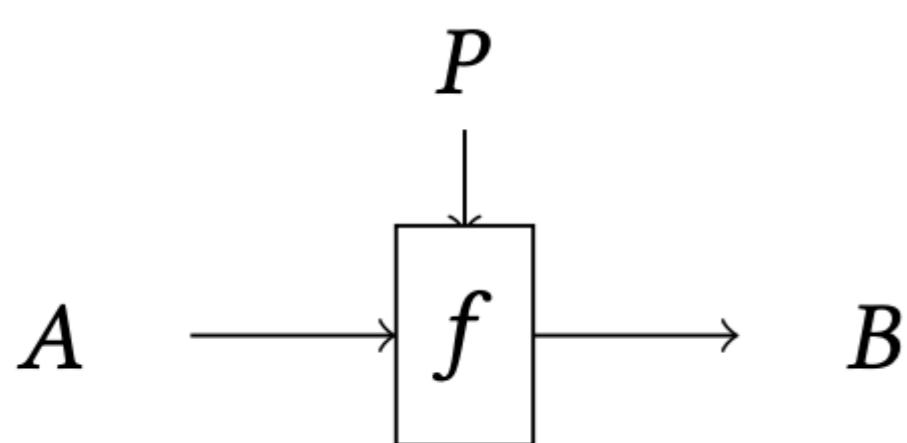


Para(C)

Parametric Lenses

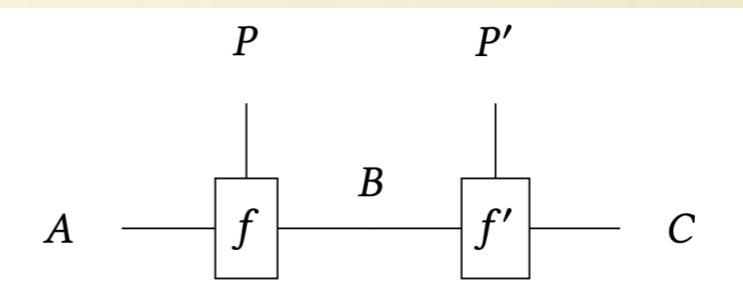
Parametric lens = parametric map with lens structure

$$f: P \times A \rightarrow B$$

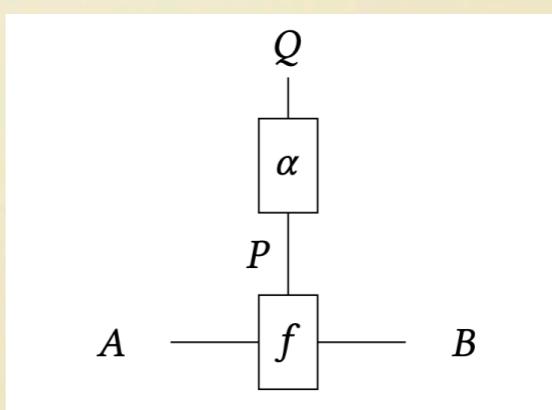


Para(C)

Composition

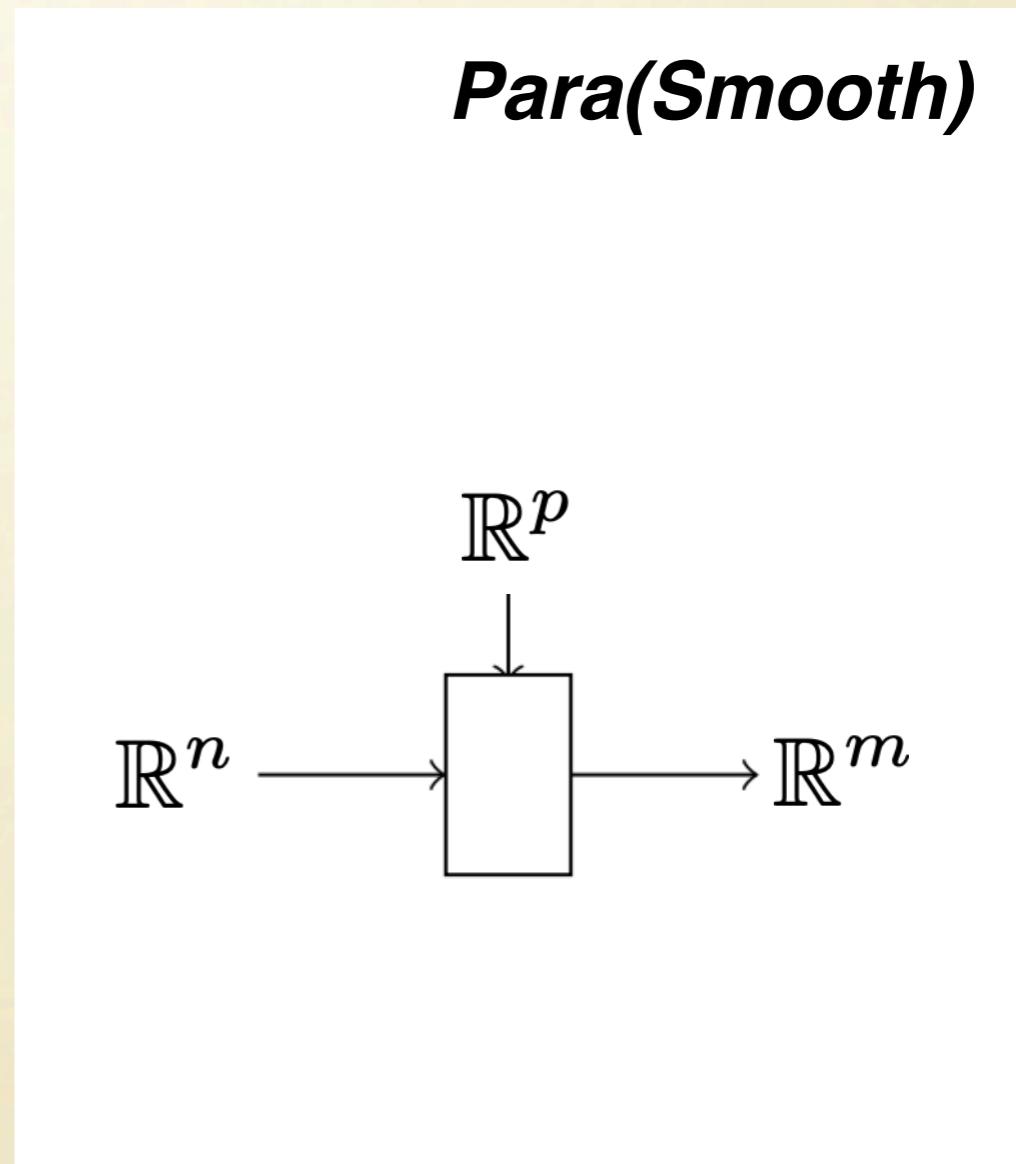
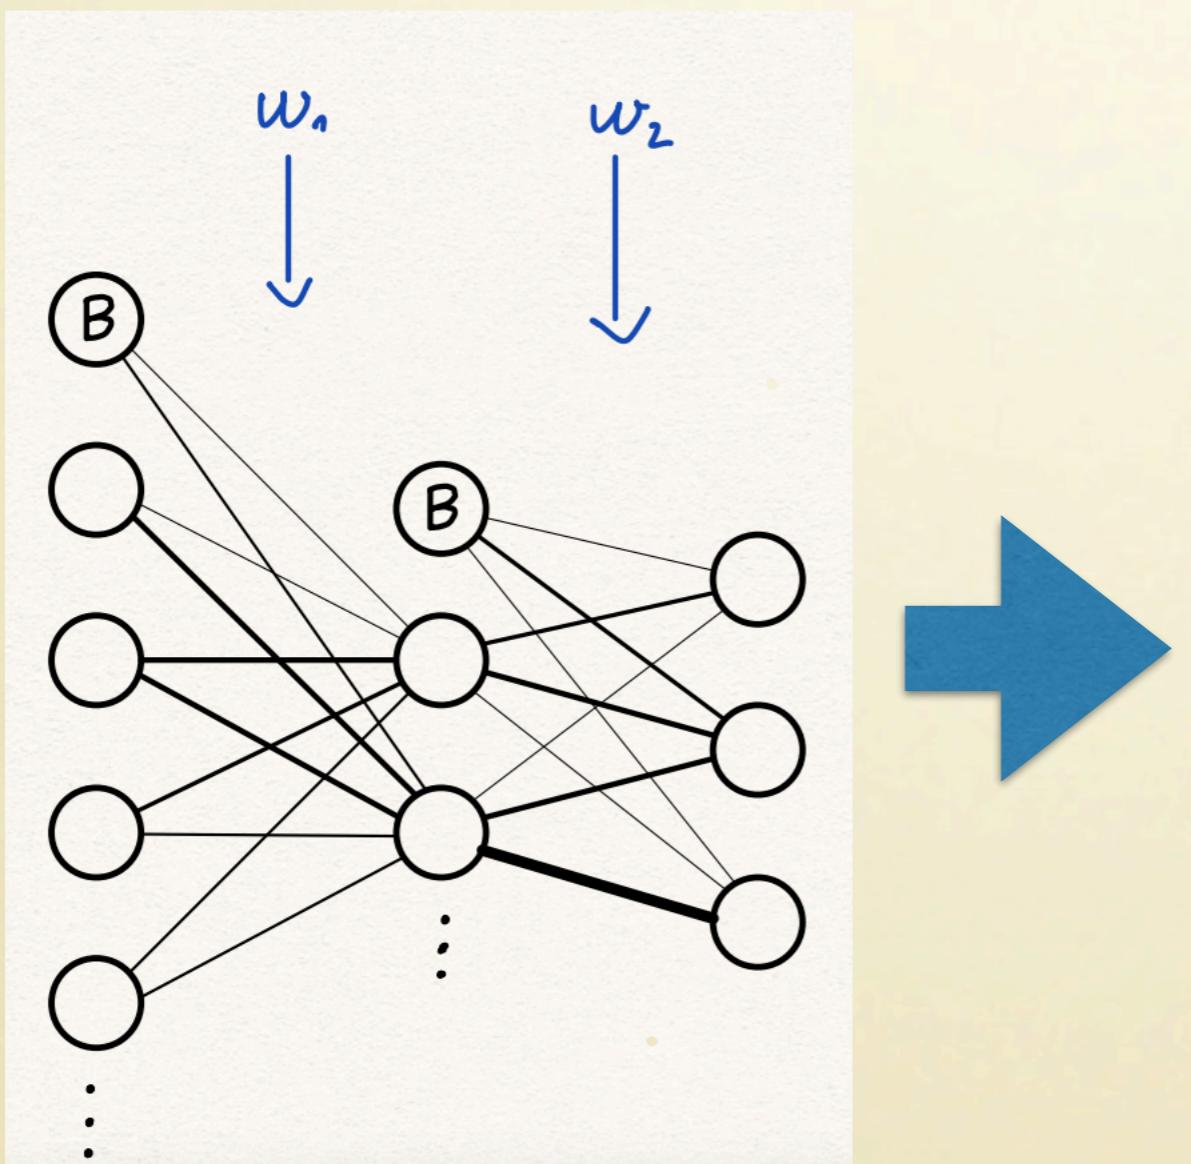


Reparametrisation



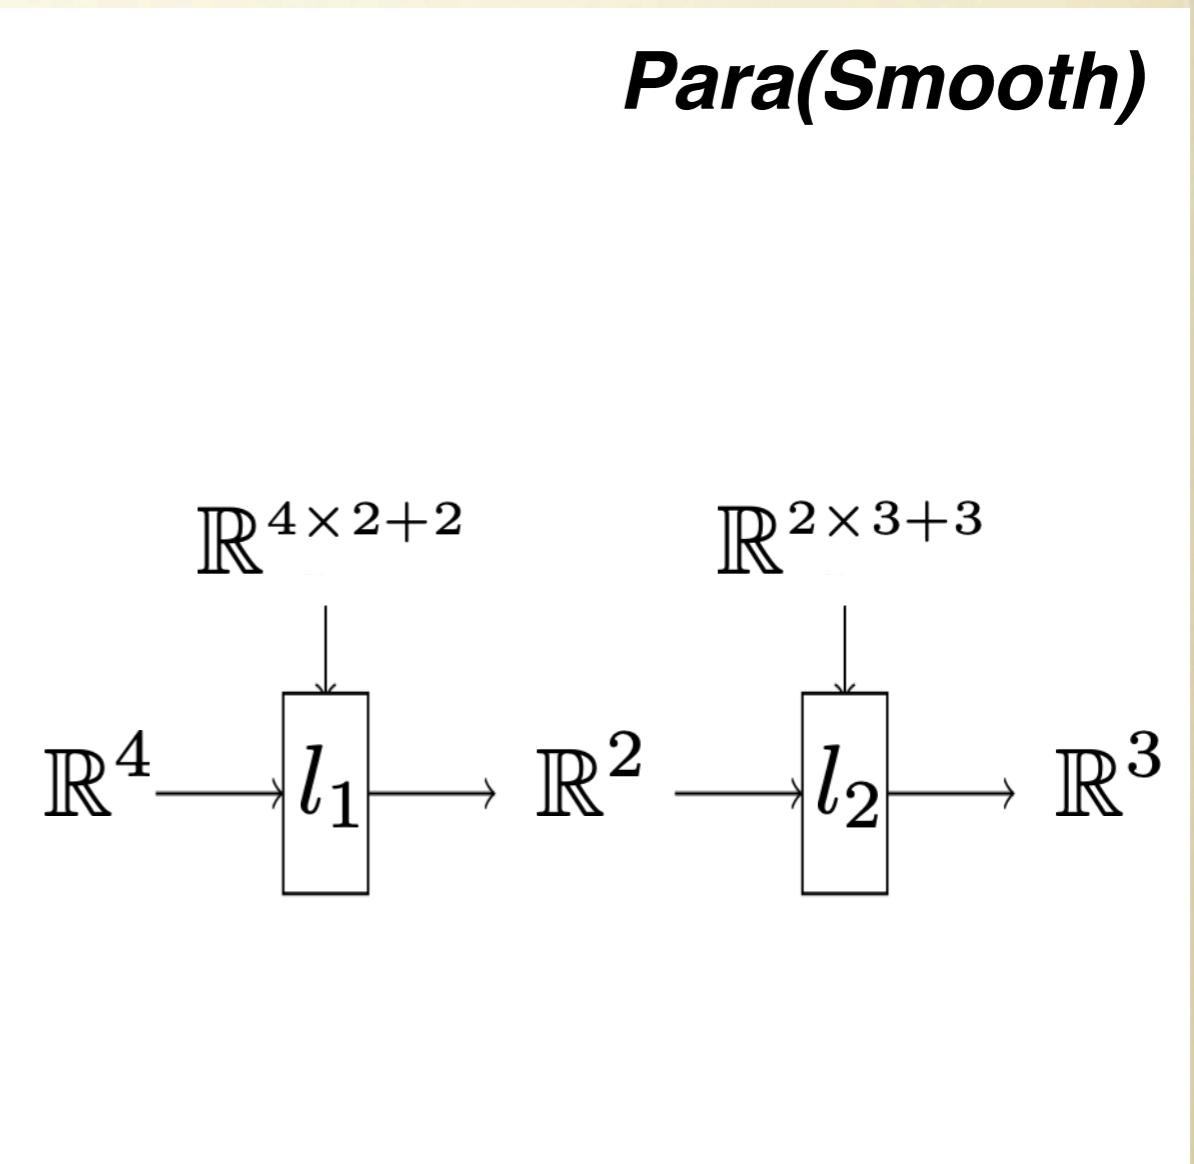
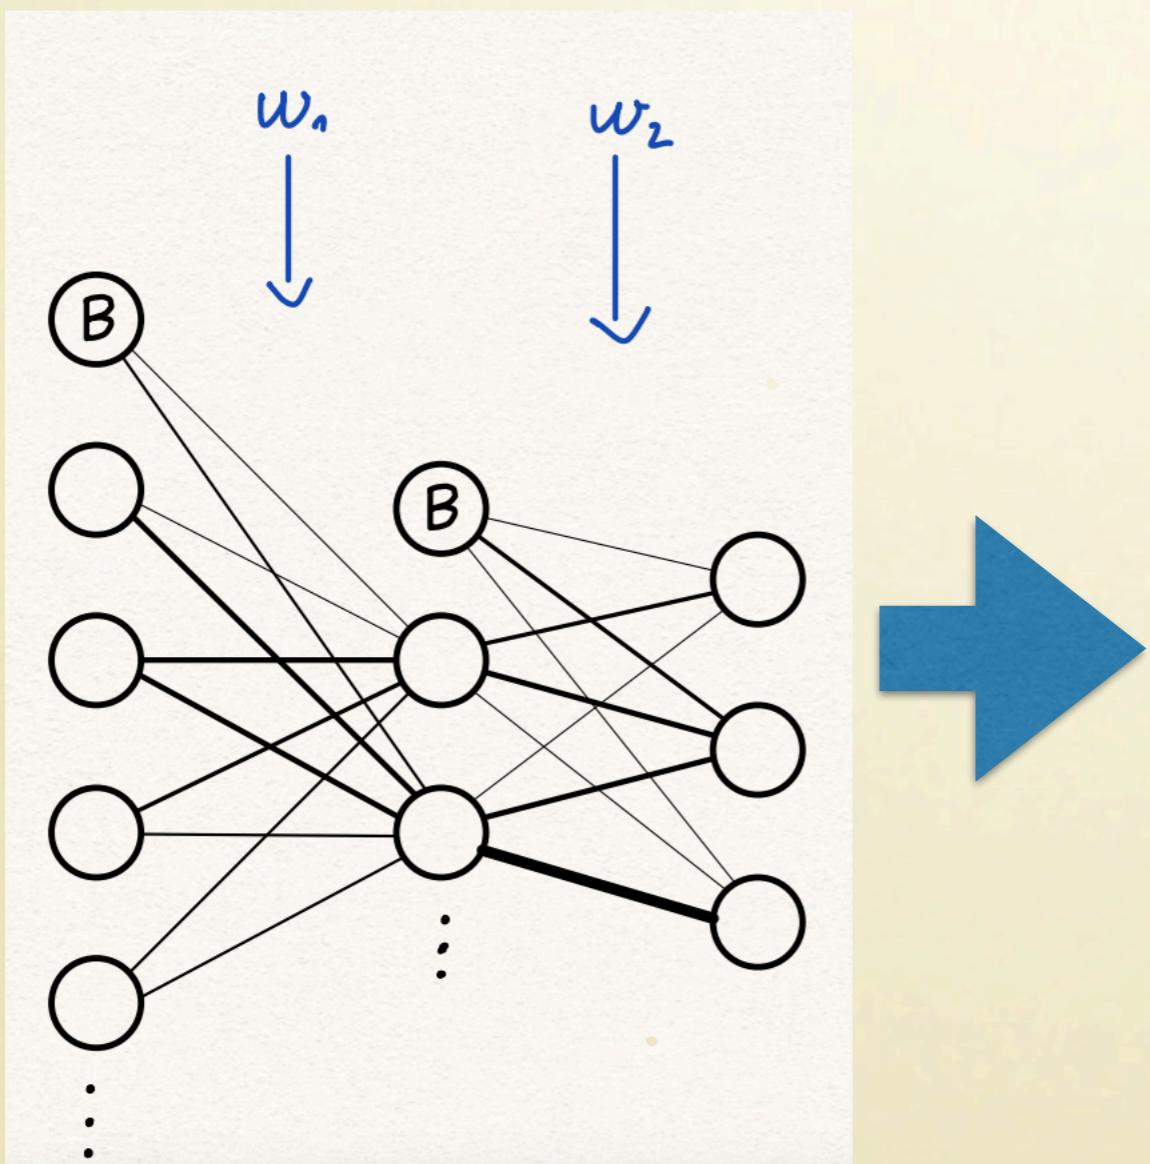
Parametric Maps

A neural network is a parametric map



Parametric Maps

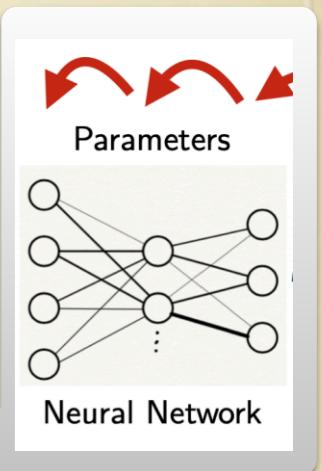
A neural network is a parametric map



Lenses

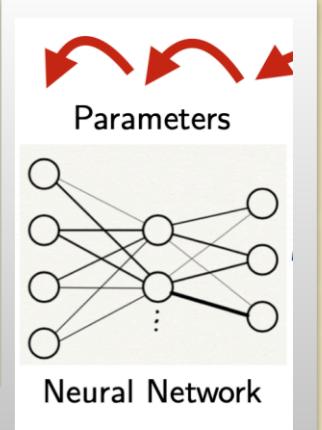
Parametric lens = parametric map with lens structure

Lenses



Parametric lens = parametric map with lens structure

Lenses

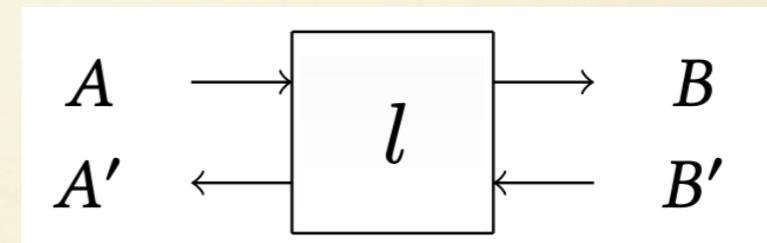


Parametric lens = parametric map with lens structure

$$l: (A, A') \rightarrow (B, B')$$

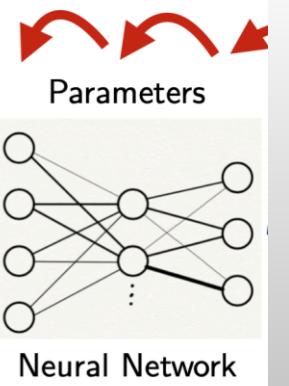
$$\text{get}_l: A \rightarrow B$$

$$\text{put}_l: A \times B' \rightarrow A'$$



Lens(C)

Lenses



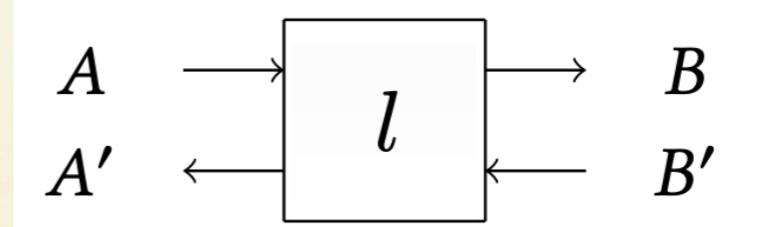
Parametric lens = parametric map with lens structure

$$l: (A, A') \rightarrow (B, B')$$

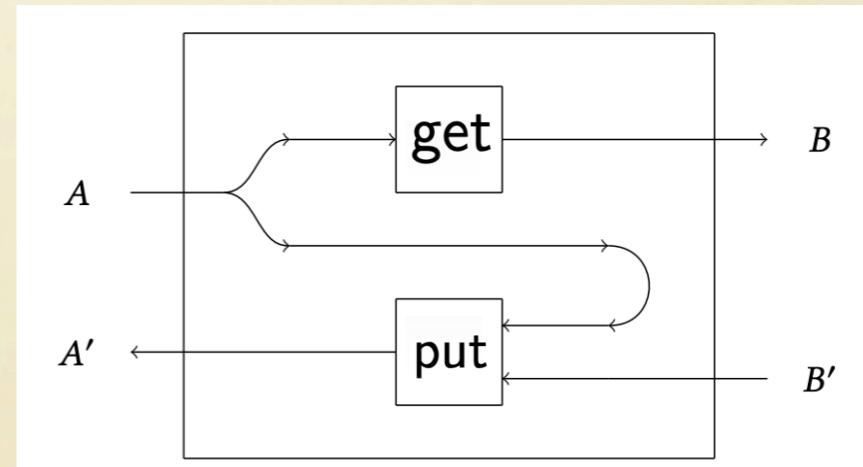
$$\text{get}_l: A \rightarrow B$$

$$\text{put}_l: A \times B' \rightarrow A'$$

Lens(C)



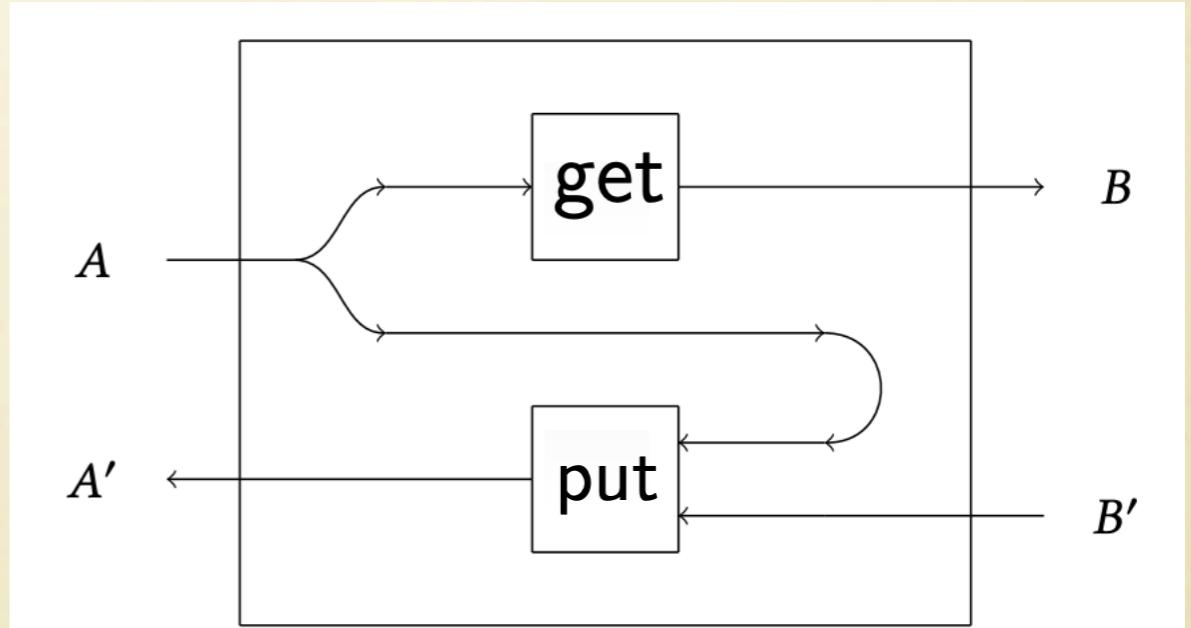
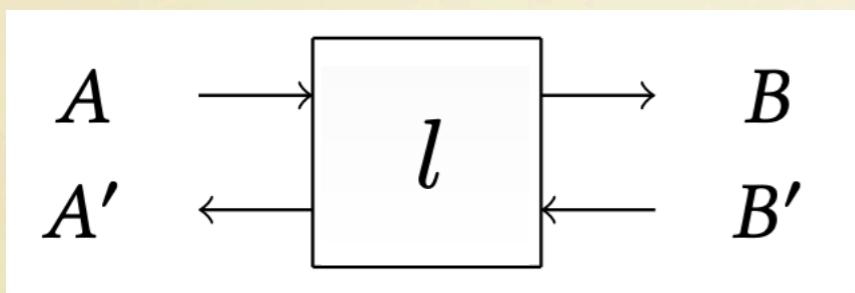
‘Opening the box’:



String Diagrams for Lenses

We use string diagrams in $\text{Tamb}(\mathcal{C})$

Faithful embedding of $\text{Lens}(\mathcal{C})^l$ in $[\text{Lens}(\mathcal{C})^{op}, \text{Set}] \cong \text{Tamb}(\mathcal{C})$



Lenses

A neural network is a lens in **Smooth**

$$l: (A, A') \rightarrow (B, B')$$

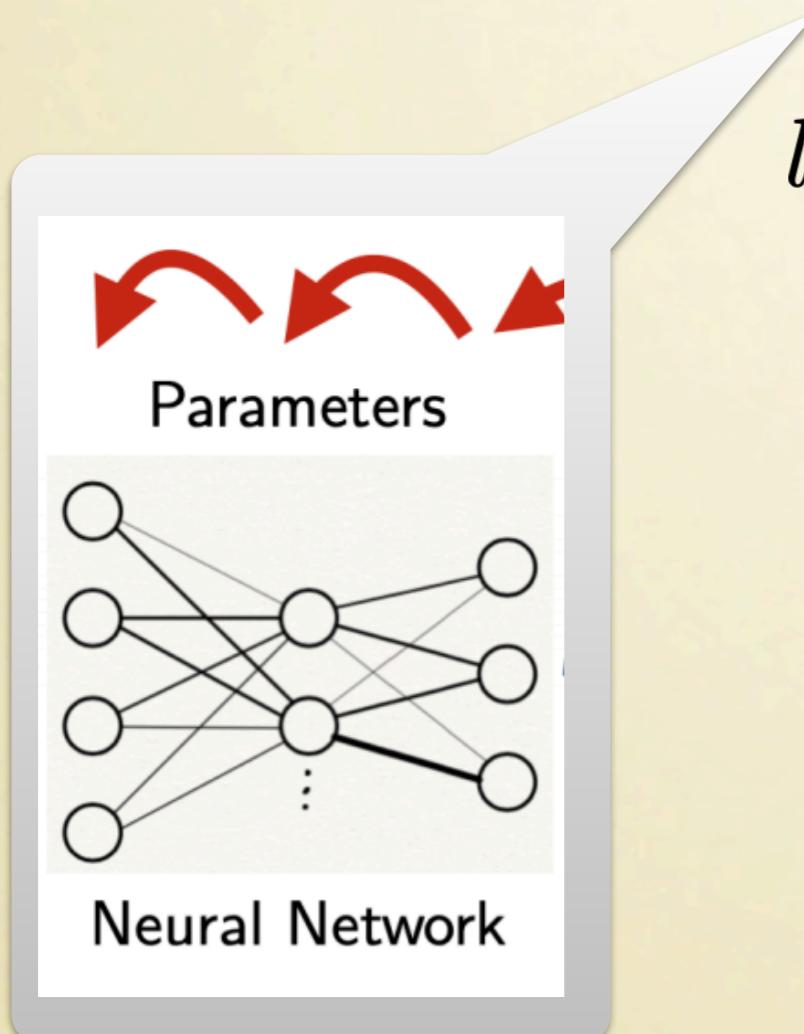
$$\overline{f: A \rightarrow B}$$

$$R[f]: A \times B' \rightarrow A'$$

$$(\vec{x}, \vec{y}) \mapsto J_f(\vec{x})^T \cdot \vec{y}$$

Lenses

A neural network is a lens in **Smooth**



$$l: (A, A') \rightarrow (B, B')$$

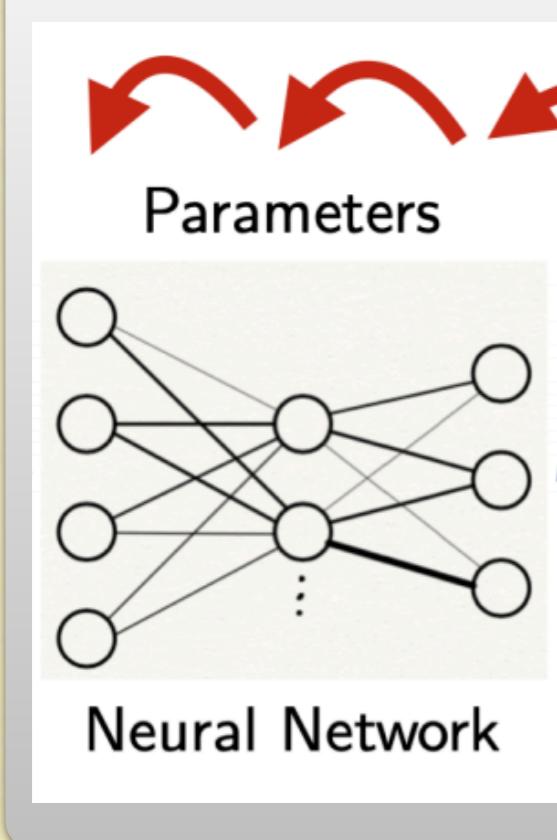
$$\overline{f: A \rightarrow B}$$

$$R[f]: A \times B' \rightarrow A'$$

$$(\vec{x}, \vec{y}) \mapsto J_f(\vec{x})^T \cdot \vec{y}$$

Lenses

A neural network is a lens in **Smooth**



$$l: (A, A') \rightarrow (B, B')$$

$$\overline{f: A \rightarrow B}$$
$$R[f]: A \times B' \rightarrow A'$$

$$(\vec{x}, \vec{y}) \mapsto J_f(\vec{x})^T \cdot \vec{y}$$

Reverse
derivative
of f

Reverse derivative categories *

Robin Cockett
University of Calgary, Department of Computer Science, Canada
robin@ucalgary.ca

Geoffrey Cruttwell
Mount Allison University, Department of Mathematics and Computer Science, Canada
gcruttwell@mta.ca

Jonathan Gallagher
Dalhousie University, Department of Mathematics and Statistics, Canada
jonathan.gallagher@dal.ca

Jean-Simon Pacaud Lemay
University of Oxford, Department of Computer Science, UK
jean-simon.lemay@kellogg.ox.ac.uk

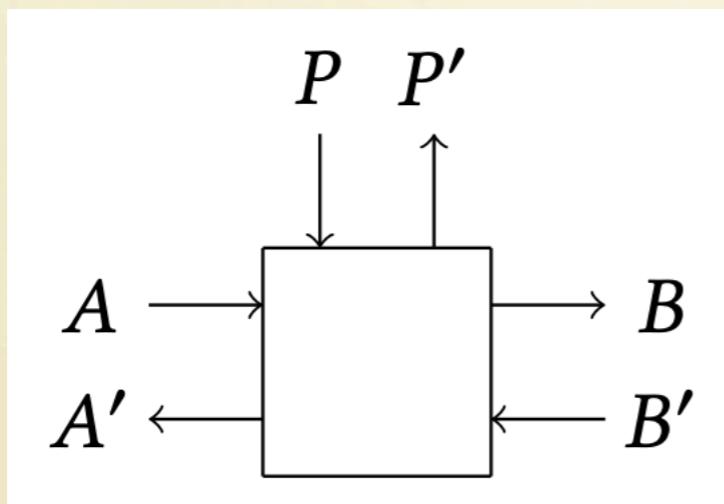
Benjamin MacAdam
University of Calgary, Department of Computer Science, Canada
benjamin.macadam@ucalgary.ca

Gordon Plotkin
Google Research
gdp@inf.ed.ac.uk

Dorette Pronk
Dalhousie University, Department of Mathematics and Statistics, Canada
dorette.pronk@dal.ca

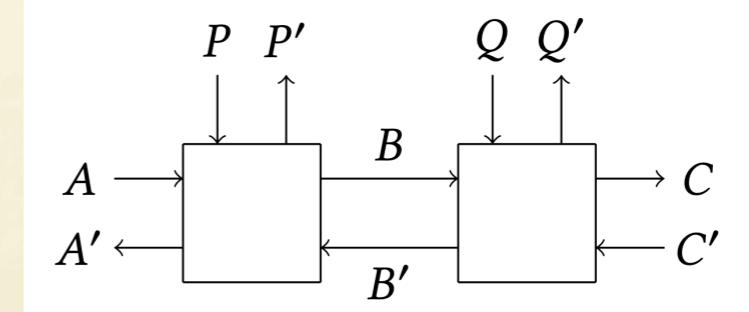
Parametric Lenses

Parametric lens = parametric map with lens structure

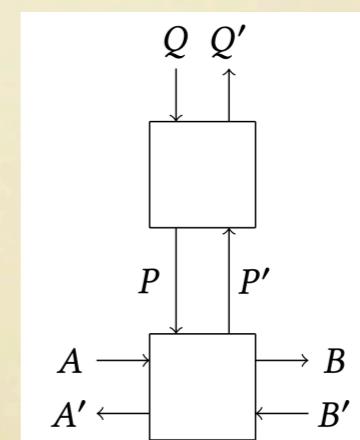


Para(Lens(C))

Composition

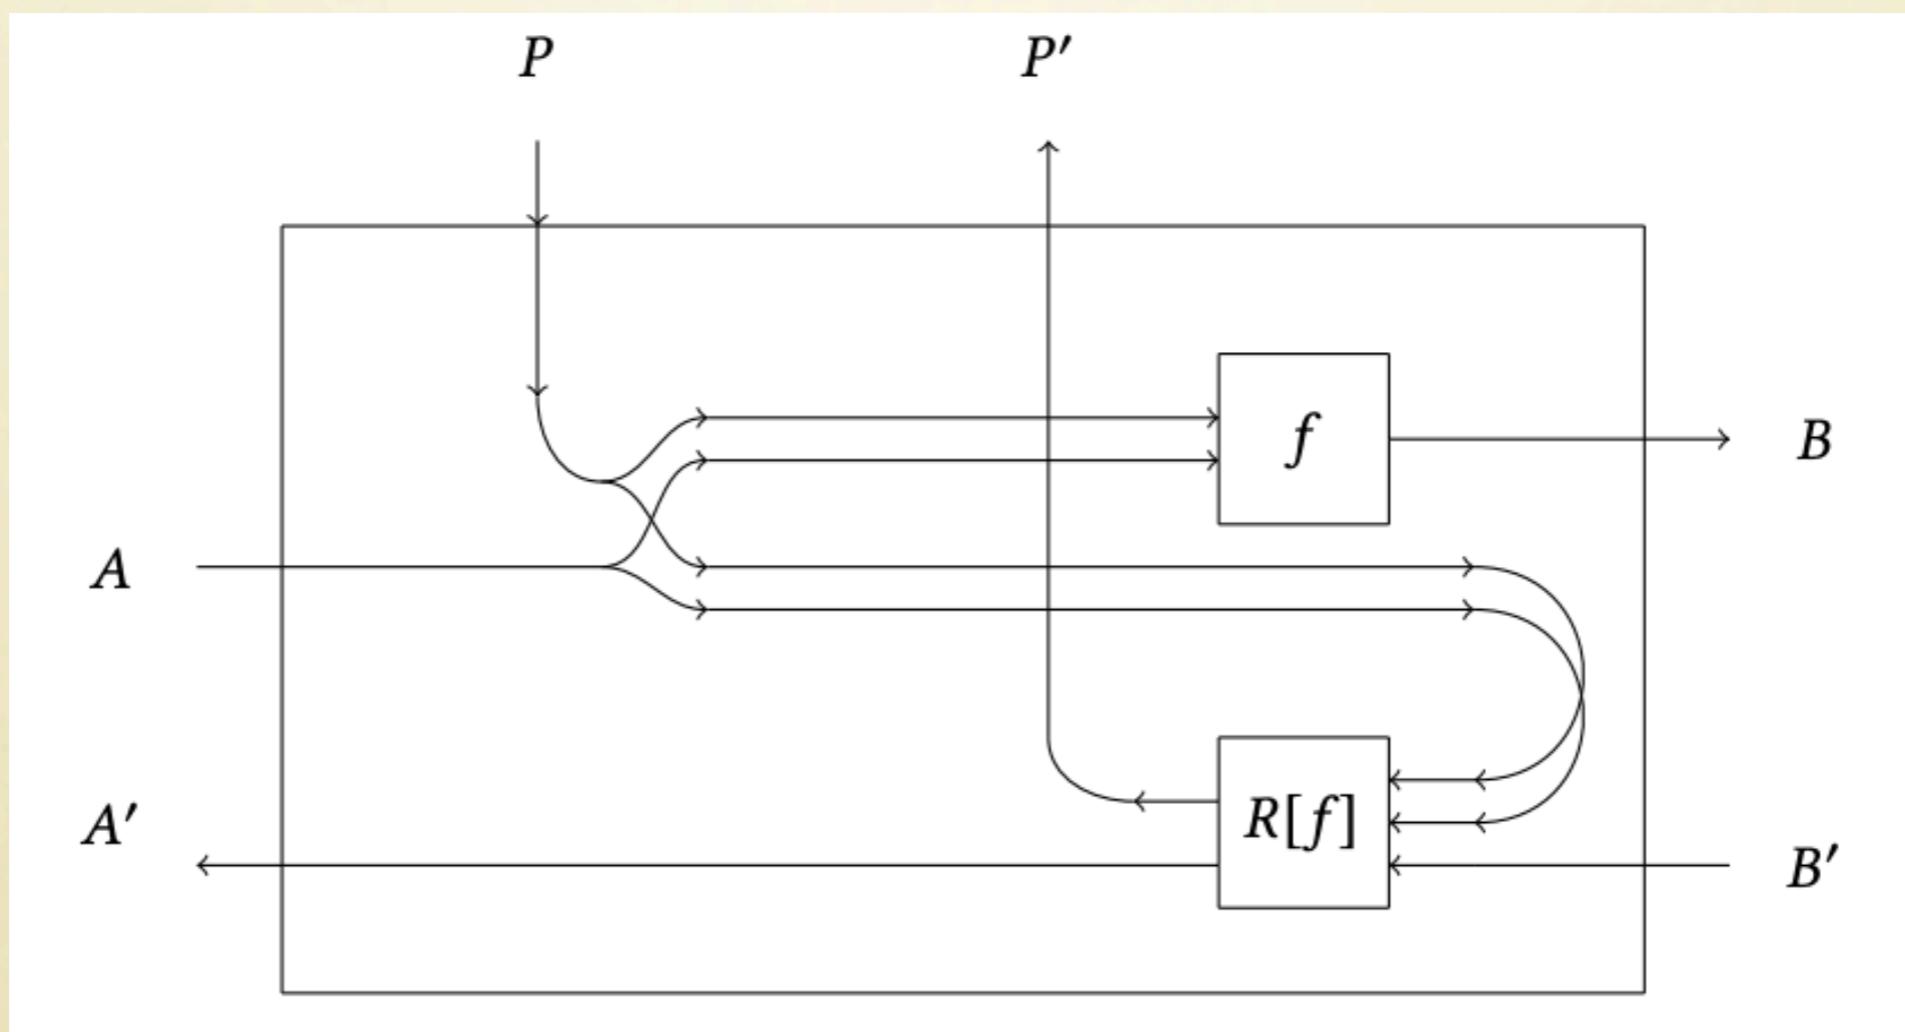


Reparametrisation



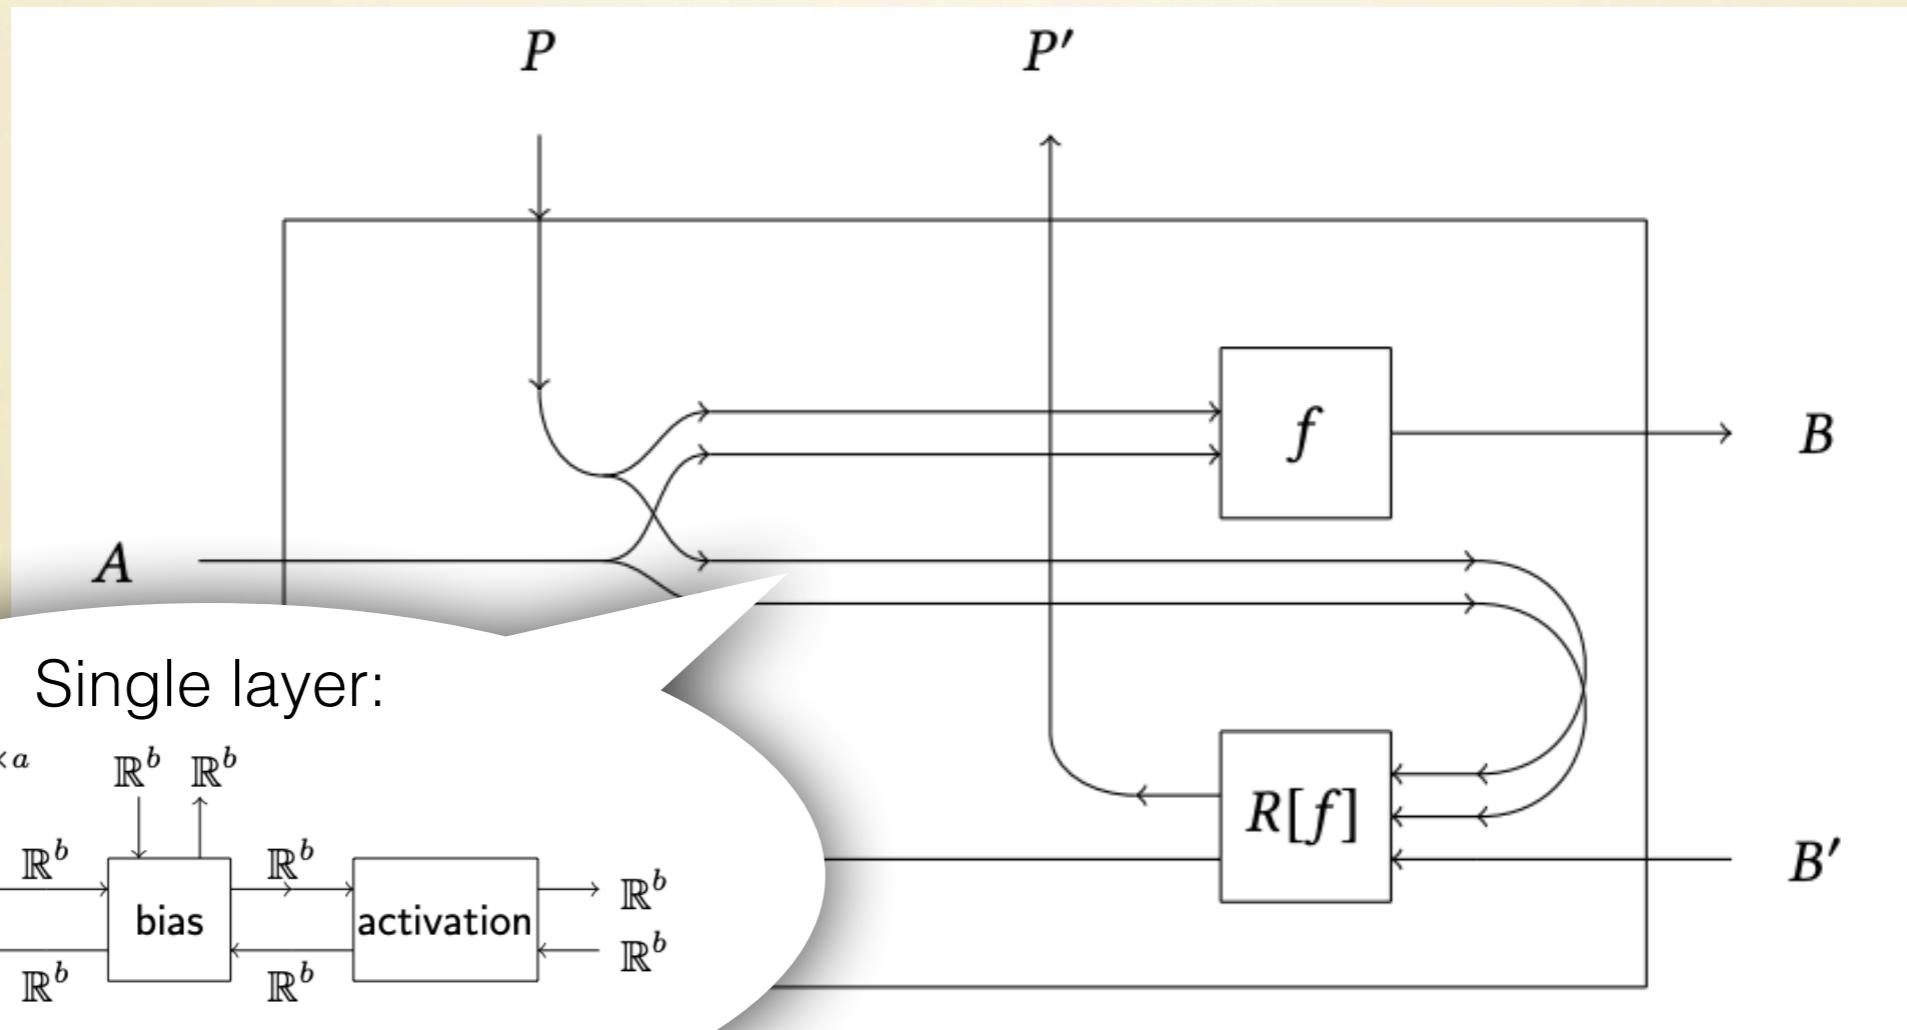
Parametric Lenses

A neural network is a parametric lens

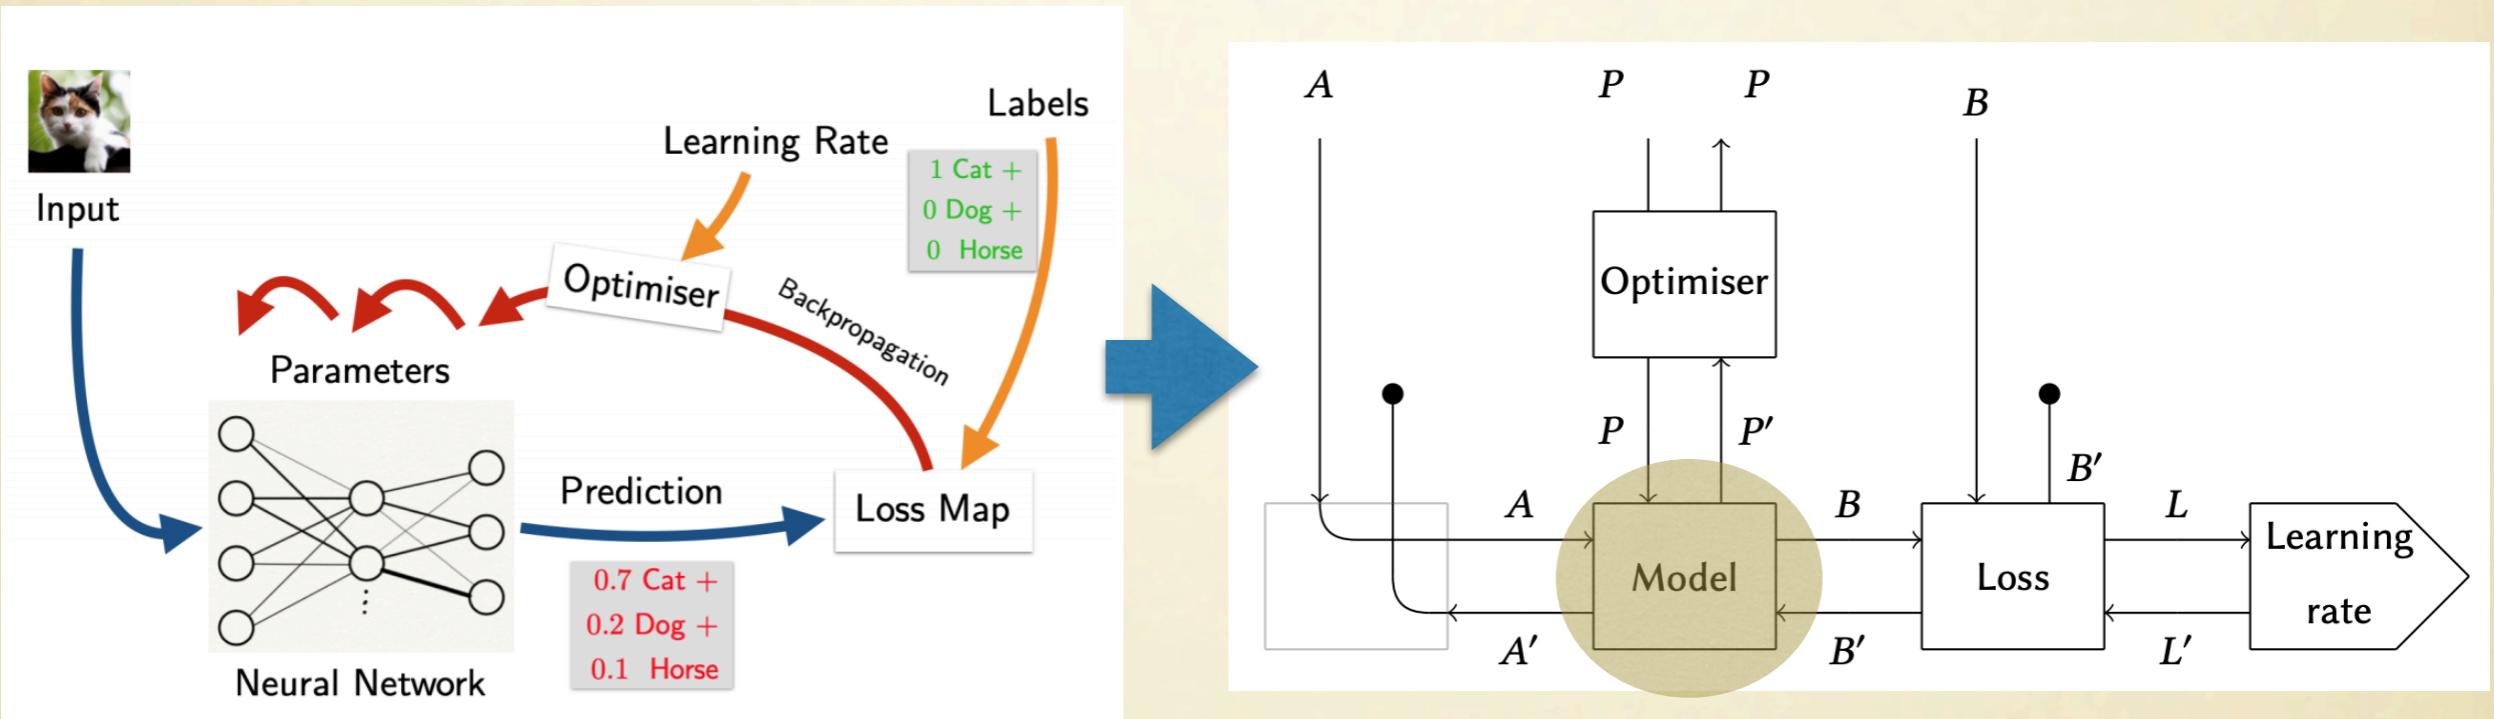


Parametric Lenses

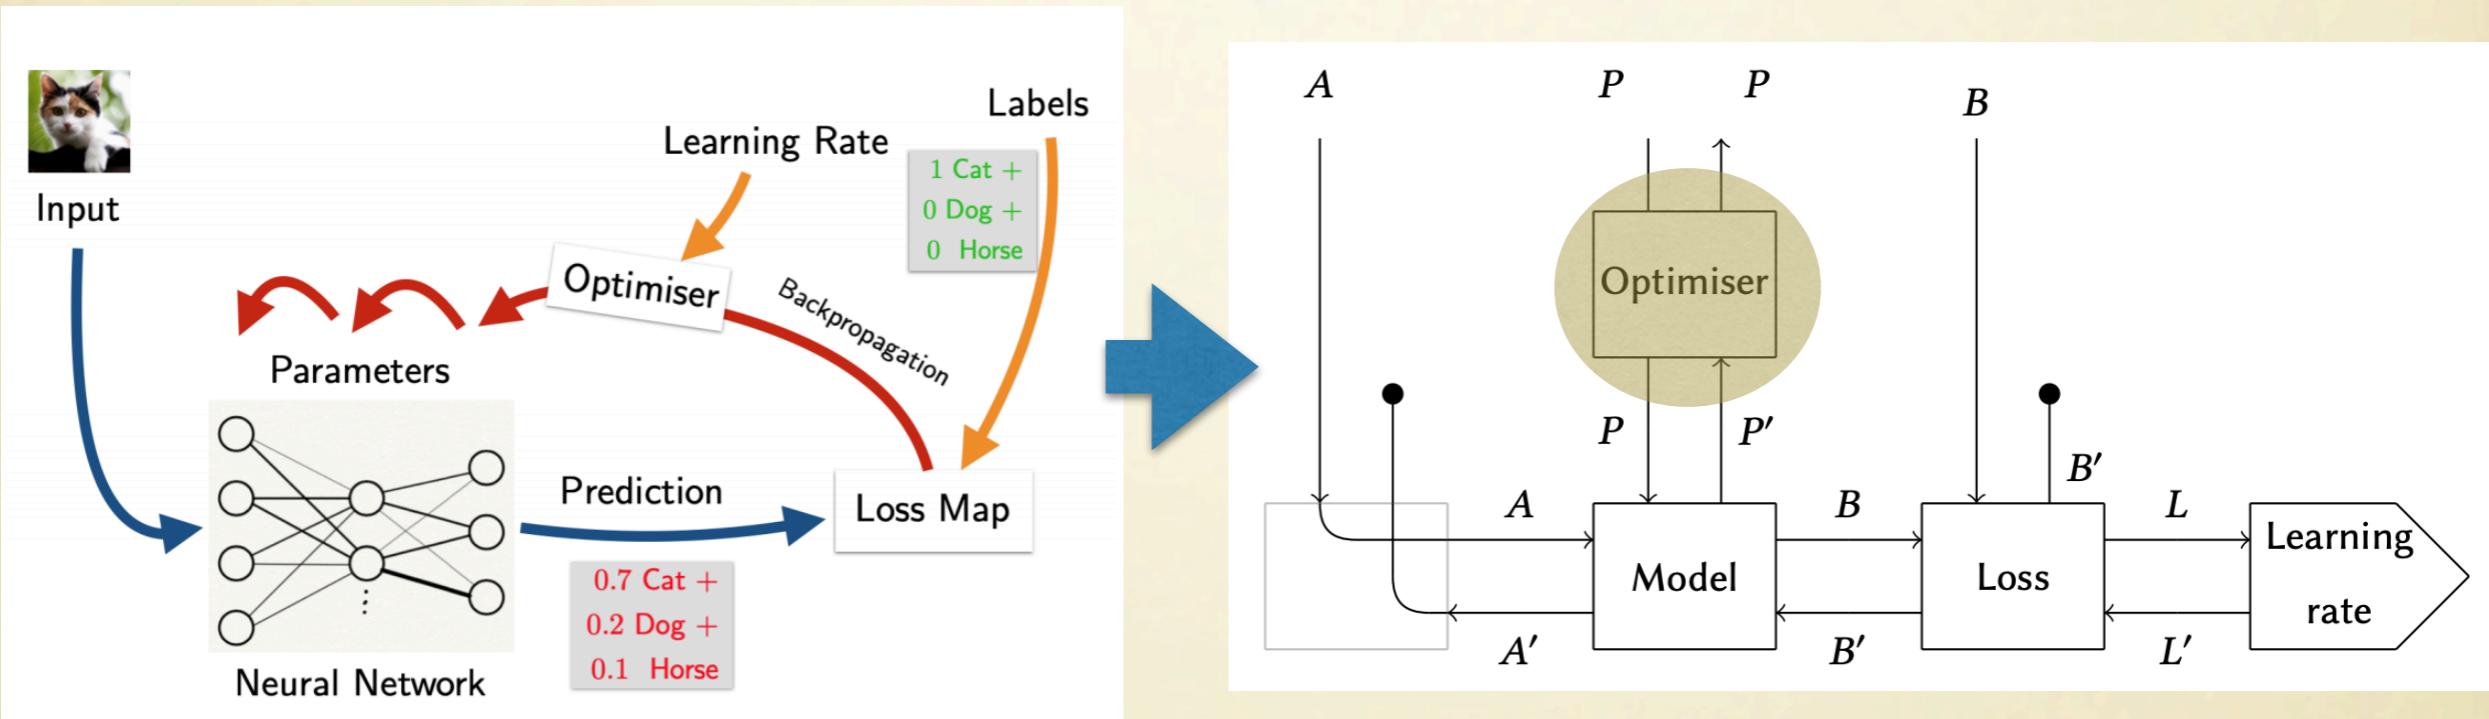
A neural network is a parametric lens



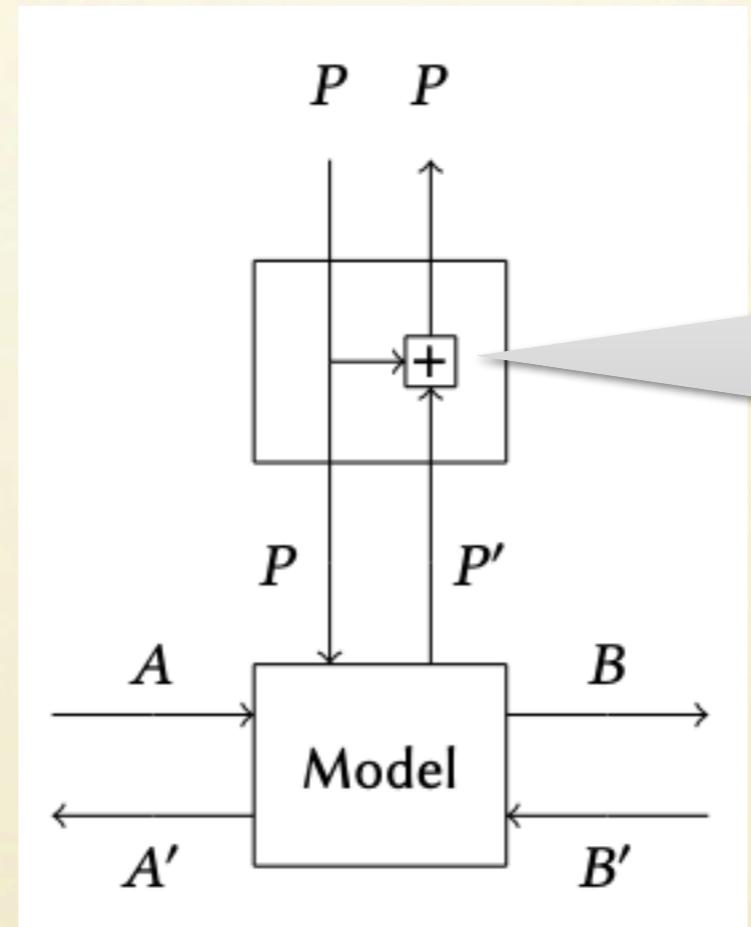
Where we are, so far



Where we are, so far



Optimisers are reparametrisations



Lens(Smooth)

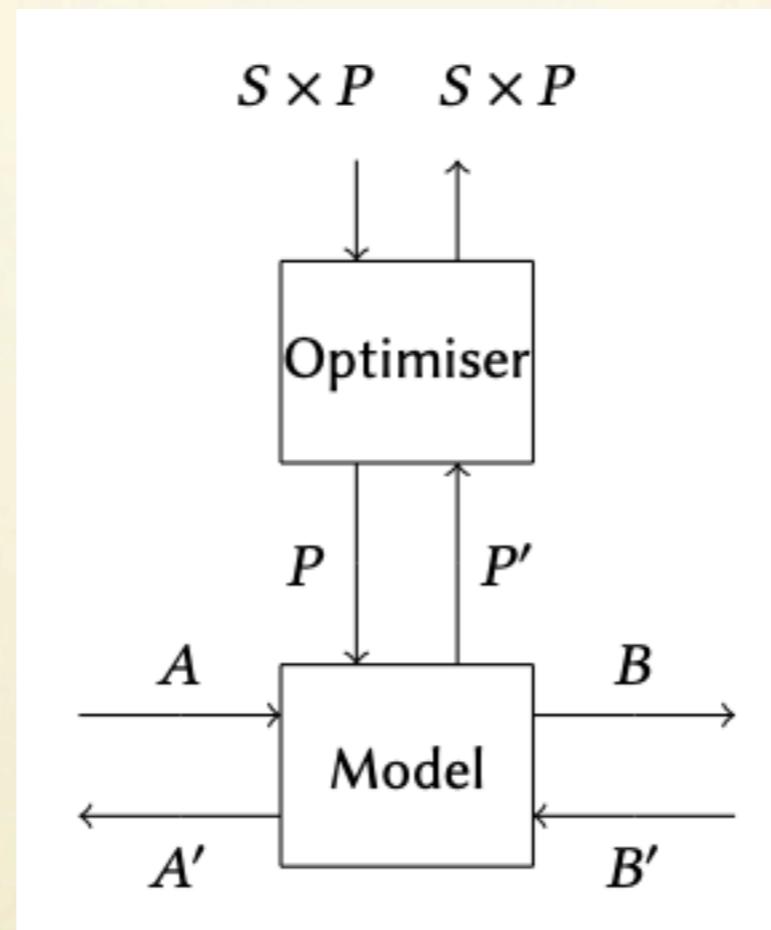
$$\text{gd}: (P, P) \rightarrow (P, P')$$

$$\text{get}_{\text{gd}}: p \mapsto p$$

$$\text{put}_{\text{gd}}: (p, p') \mapsto p + p'$$

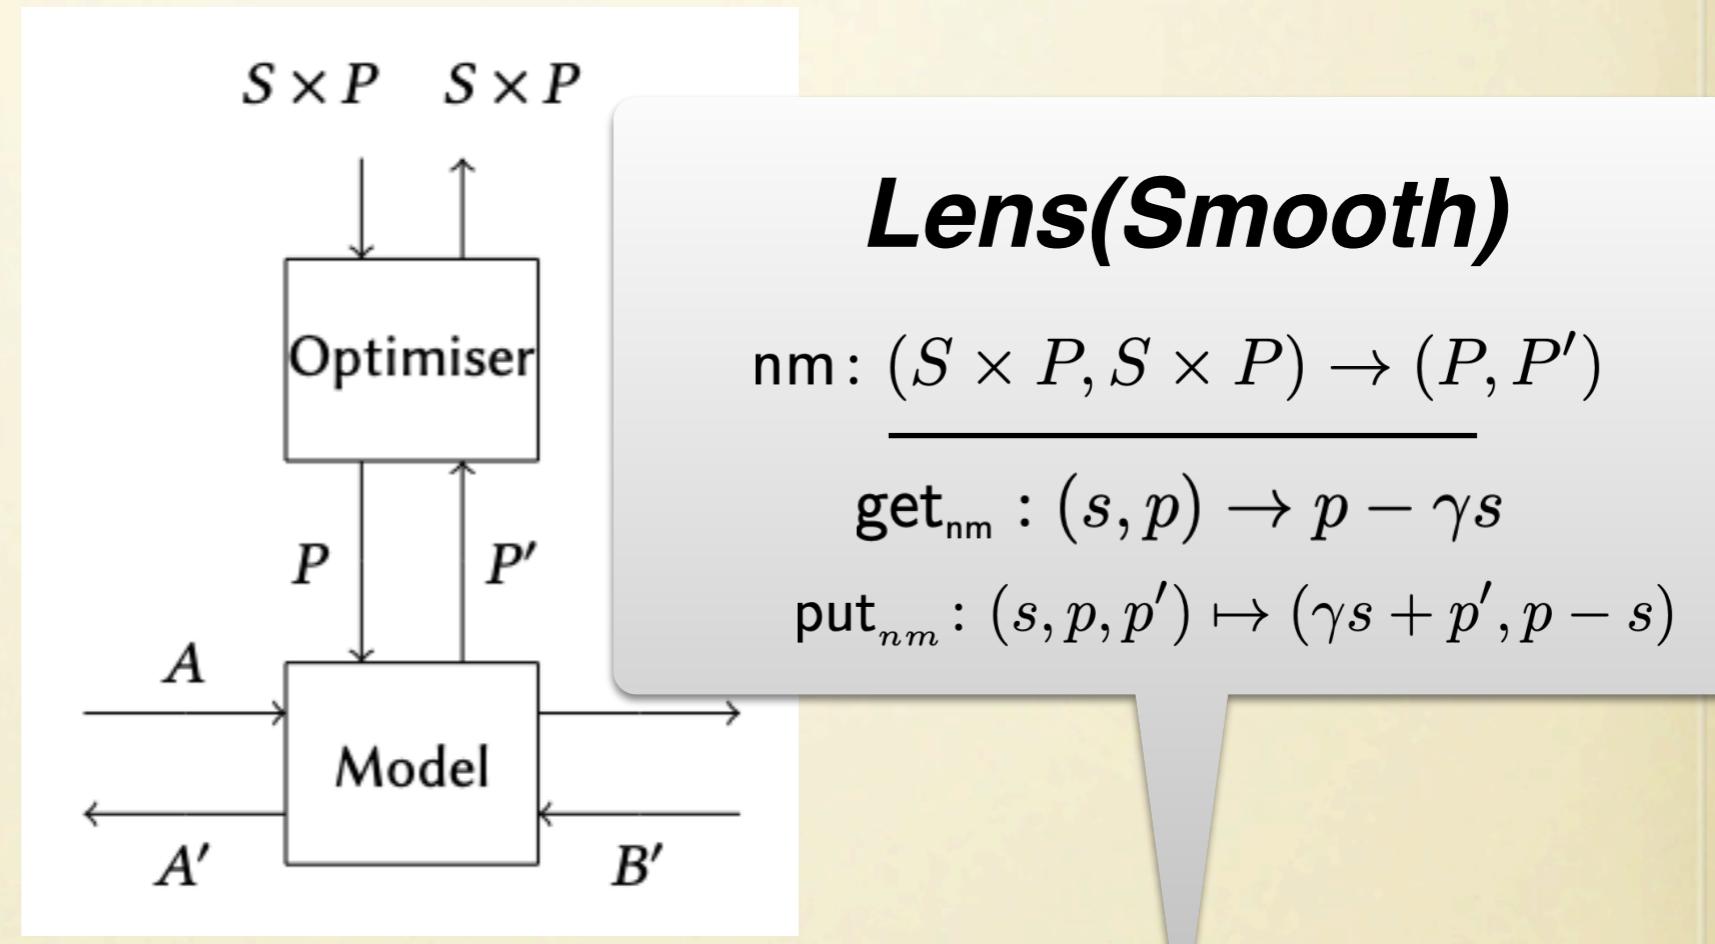
Basic gradient descent

Optimisers are reparametrisations



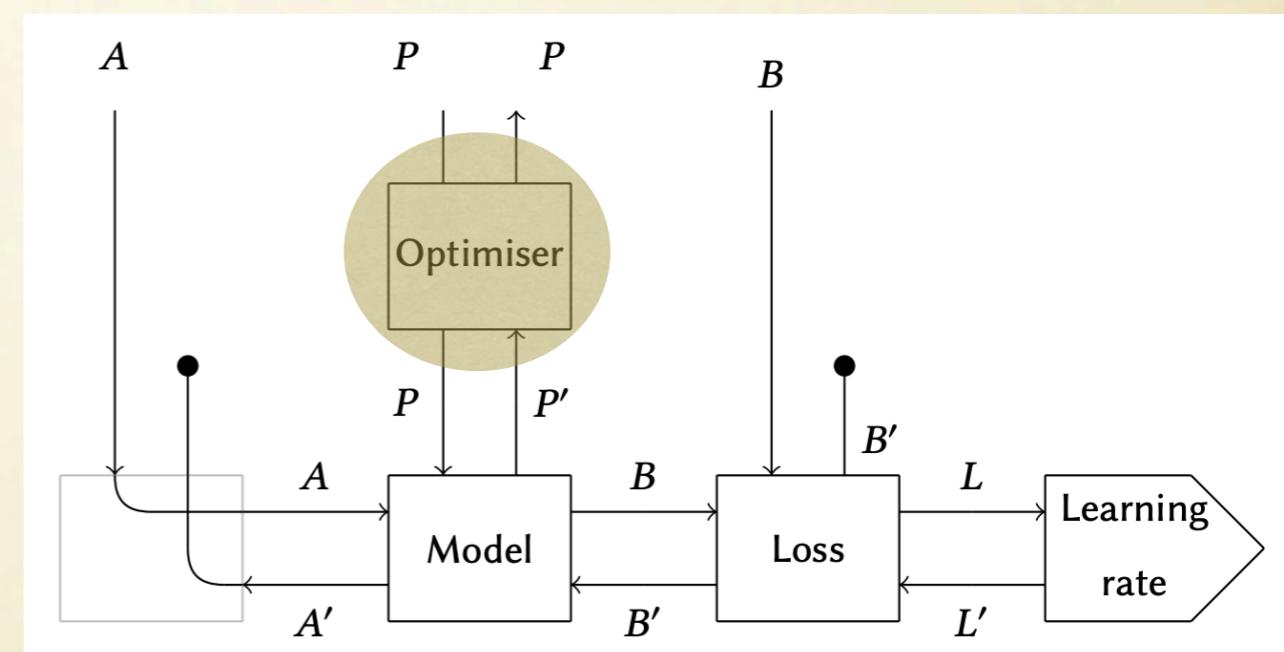
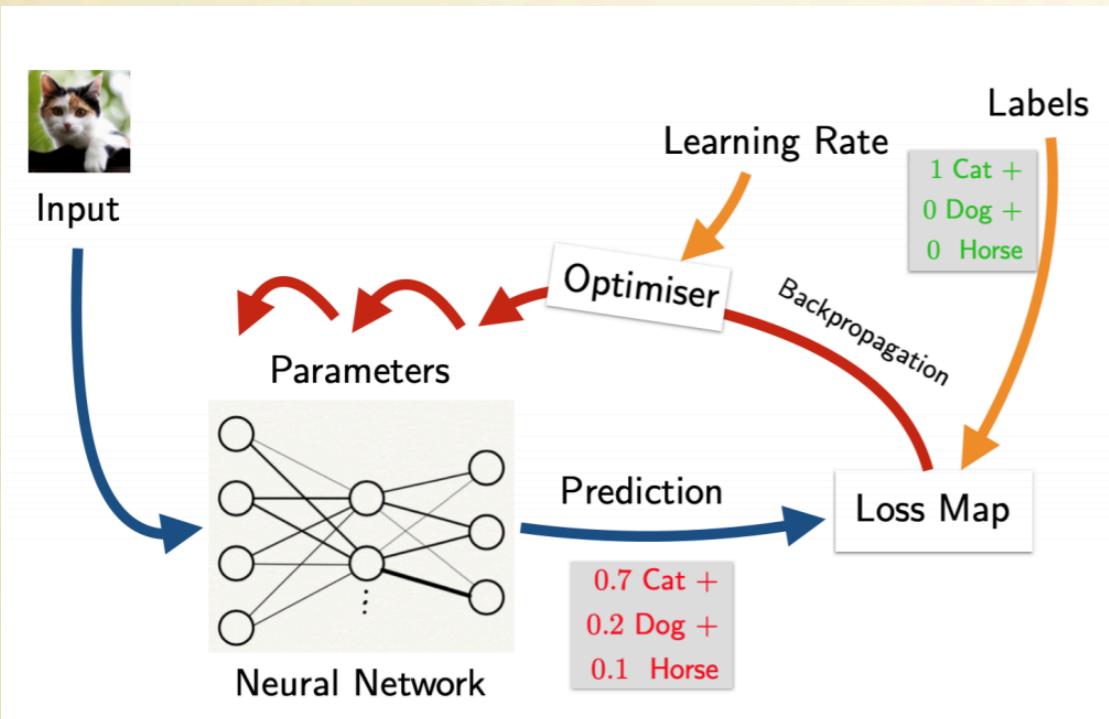
‘Stateful’ optimiser such as **Adagrad**, **Adam**, **Nesterov**, **Momentum**, etc. can be all formalised uniformly as reparametrisations.

Optimisers are reparametrisations

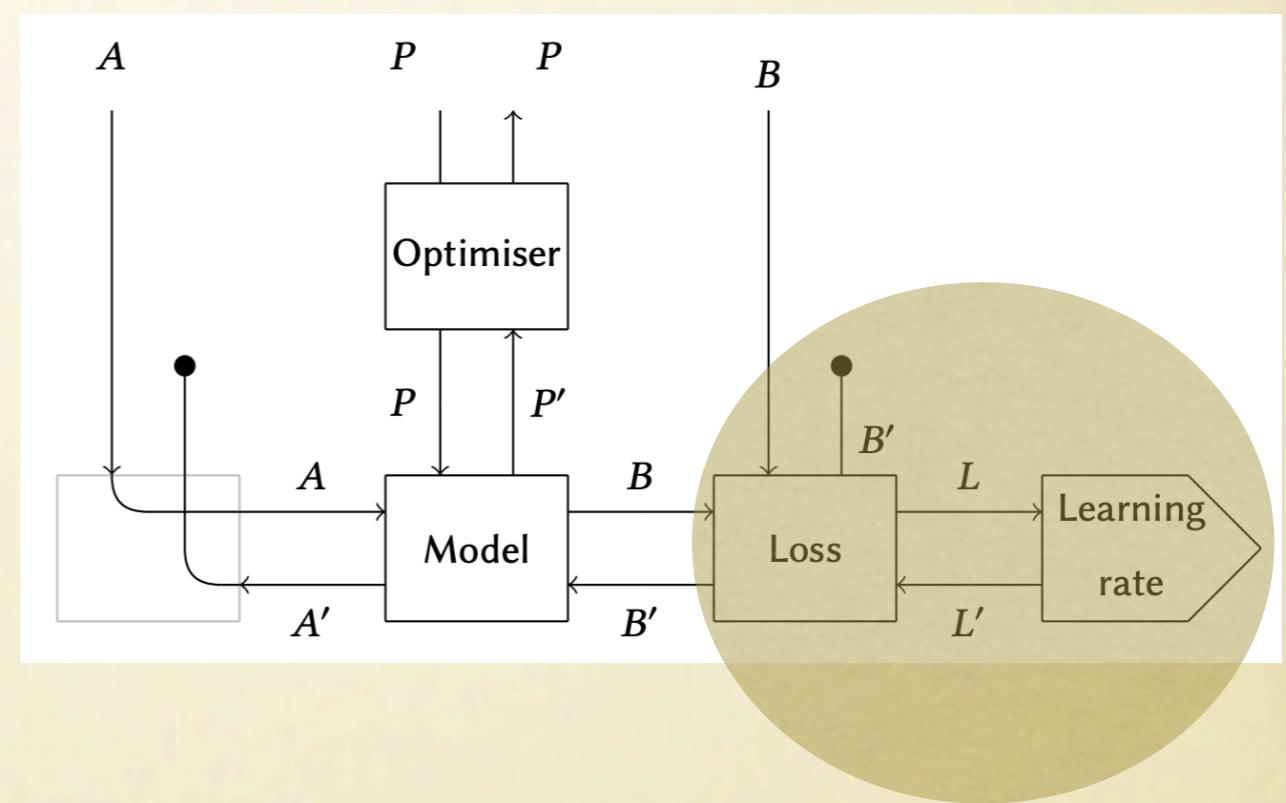
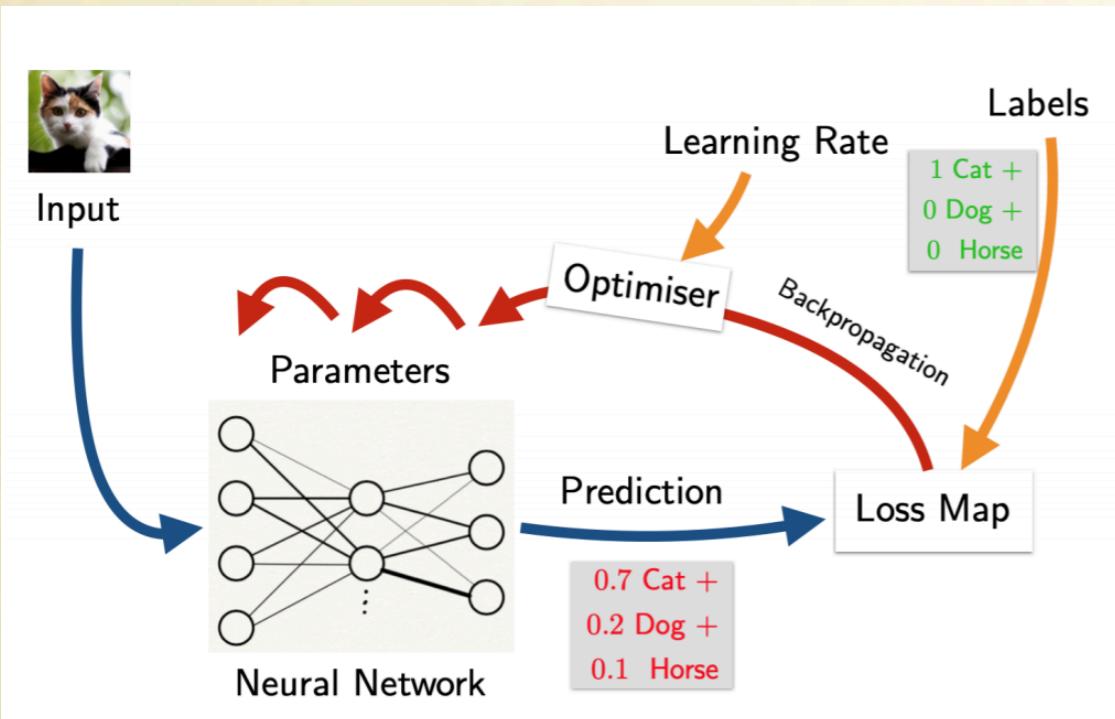


‘Stateful’ optimiser such as **Adagrad**, **Adam**, **Nesterov**, **Momentum**, etc. can be all formalised uniformly as reparametrisations.

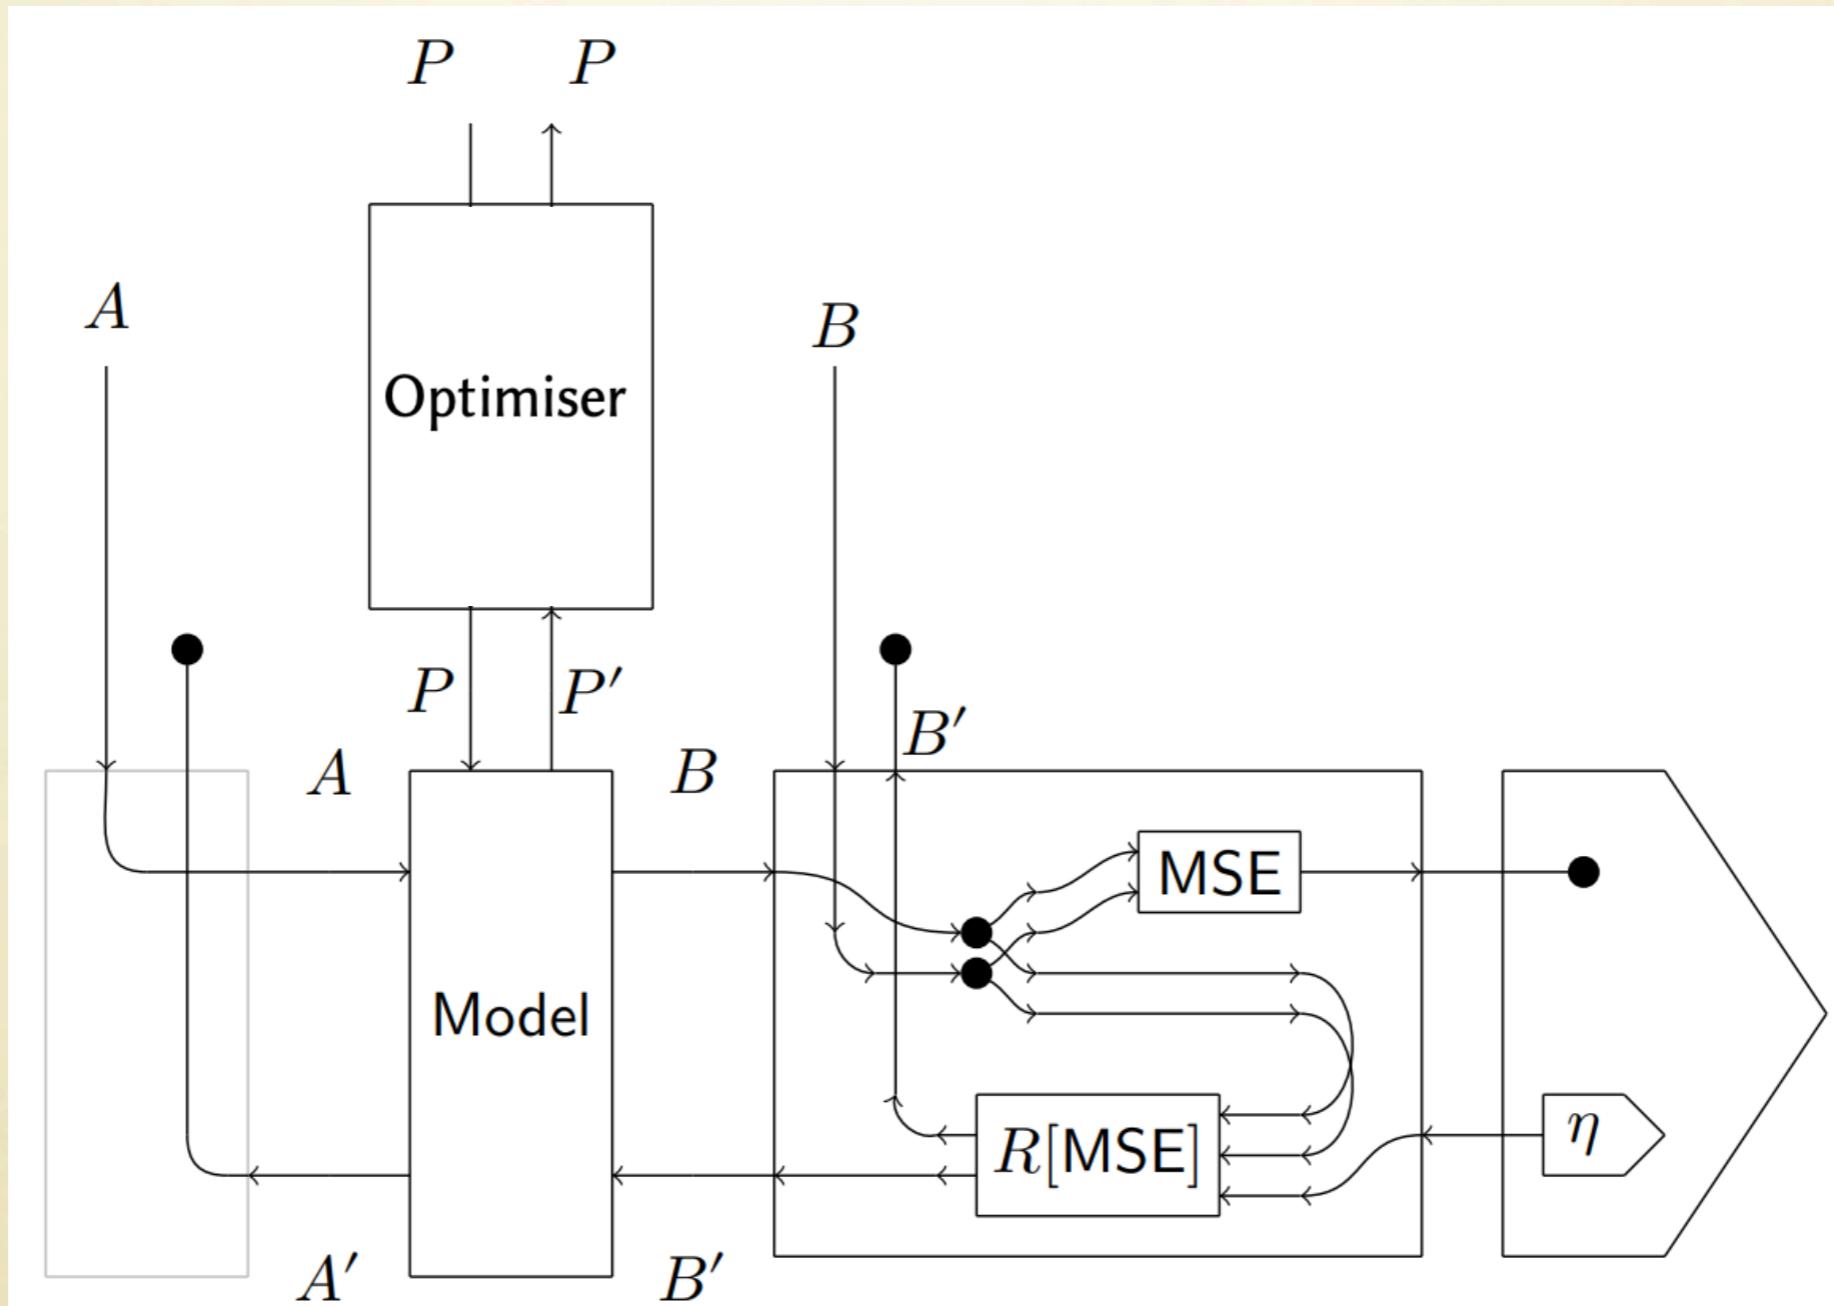
Where we are, so far



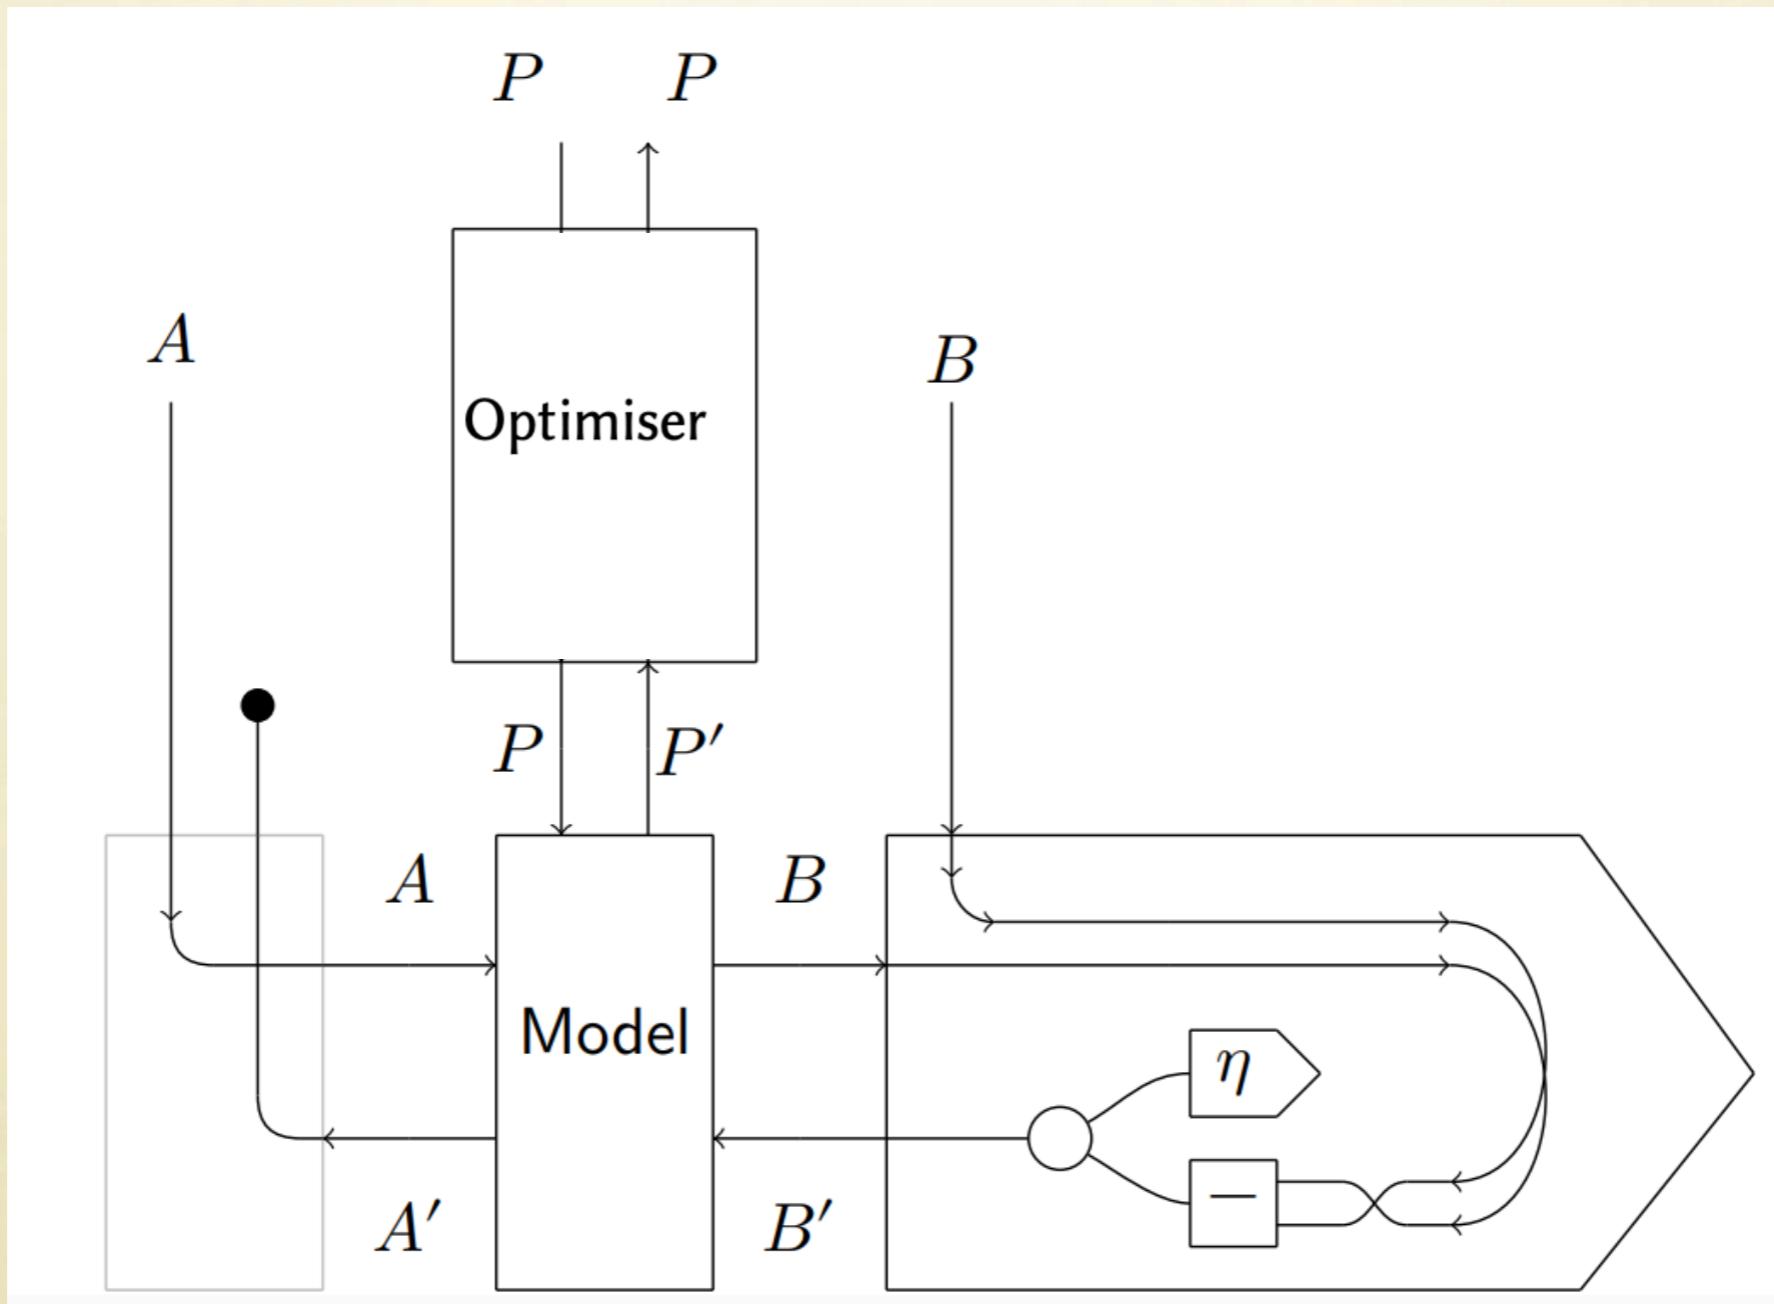
Where we are, so far



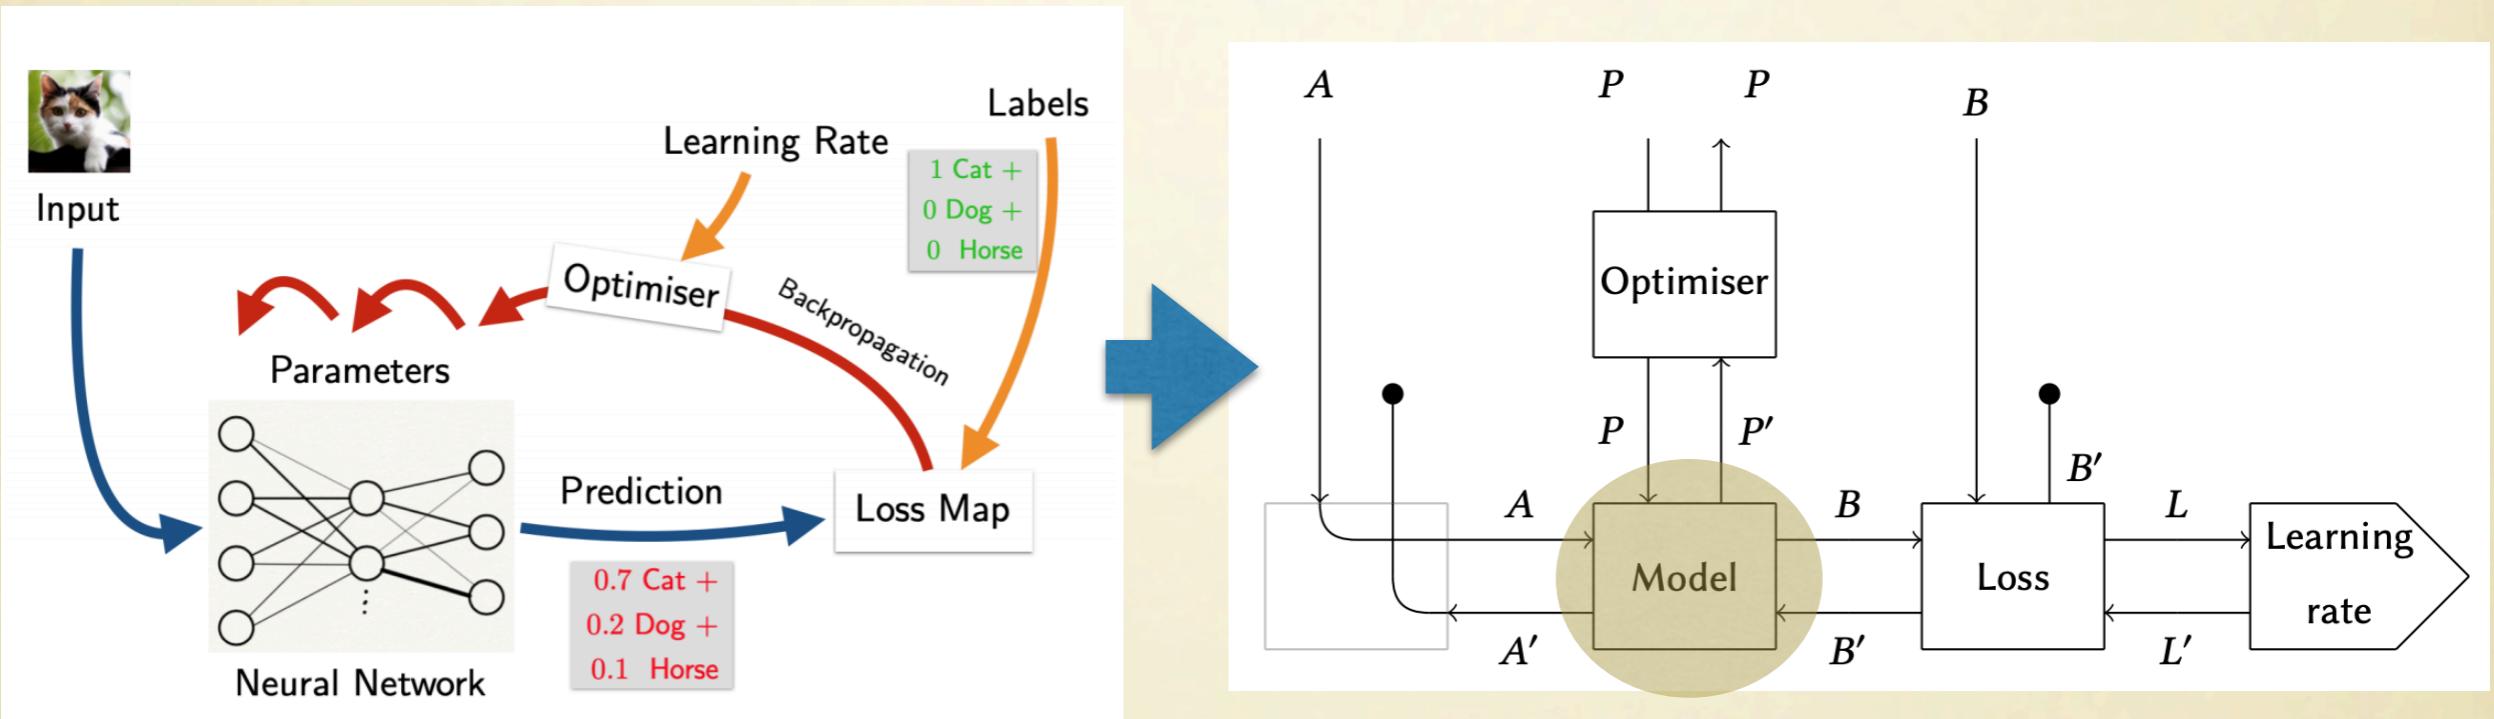
Loss map as parametric lens



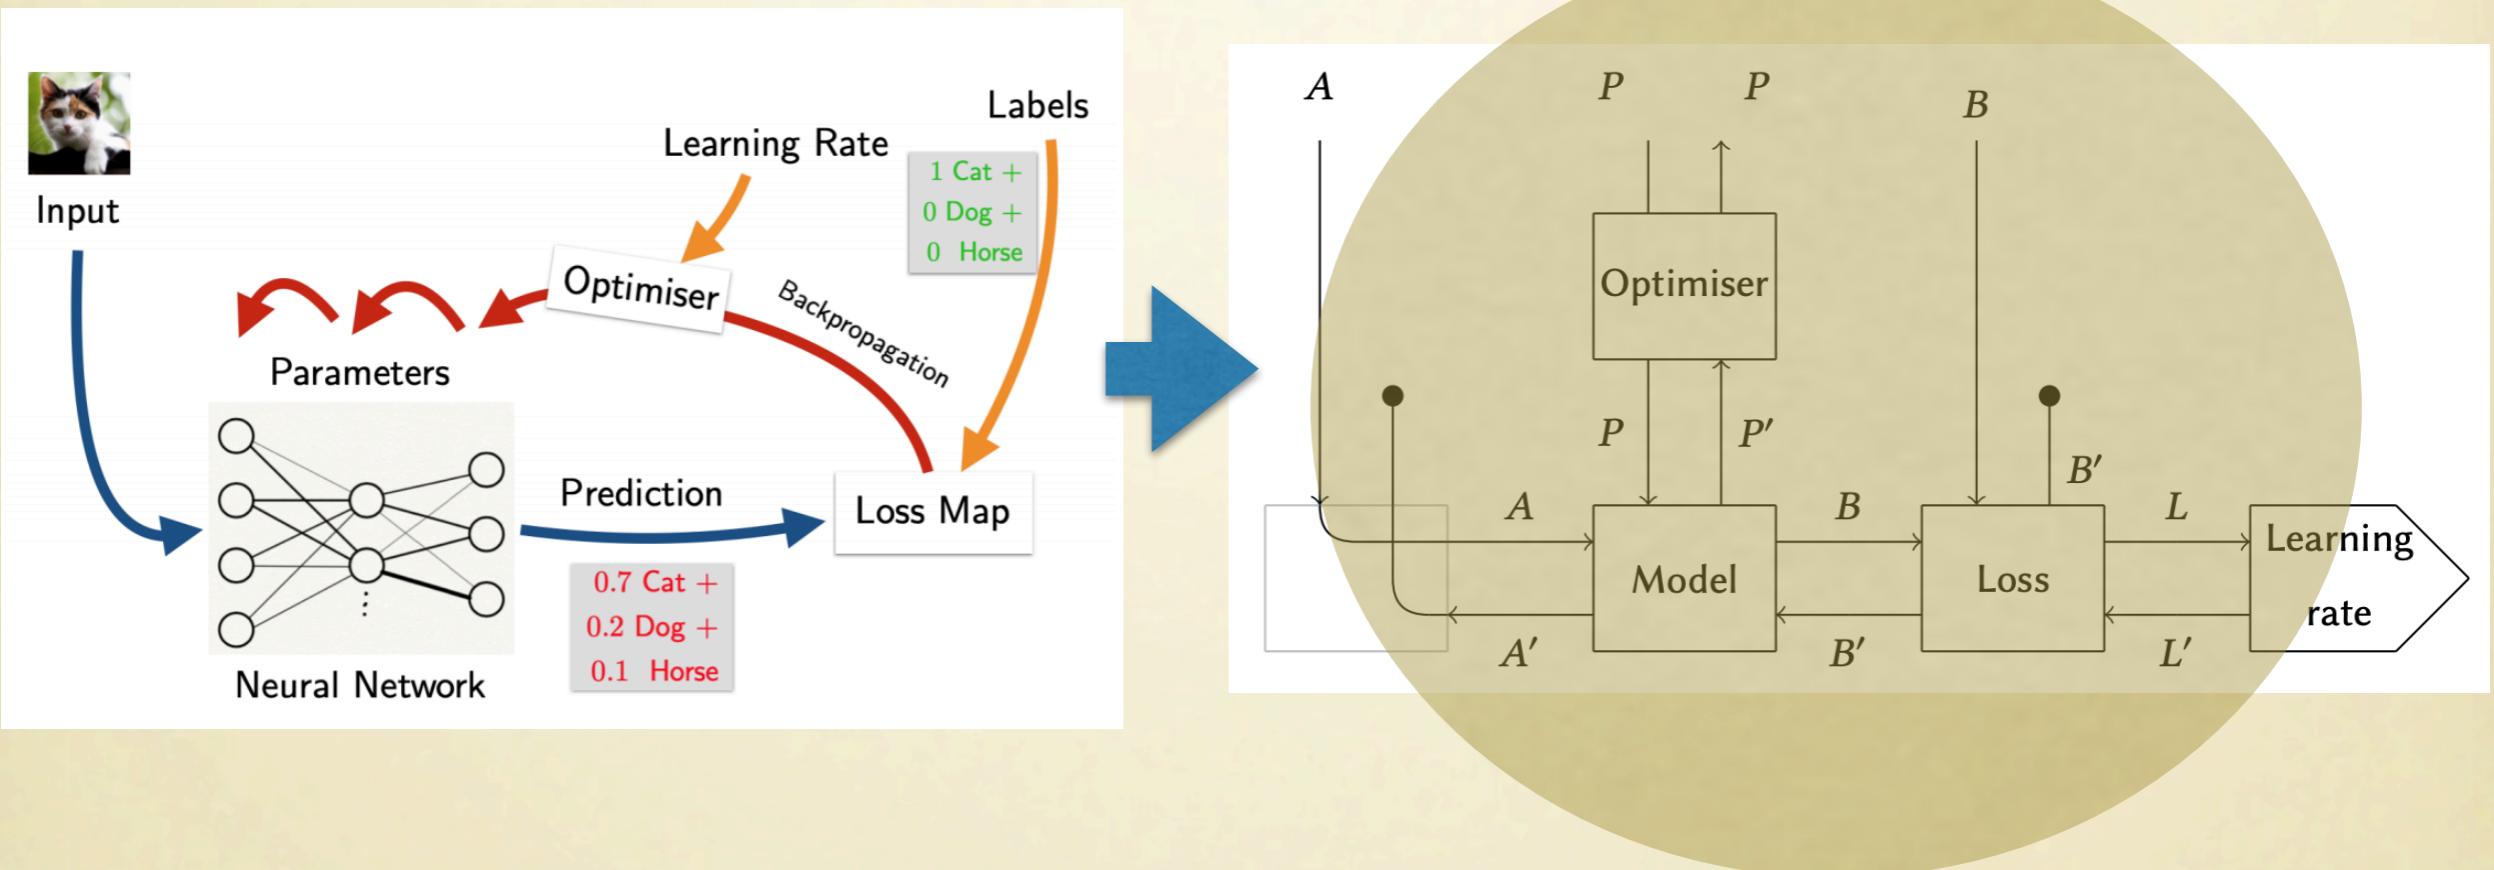
Loss map as parametric lens



Where we are, so far



Where we are, so far



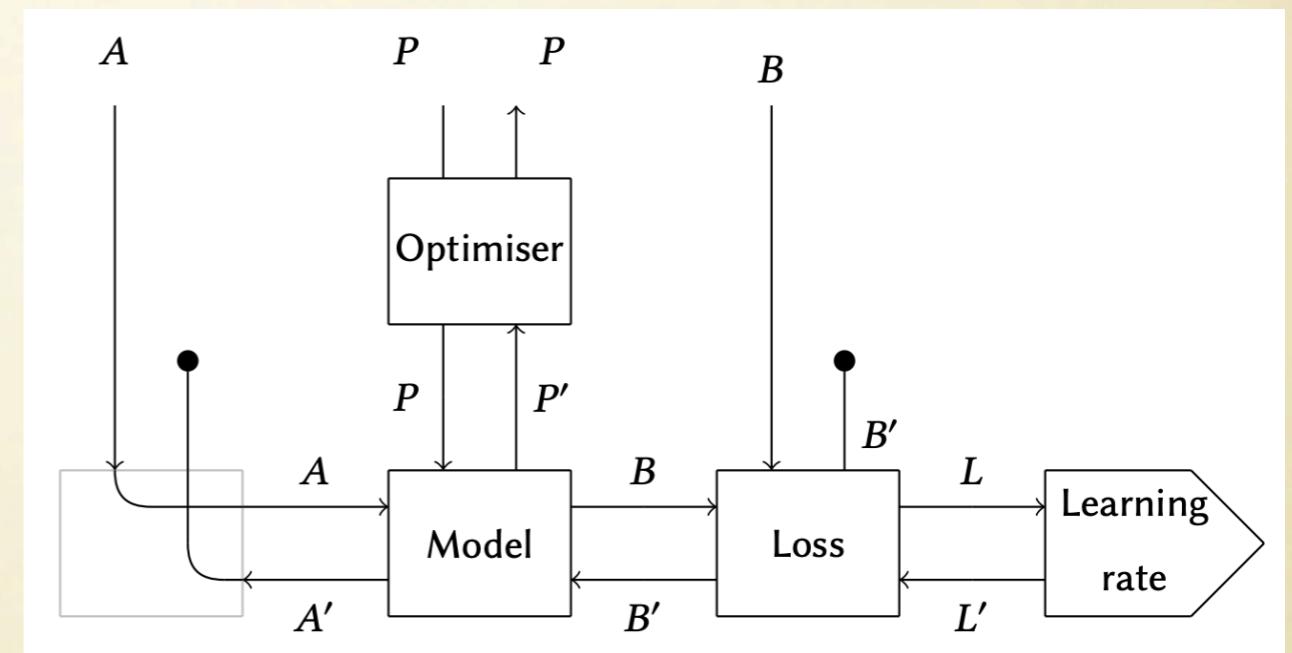
Putting it all together

A parametric lens of type

$$(1, 1) \rightarrow (1, 1)$$

A lens of type

$$(A \times P \times B, A' \times P' \times B') \rightarrow (1, 1)$$



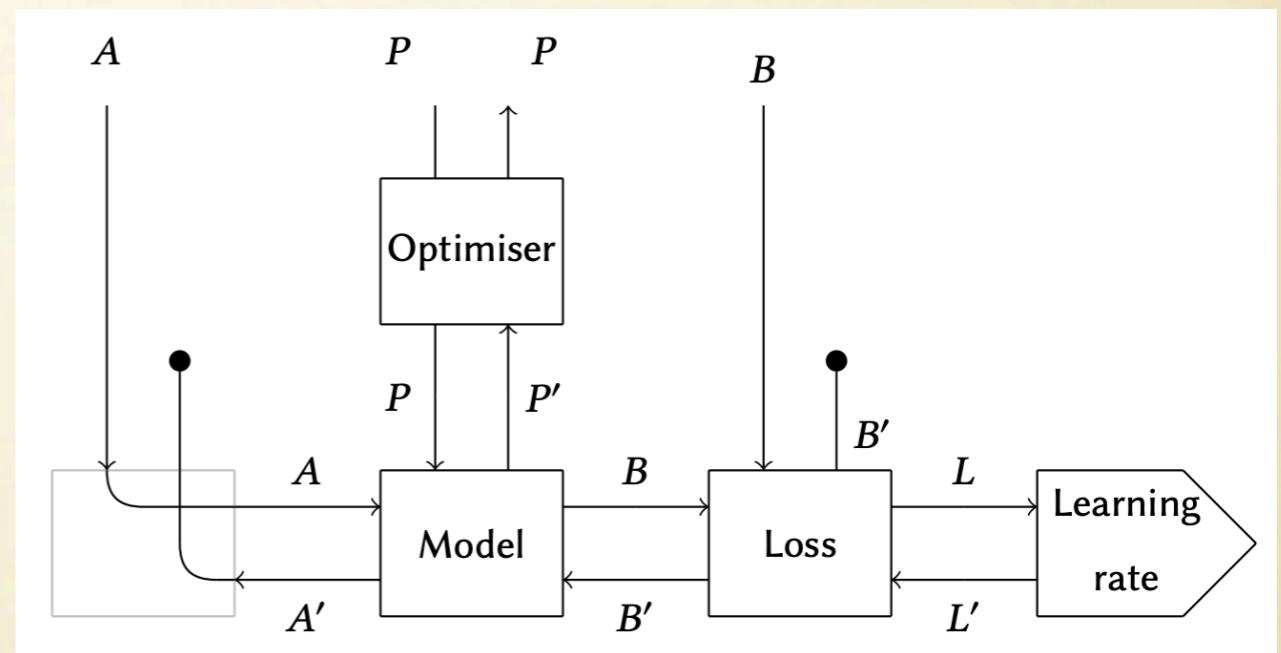
Putting it all together

A parametric lens of type

$$(1, 1) \rightarrow (1, 1)$$

A lens of type

$$(A \times P \times B, A' \times P' \times B') \rightarrow (1, 1)$$

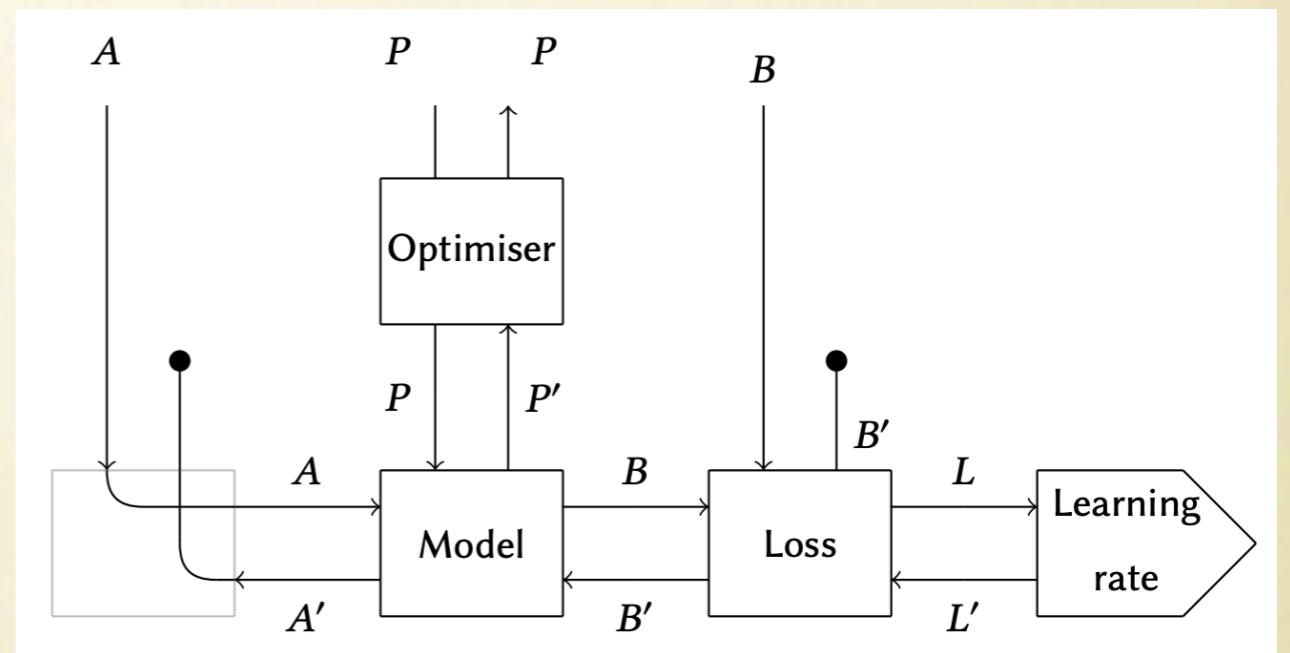


put $A \times P \times B \rightarrow A' \times P' \times B'$

put(a, p, b_t) = (a', p', b'_t)

$b_p = f(p, a)$ $(b'_t, b'_p) = R[\text{loss}](b_t, b_p, \alpha(\text{loss}(b_t, b_p)))$ $(p', a') = R[f](p, a, b'_p)$

Learning as iterative process

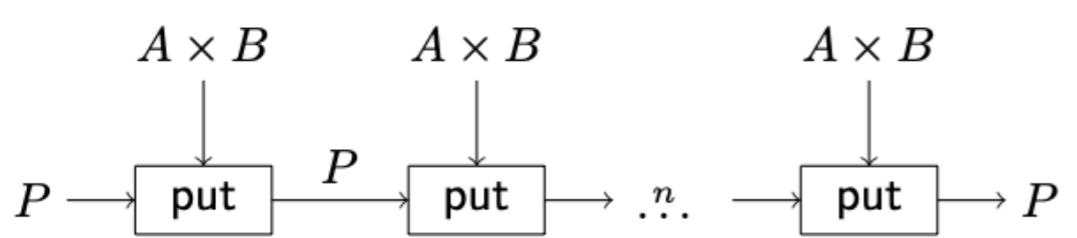


put $A \times P \times B \rightarrow A' \times P' \times B'$

$$\text{put}(a, p, b_t) = (a', p', b'_t)$$

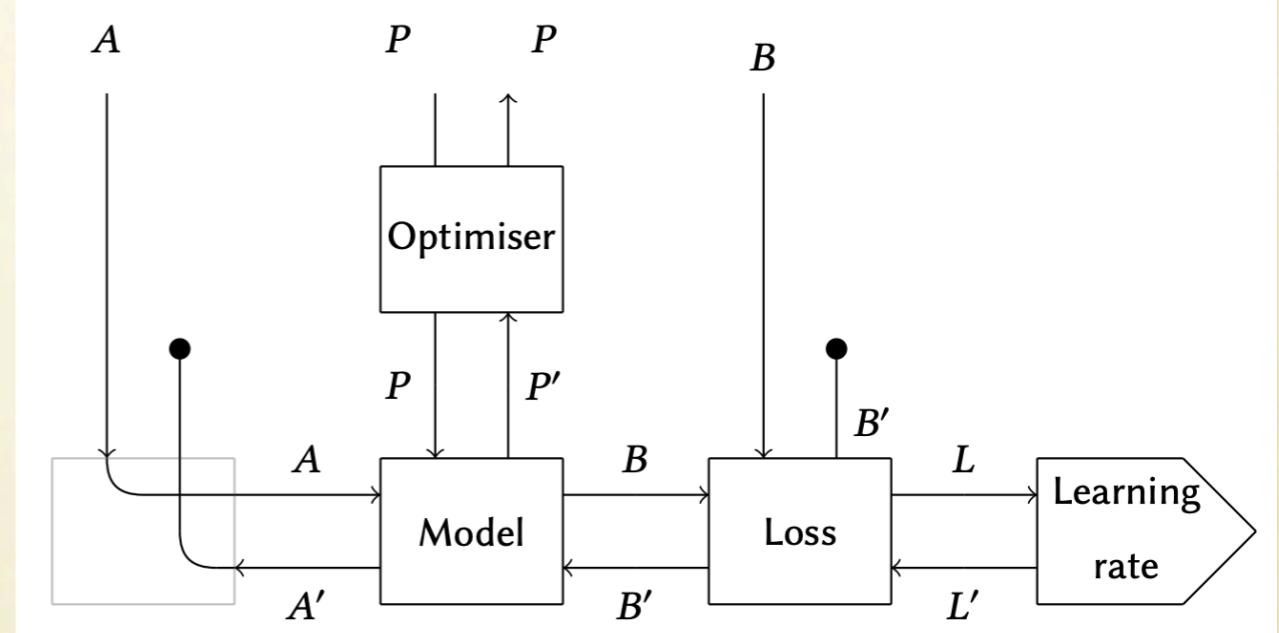
$$b_p = f(p, a) \quad (b'_t, b'_p) = R[\text{loss}](b_t, b_p, \alpha(\text{loss}(b_t, b_p))) \quad (p', a') = R[f](p, a, b'_p)$$

Learning as iterative process



Dataset becomes parameters

Parameters become inputs/outputs

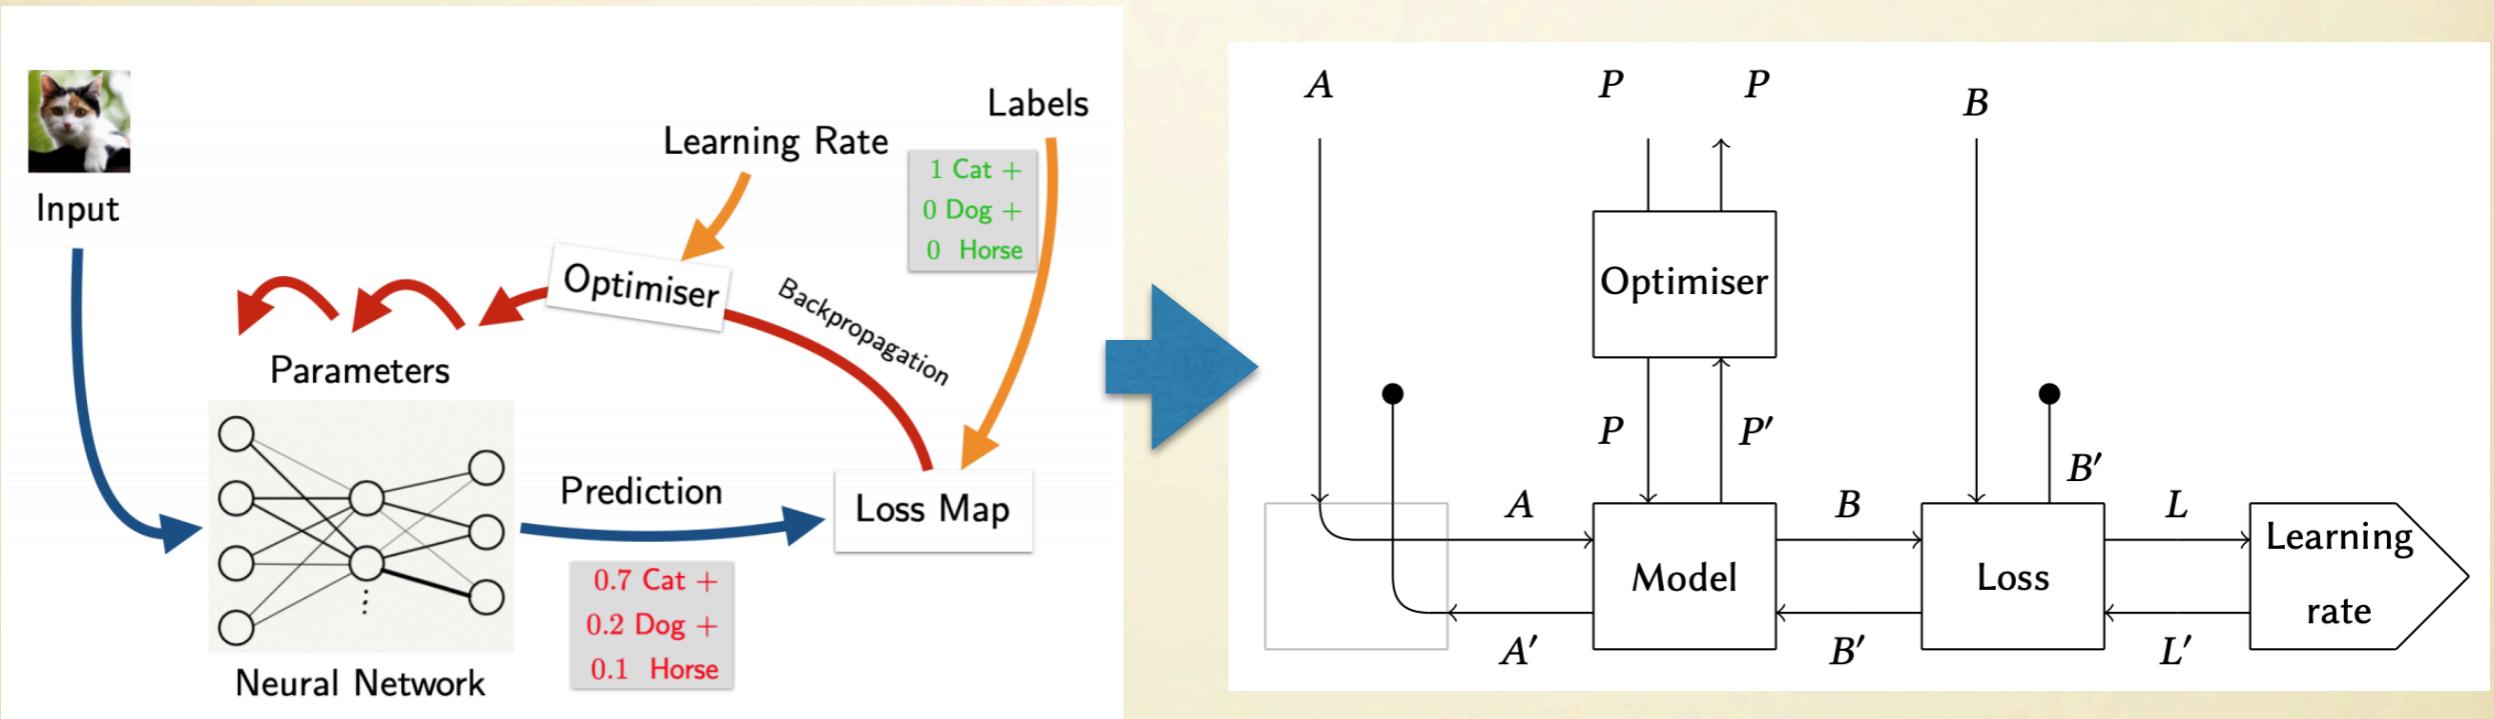


$$\text{put} : A \times P \times B \rightarrow A' \times P' \times B'$$

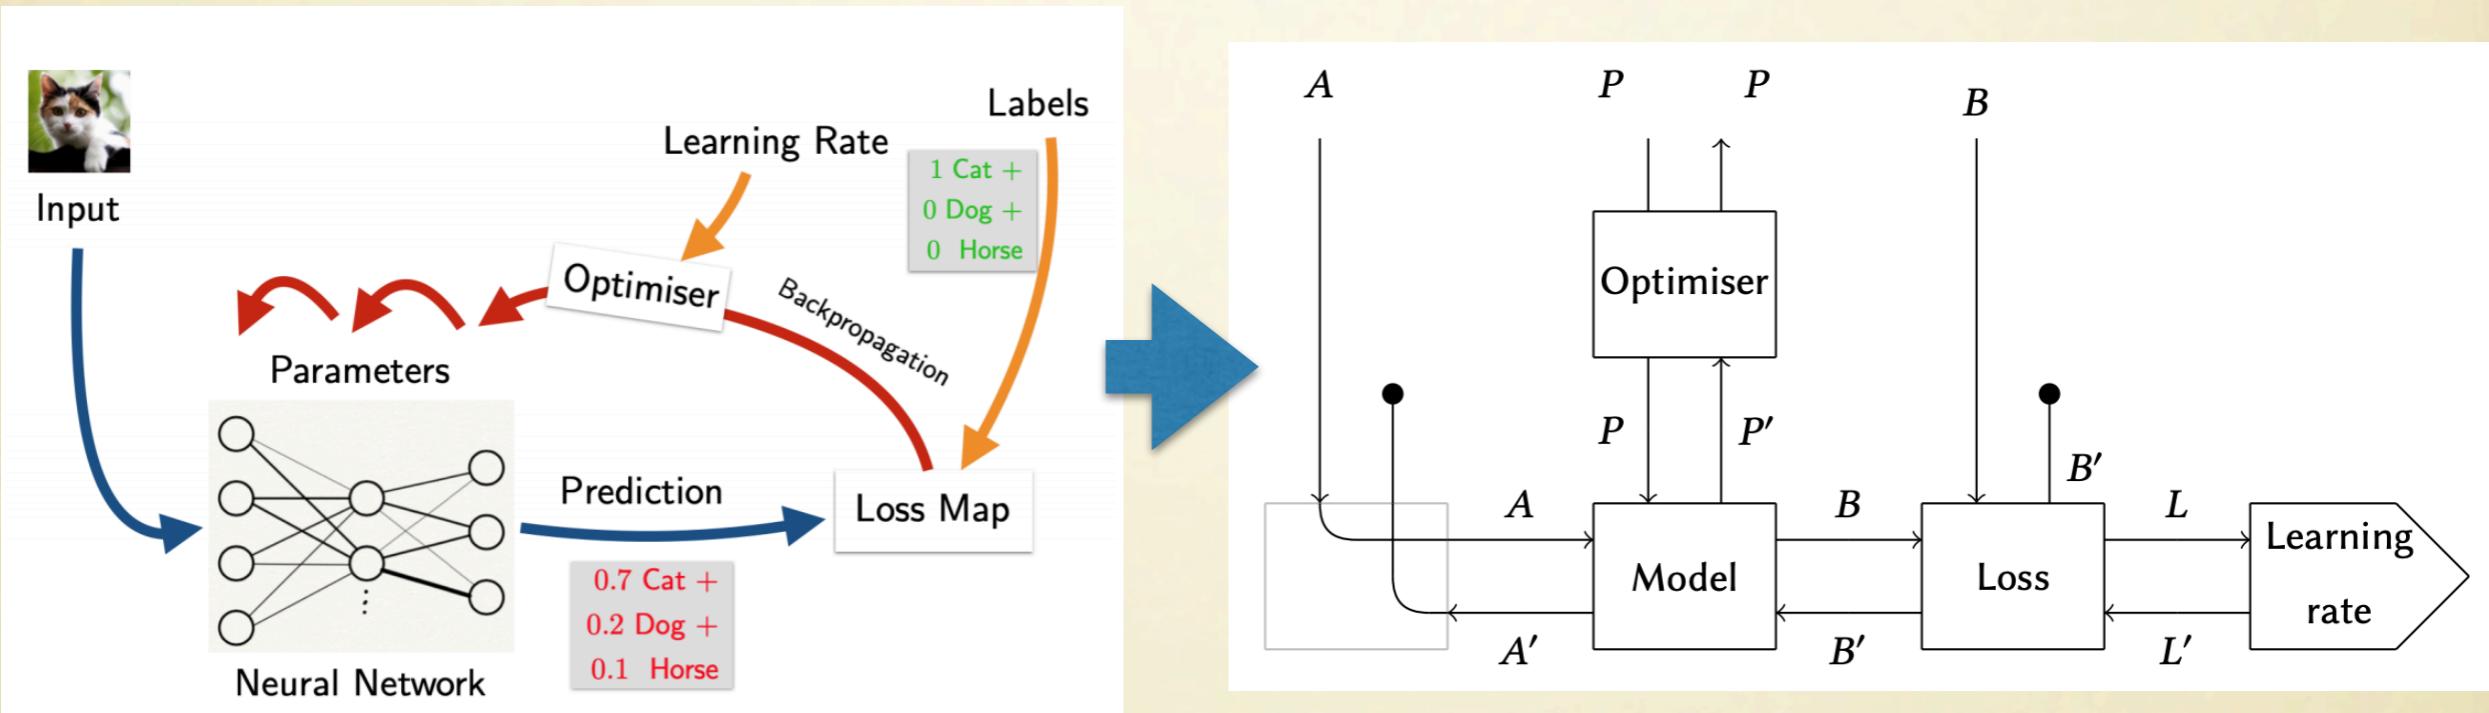
$$\text{put}(a, p, b_t) = (a', p', b'_t)$$

$$b_p = f(p, a) \quad (b'_t, b'_p) = R[\text{loss}](b_t, b_p, \alpha(\text{loss}(b_t, b_p))) \quad (p', a') = R[f](p, a, b'_p)$$

Where we are, so far



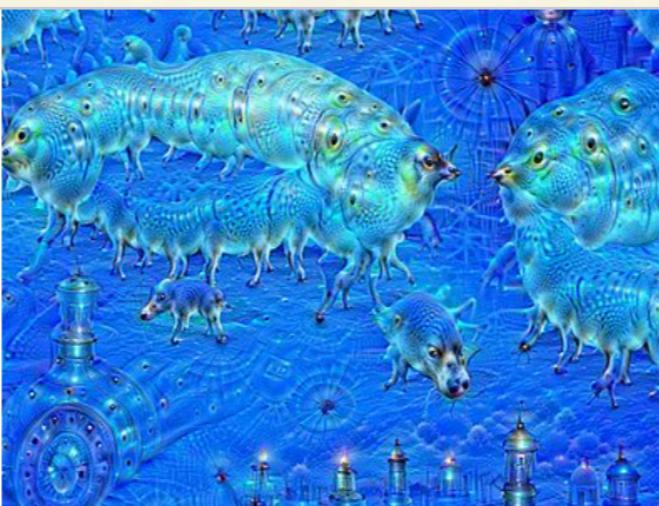
Where we are, so far



Variations I: Deep Dreaming



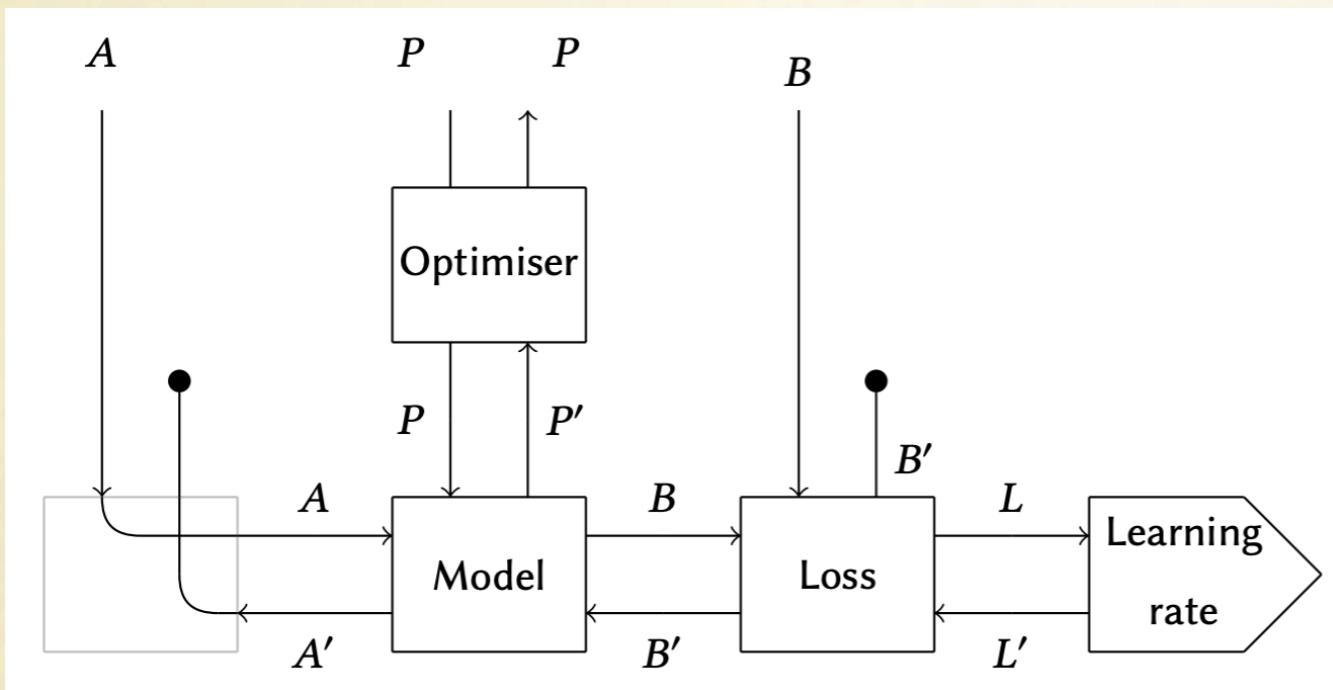
Input
Image



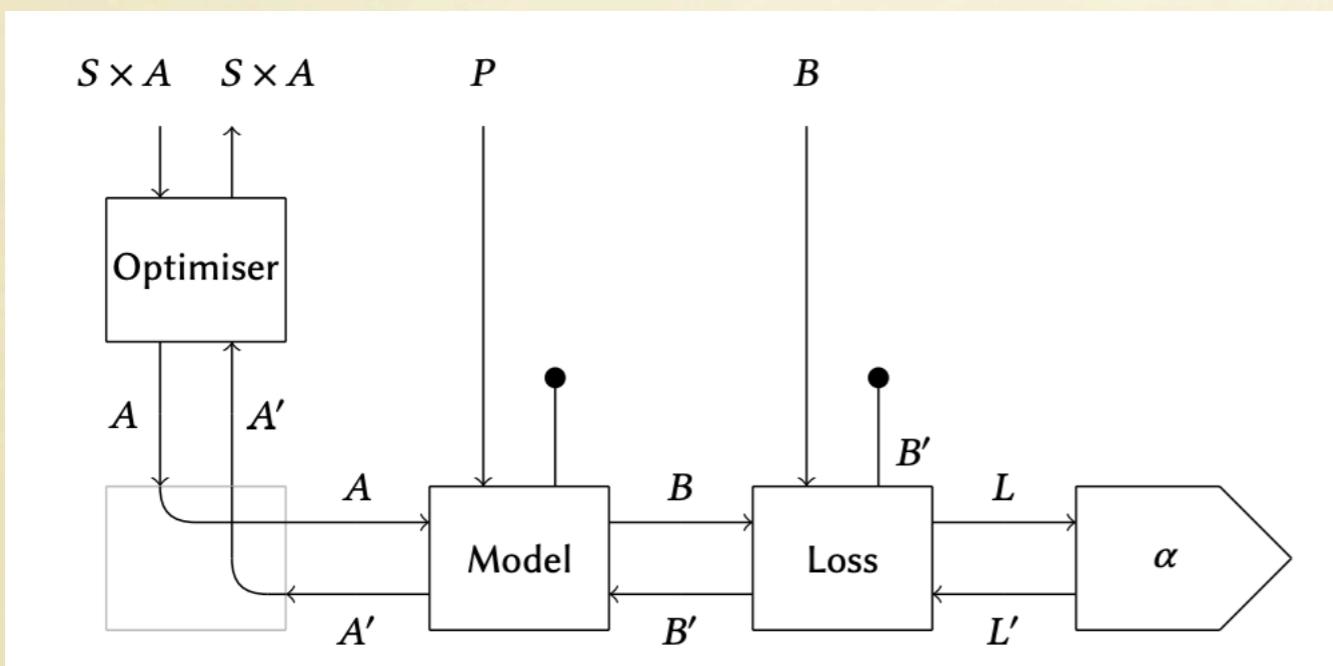
``Over-
processed''
Image



Variations I: Deep Dreaming

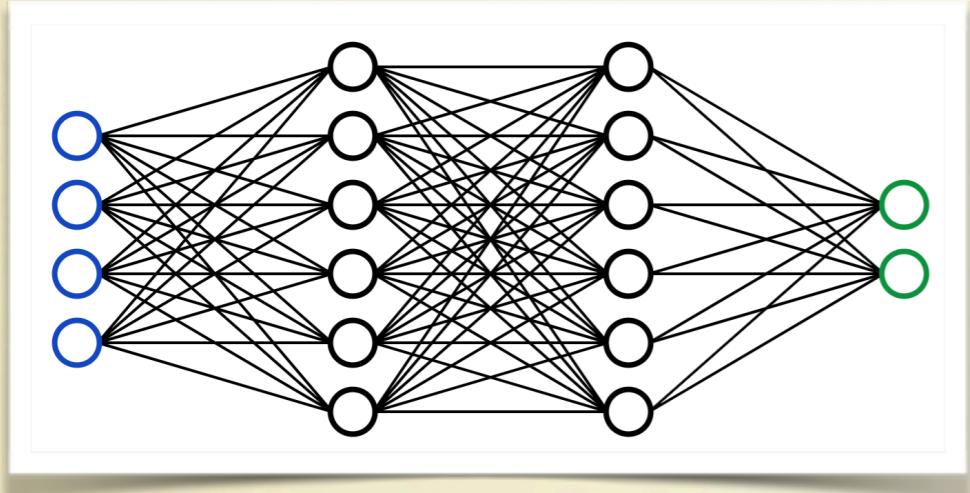


Parameter
Learning

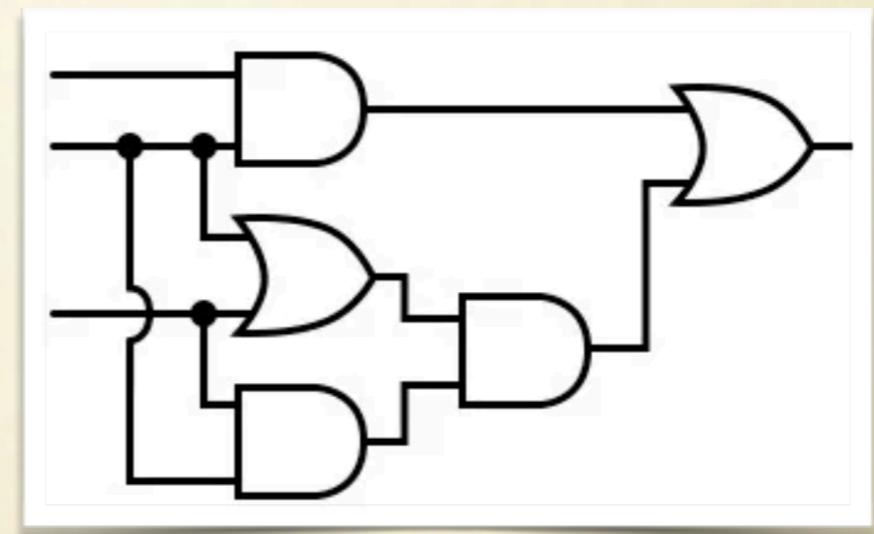


Deep
dreaming

Variations II: Learning Boolean circuits



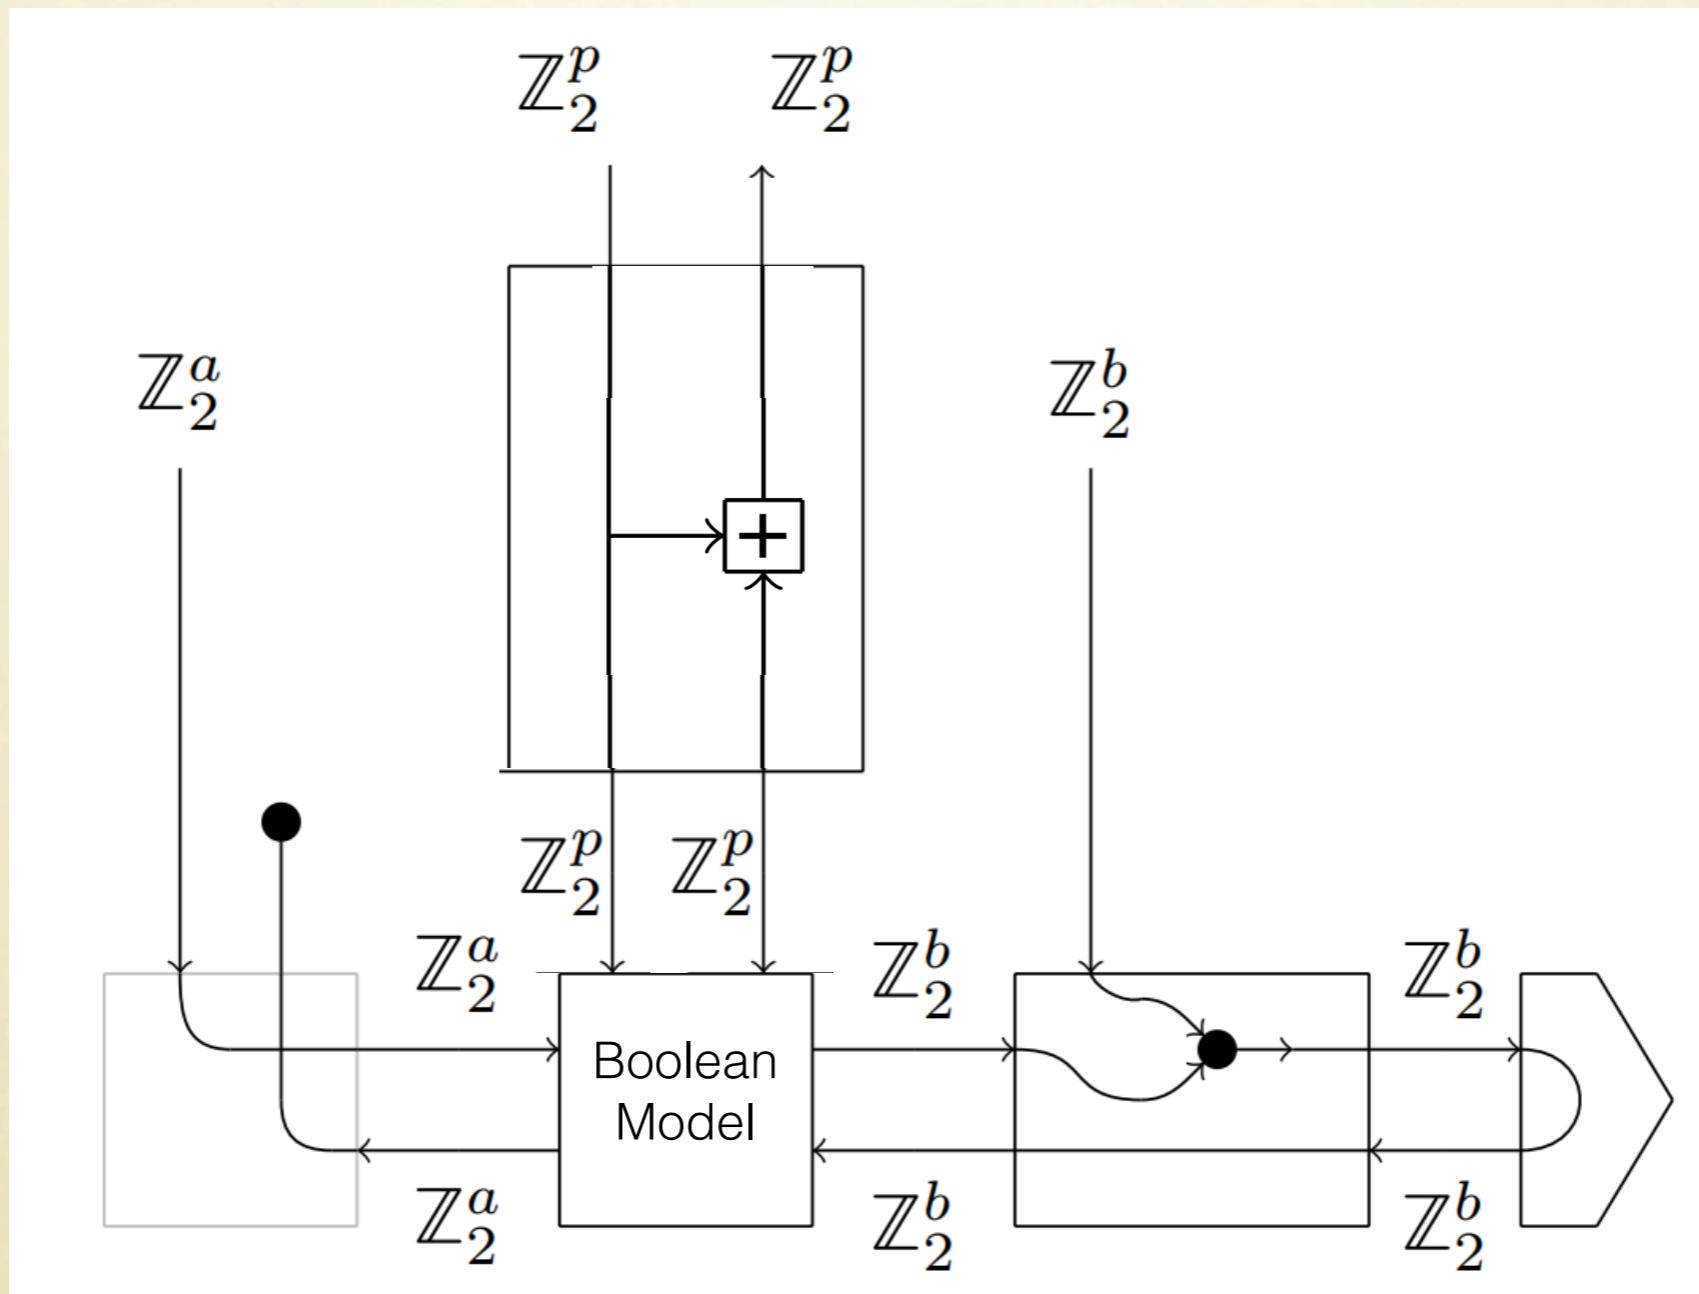
Neural networks



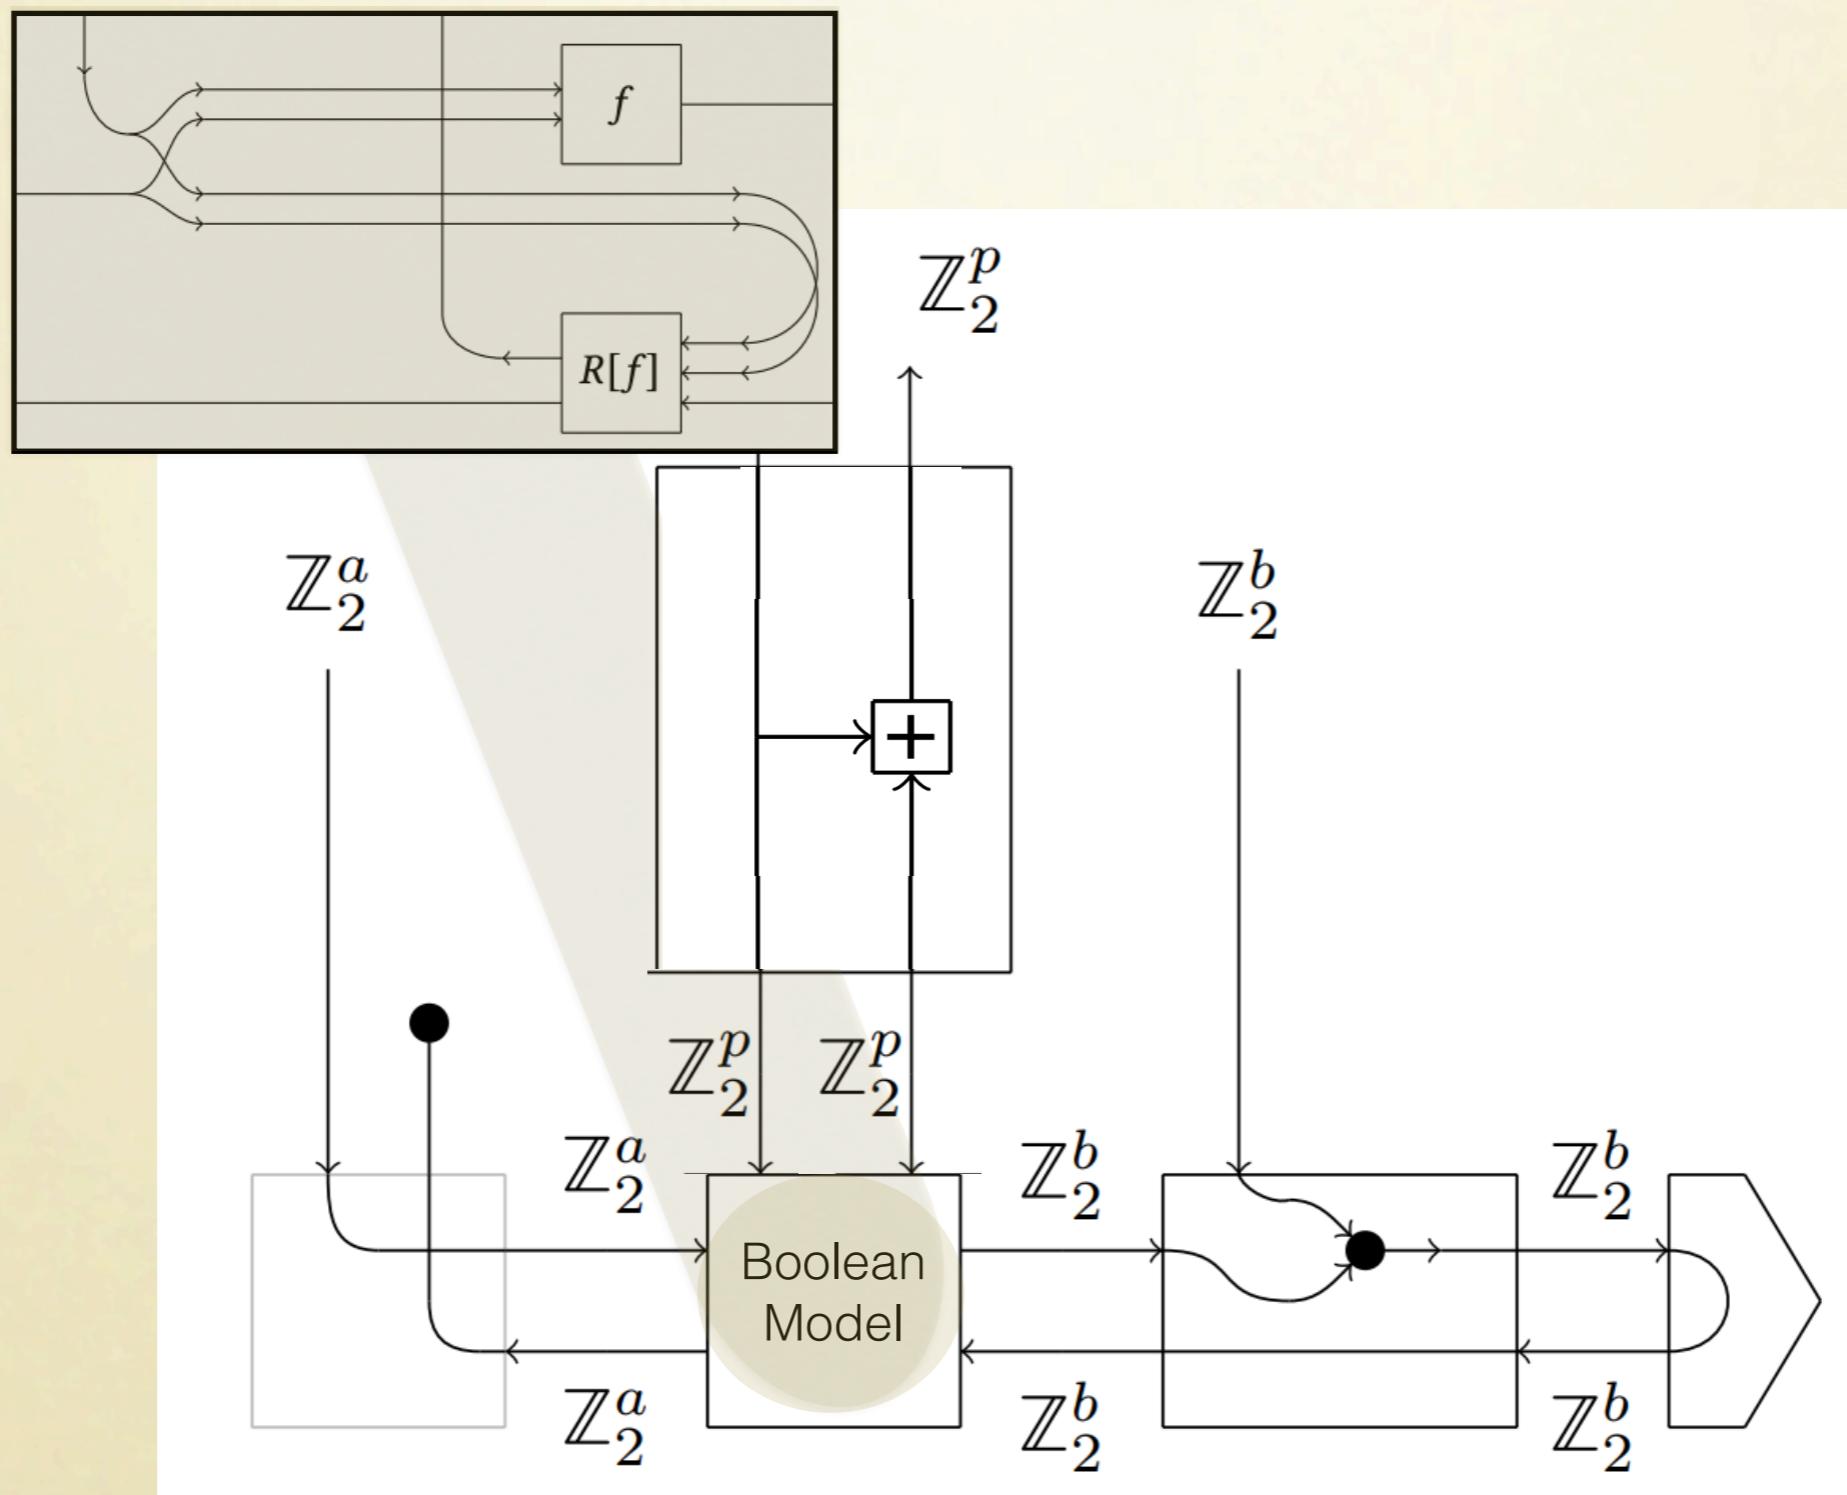
Boolean circuits

Instead of ***Smooth*** we work in ***Bool***

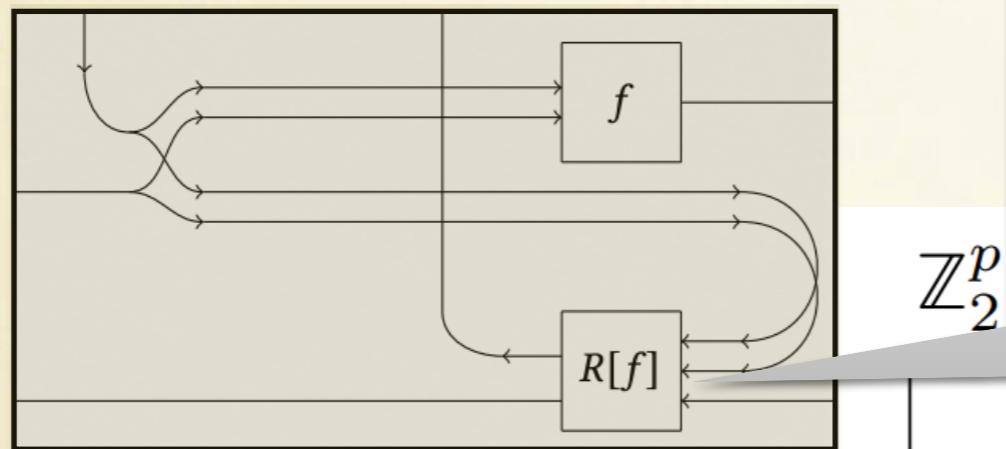
Variations II: Learning Boolean circuits



Variations II: Learning Boolean circuits

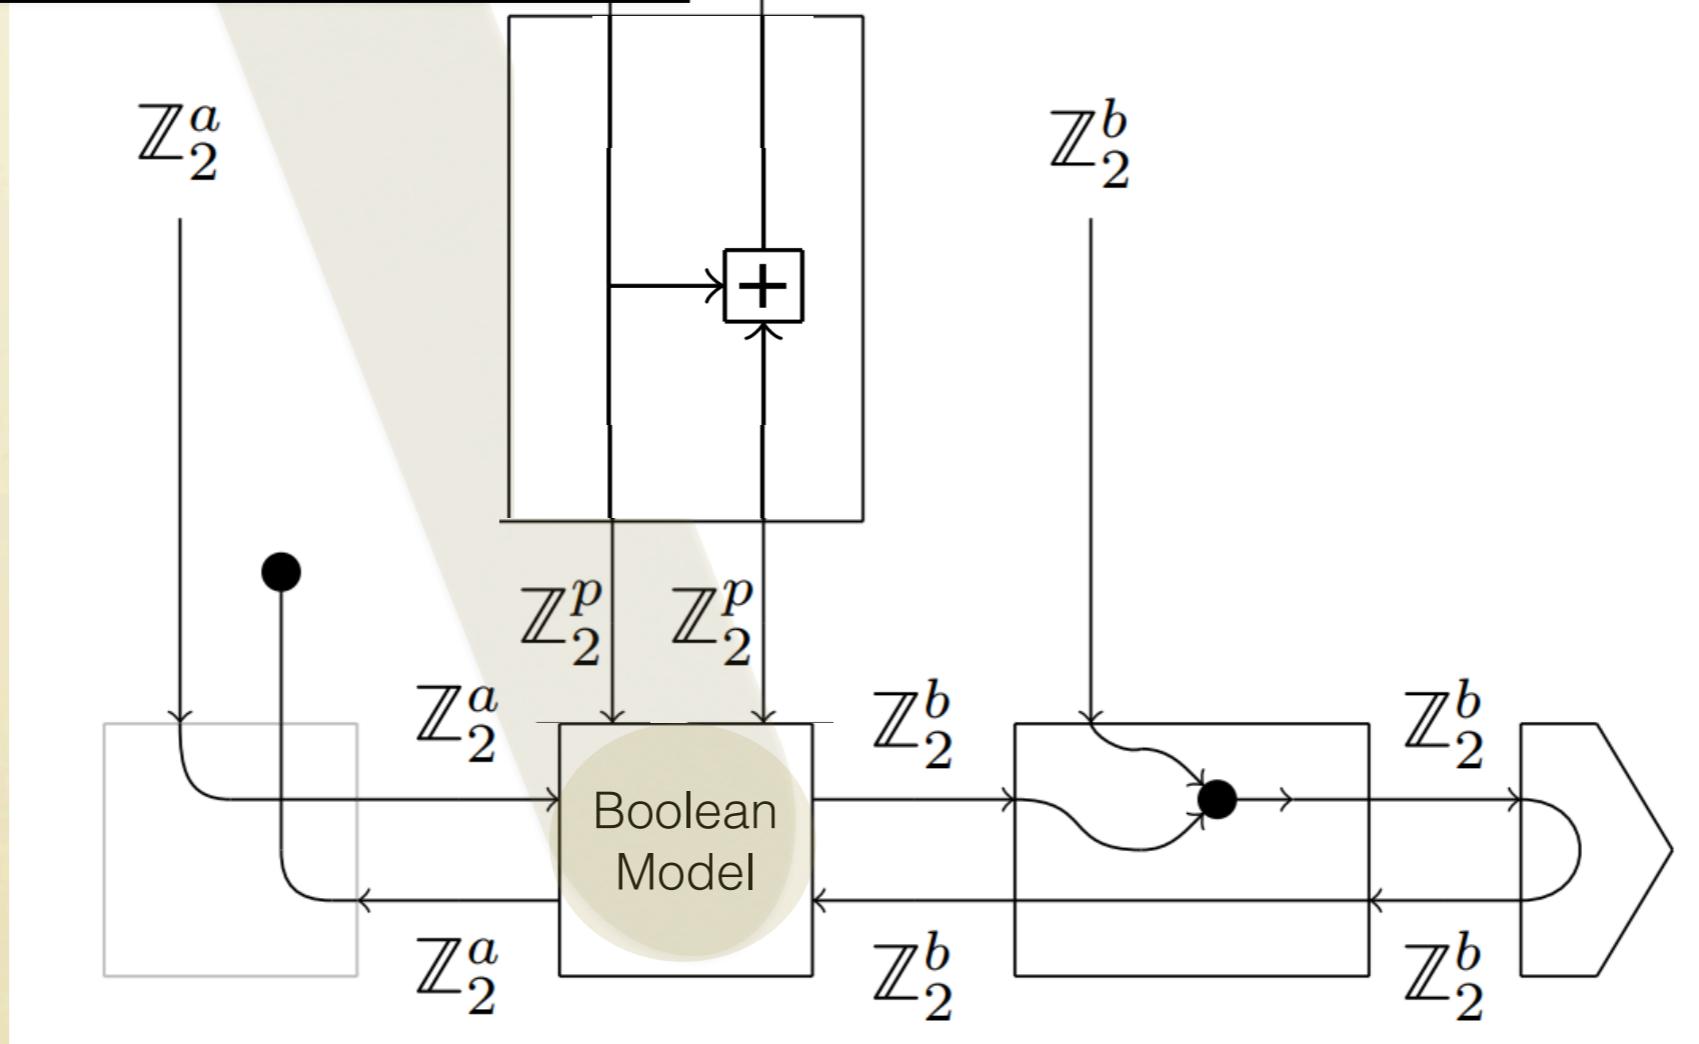


Variations II: Learning Boolean circuits



$R[f]$ is defined **inductively** on circuit syntax:

$\boxed{f}, \boxed{d} ::= \bullet | \overline{\bullet} | \wedge | \vee | \oplus | \ominus | \top | \perp | \dots | \infty | \boxed{f} | \boxed{d}$



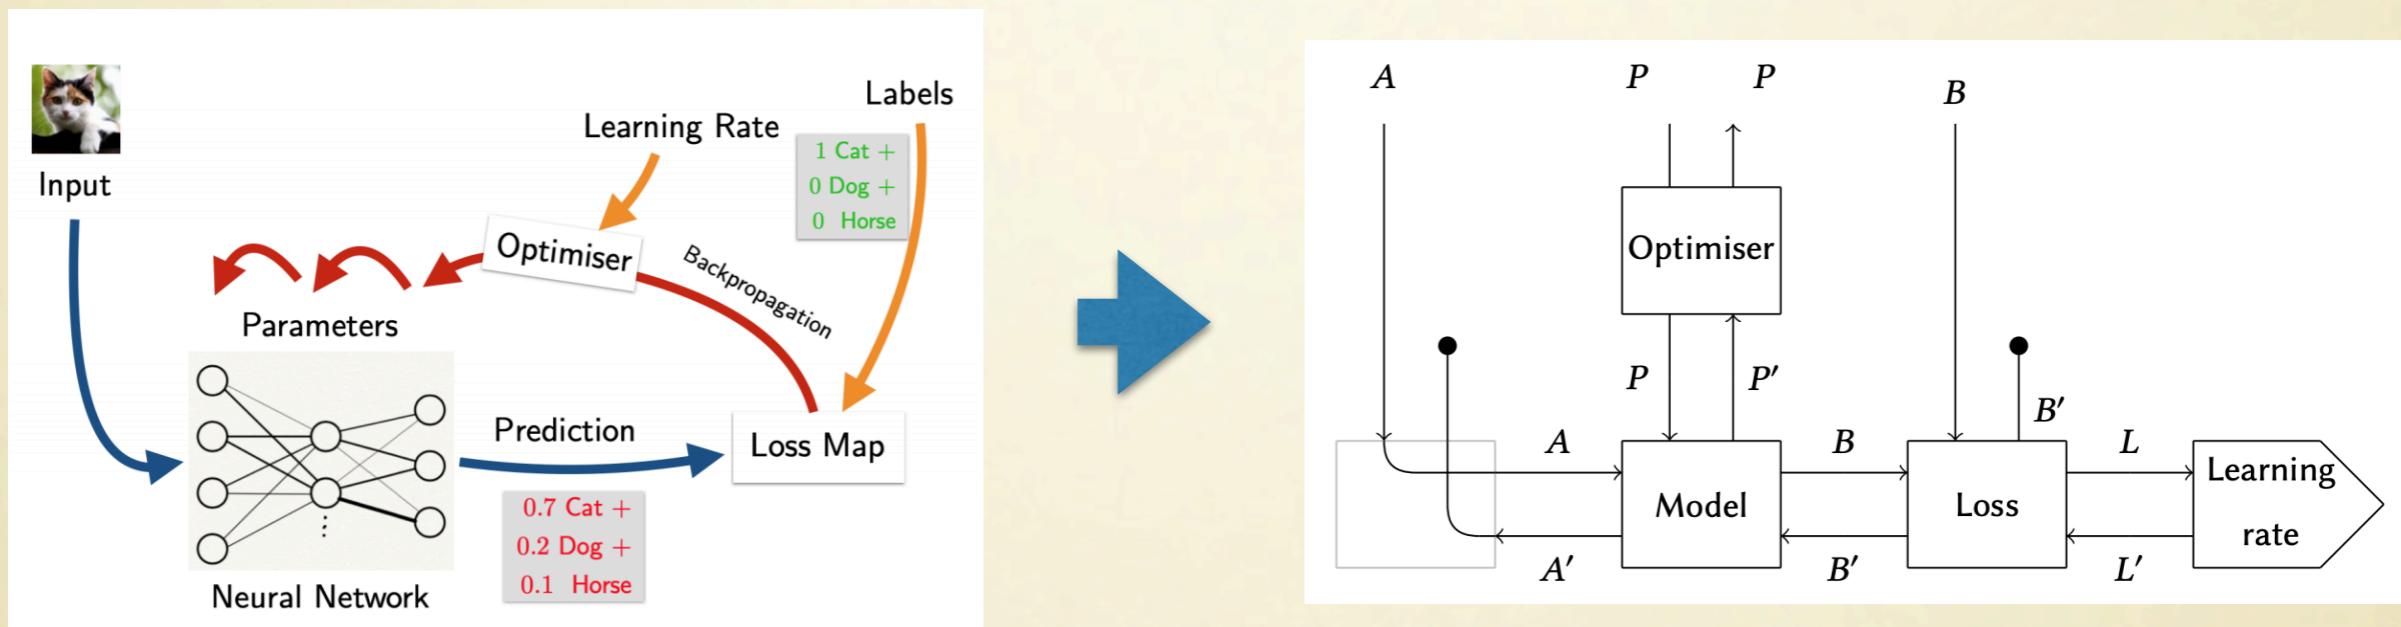
... and more

Generative Adversarial Networks

Convolutional Neural Networks

...

Discussion



A lens-theoretic semantics for gradient-based learning.

Discussion

Uniformity

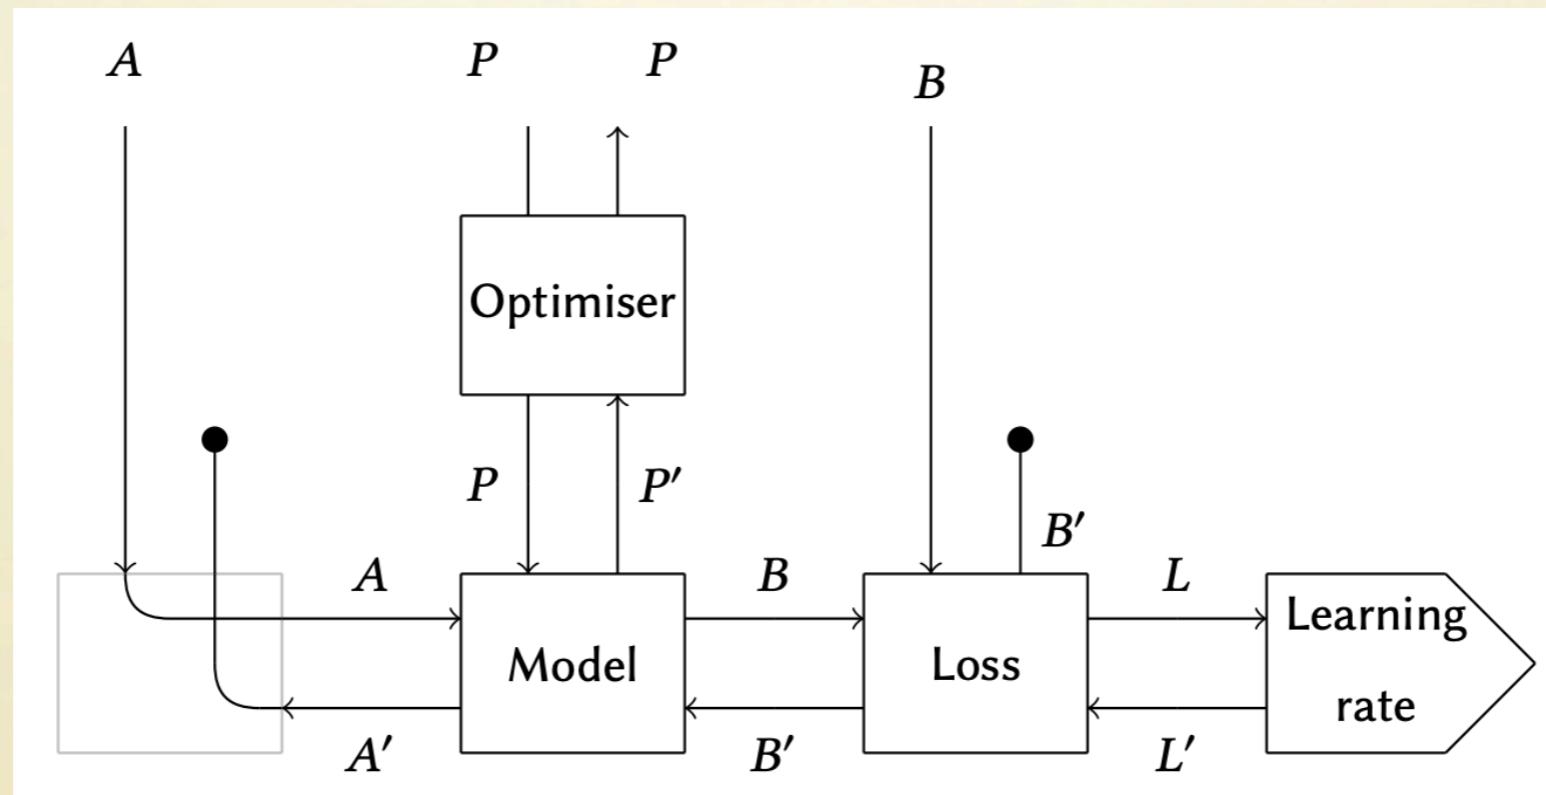
All components at work in the learning process are reduced to a single concept: parametric lenses.

Abstraction

A wide range of optimisers, models, loss maps, etc. are instances of the framework.

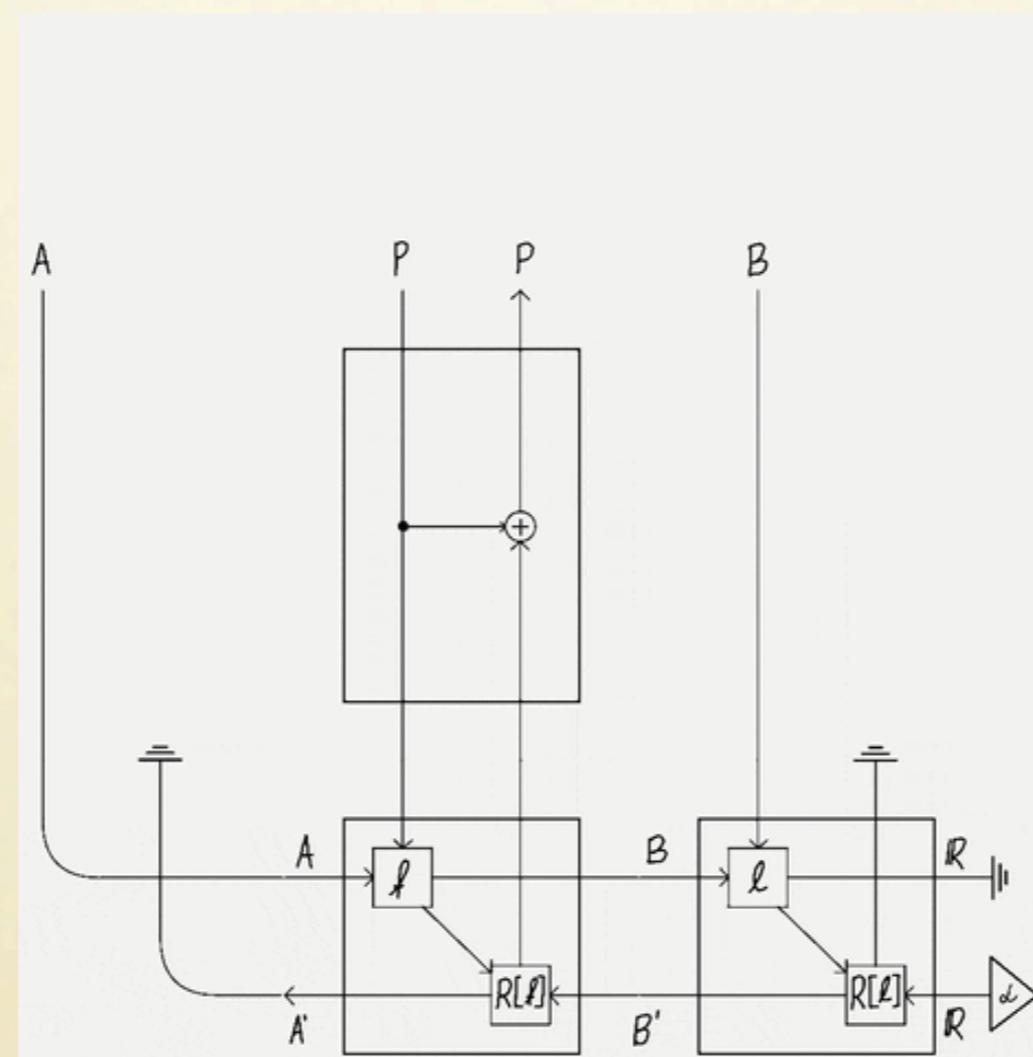
Discussion

Compositional reasoning



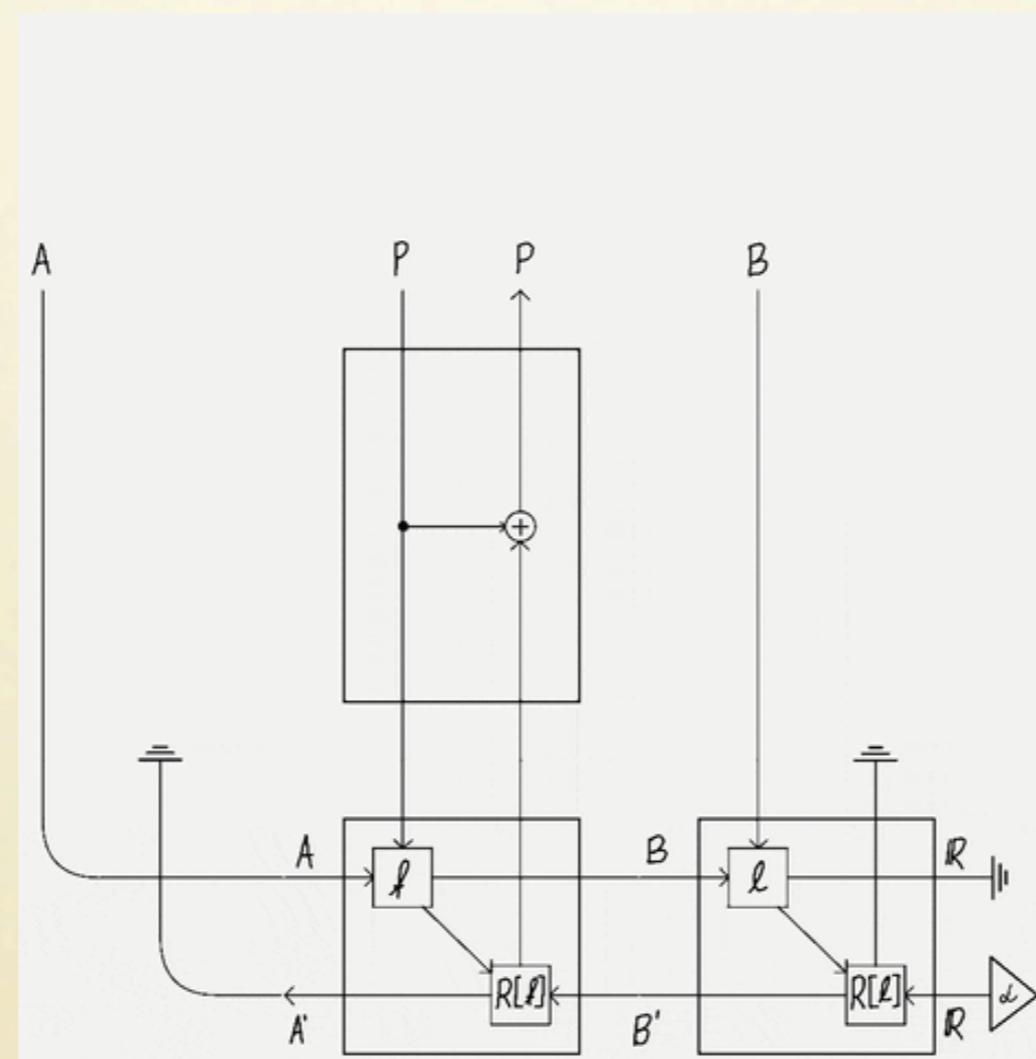
Discussion

Resource-sensitive Formalism



Discussion

Resource-sensitive Formalism



Going Forward

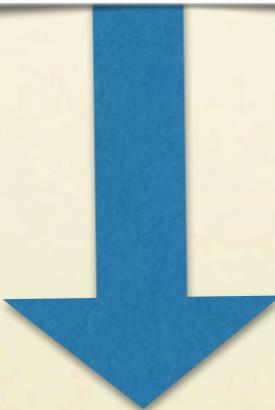
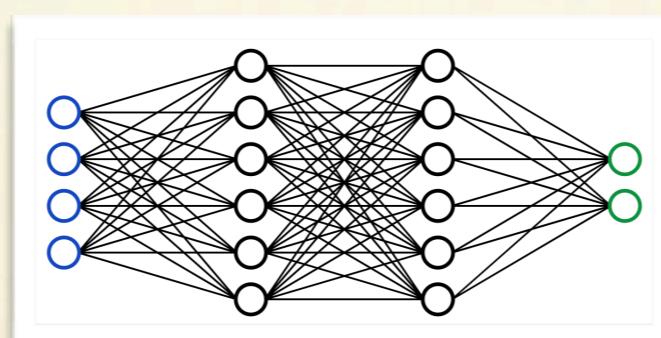
Going Forward

Learning over simpler models

Performance



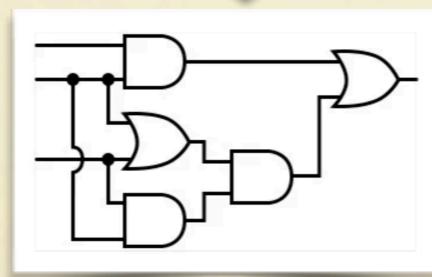
Accuracy



Performance



Accuracy



Boolean circuits

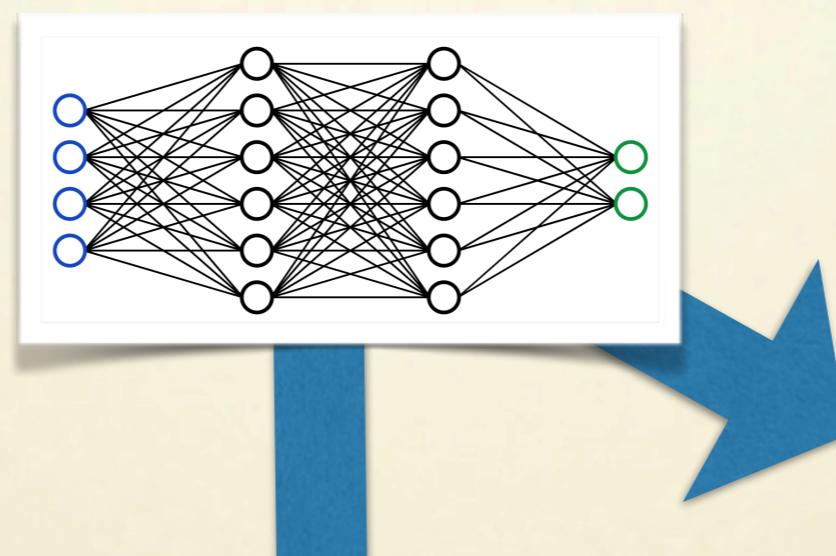
Going Forward

Learning over simpler models

Performance



Accuracy



Performance



Accuracy



Performance



Accuracy



Boolean circuits

Polynomial circuits
cf. saturation arithmetic,
zero-knowledge proofs

Going Forward

Programming with lenses?

<https://github.com/statusfailed/numeric-optics-python/>

<http://catgrad.com/p/reverse-derivative-ascent>

Going Forward

Programming with lenses?

<https://github.com/statusfailed/numeric-optics-python/>

<http://catgrad.com/p/reverse-derivative-ascent>

```
def dense(a, b, activation):
    return linear(a, b) >> bias(b) >> activation
```

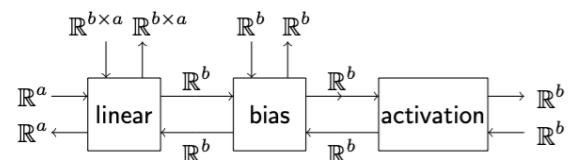
Going Forward

Programming with lenses?

<https://github.com/statusfailed/numeric-optics-python/>

<http://catgrad.com/p/reverse-derivative-ascent>

```
def dense(a, b, activation):
    return linear(a, b) >> bias(b) >> activation
```



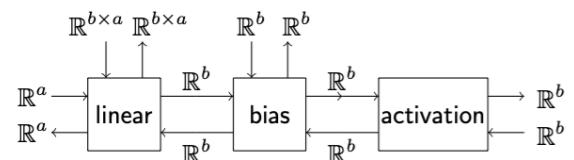
Going Forward

Programming with lenses?

<https://github.com/statusfailed/numeric-optics-python/>

<http://catgrad.com/p/reverse-derivative-ascent>

```
def dense(a, b, activation):
    return linear(a, b) >> bias(b) >> activation
```



```
apply_update(basic_update, dense) >> loss >> learning_rate(epsilon)
```

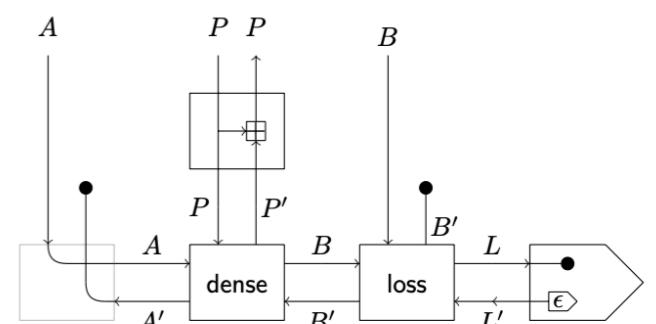
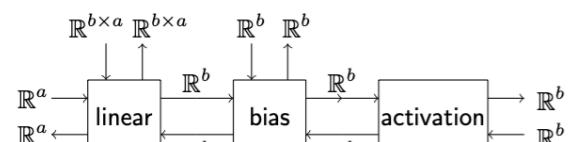
Going Forward

Programming with lenses?

<https://github.com/statusfailed/numeric-optics-python/>

<http://catgrad.com/p/reverse-derivative-ascent>

```
def dense(a, b, activation):
    return linear(a, b) >> bias(b) >> activation
```



```
apply_update(basic_update, dense) >> loss >> learning_rate(ε)
```

Going Forward

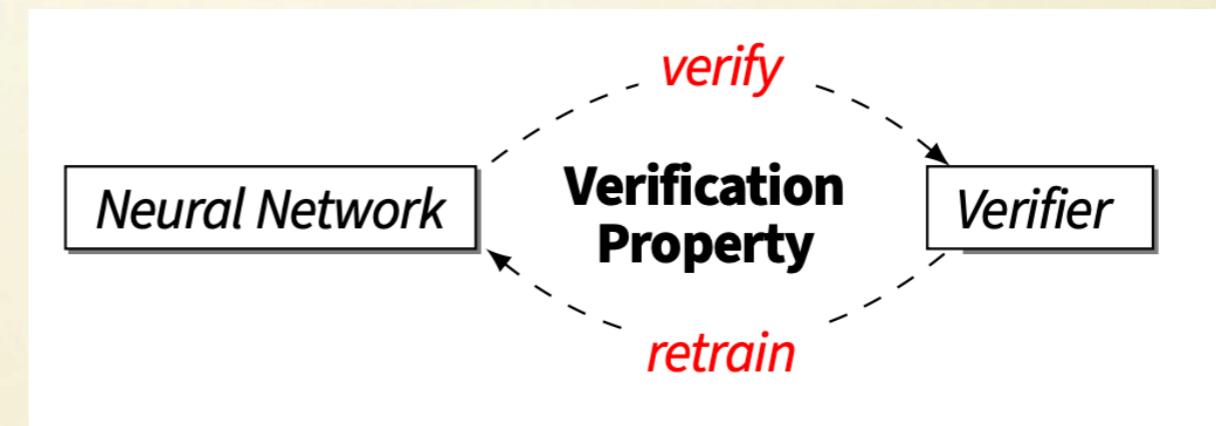
Formal verification with lenses

Going Forward

Formal verification with lenses

Continuous verification:

Loss function encapsulates
the property to verify



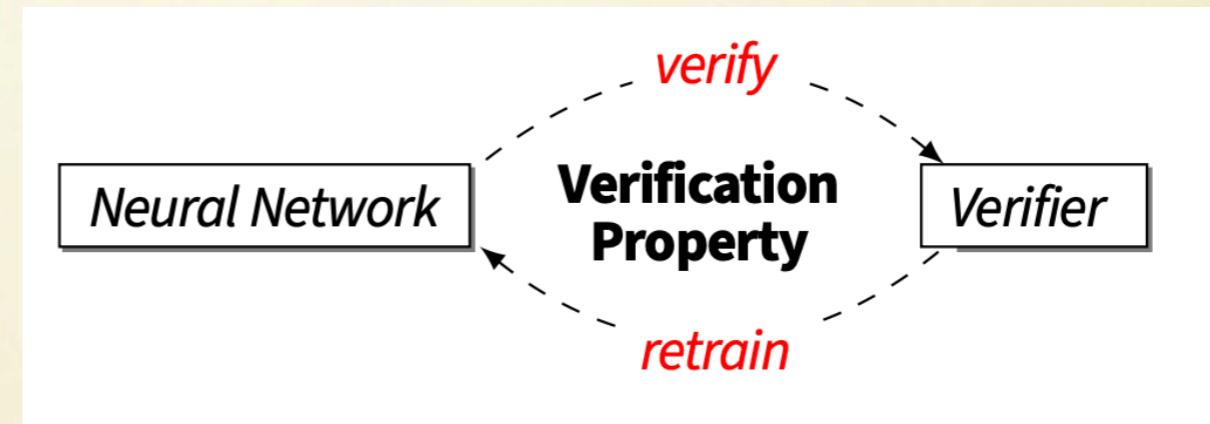
Going Forward

Formal verification with lenses

Continuous verification:

Loss function encapsulates
the property to verify

Specification language:
Differentiable logics



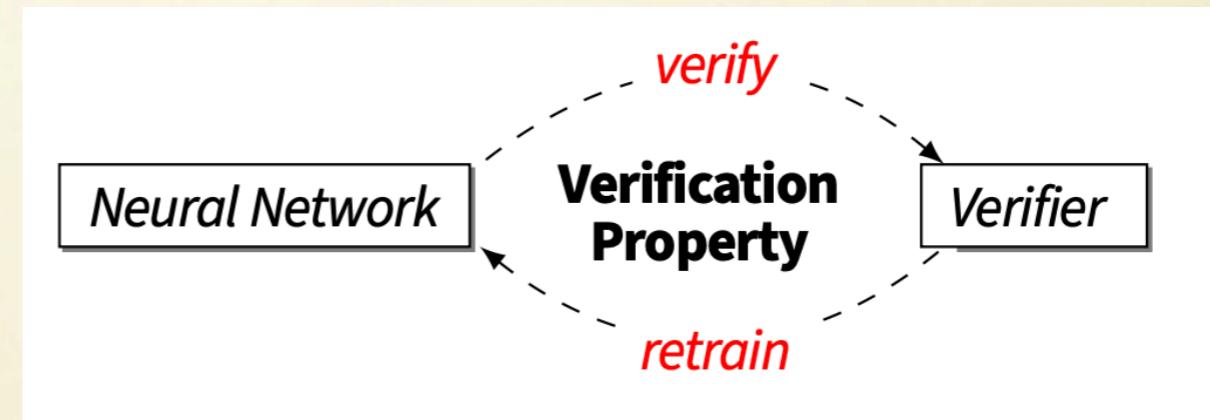
Going Forward

Formal verification with lenses

Continuous verification:

Loss function encapsulates
the property to verify

Specification language:
Differentiable logics



Lens-based semantics
of the continuous verification cycle

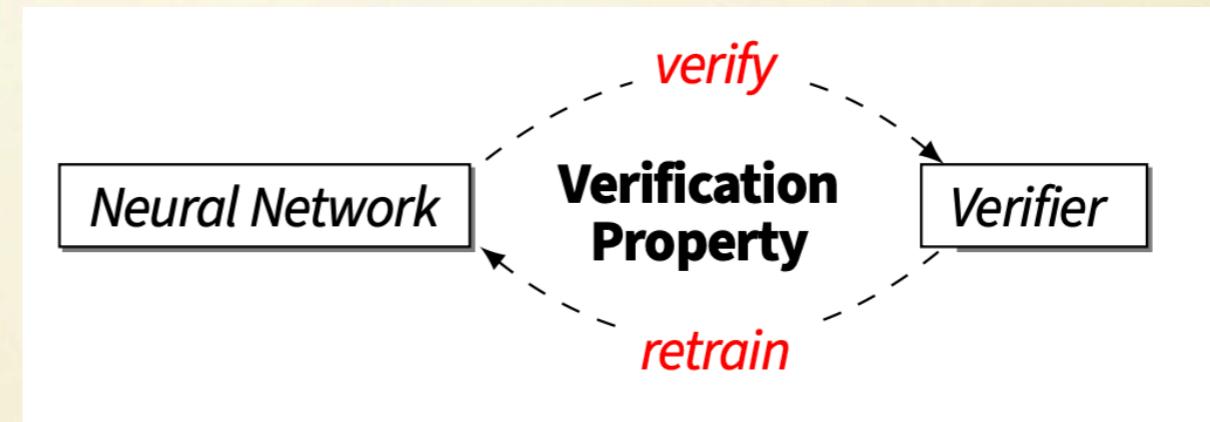
Going Forward

Formal verification with lenses

Continuous verification:

Loss function encapsulates
the property to verify

Specification language:
Differentiable logics



Lens-based semantics
of the continuous verification cycle

Lens-based specification language for
structural properties to be verified

References

Implementation

<https://github.com/statusfailed/numeric-optics-python/>

<http://catgrad.com/p/reverse-derivative-ascent>

<https://yarrow.id/> (Ethereum Foundation Grant)

(Partial) bibliography

G. Cruttwell, B. Gavranovic, N. Ghani, P. Wilson, F. Zanasi - *Categorical Foundations of Gradient-Based Learning*, ESOP 2022

P. Wilson, F. Zanasi - *Reverse Derivative Ascent: a Categorical Approach to Learning Boolean Circuits*, ACT 2020

P. Wilson, F. Zanasi - *Categories of Differentiable Polynomial Circuits for Machine Learning*, ICGT 2022